

Parallel Computation

Mark Greenstreet


CpSc 418 – Jan. 4, 2017

Outline:

- Why Parallel Computation Matters
- Course Overview
- Our First Parallel Program
- The next month
- Table of Contents



Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license <https://creativecommons.org/licenses/by/4.0/>



Clock Speed and Power of Intel Processors vs. Year Released

[\[Wikipedia CPU-Power, 2011\]](#)

- In the good-old-days, processor performance doubled roughly every 1.5 years.
- Single thread performance has seen small gains in the past 14 years.
 - Too bad. If it had, we would have 1000GHz CPUs today. ☹
- Need other ways to increase performance.

Power

- CPUs with faster clocks use more energy per operation than slower ones.
- For mobile devices: high power limits battery life.
- For desktop computers and gaming consoles: cooling high-power chips requires expensive hardware.
- For large servers and clouds, the power bill is a large part of the operating cost.

Parallel Computers

- Mobile devices:
 - multi-core to get good performance on apps and reasonable battery life.
 - many dedicated “accelerators” for graphics, WiFi, networking, video, audio, ...
- Desktop computers
 - multi-core for performance
 - separate GPU for graphics
- Commercial servers
 - multiple, multi-core processors with shared memory.
 - large clusters of machines connected by dedicated networks.

Outline

- Why Does Parallel Computation Matter?
- Course Overview
 - Topics
 - Syllabus
 - The instructor and TAs
 - The textbook(s)
 - Grades
 - Homework: 35%
 - Midterm: 25%
 - Final: 40%
 - Mini-Assignments: see description on [slide 19](#)
 - Bug Bounties: see description on [slide 20](#)
 - Plagiarism – please don't
 - Learning Objectives
- Our First Parallel Program

Topics

- Parallel Architectures
- Parallel Performance
- Parallel Algorithms
- Parallel Programming Frameworks

Parallel Architectures

- There isn't one, standard, parallel architecture for everything. We have:
 - Multi-core CPUs with a shared-memory programming model. Used for mobile device application processors, laptops, desktops, and many large data-base servers.
 - Networked clusters, typically running linux. Used for web-servers and data-mining. Scientific supercomputers are typically huge clusters with dedicated, high-performance networks.
 - Domain specific processors
 - GPUs, video codecs, WiFi interfaces, image and sound processing, crypto engines, network packet filtering, and so on.
- As a consequence, **there isn't one, standard, parallel programming paradigm.**

Parallel Performance

The incentive for parallel computing is to do things that wouldn't be practical on a single processor.

- Performance matters.
- We need good models:
 - Counting operations can be very misleading – “adding is free.”
 - Communication and coordination are often the dominant costs.
- We need to measure actual execution times of real programs.
 - There isn't a unified framework for parallel program performance analysis that works well in practice.
 - It's important to measure actual execution time and identify where the bottlenecks are.
- Key concepts with performance:
 - Amdahl's law, linear speed up, overheads.

Parallel Algorithms

- We'll explore some old friends in a parallel context:
 - Sum of the elements of an array
 - matrix multiplication
 - dynamic programming.
- And we'll explore some uniquely parallel algorithms:
 - Bitonic sort
 - mutual exclusion
 - producer consumer

Parallel Programming Frameworks

- Erlang: functional, message passing parallelism
 - Avoids many of the common parallel programming errors: races and side-effects. You can write Erlang programs with such bugs, but it takes extra effort (esp. for the examples we consider).
 - Allows a simple presentation of many ideas.
 - But it's slow, for many applications, when compared with C or C++.
 - OTOH, it finds real use in large-scale distributed systems.
- CUDA: your graphics card is a super-computer
 - Excellent performance on the “right” kind of problem.
 - The data-parallel model is simple, and useful.

Syllabus

- January: Erlang
 - Jan. 4– 9: Course overview, intro. to Erlang programming.
 - Jan. 11–18: Parallel programming in Erlang, reduce and scan.
 - Jan. 20–27: Parallel architectures
 - Jan. 29–Feb. 6: Performance analysis
- February: Erlang, Midterm
 - Feb. 8–17: Sorting
 - Feb. 20–19: Midterm break.
 - Feb. 27: Midterm Review
 - Mar. 1: Midterm
- March: CUDA and other topics
 - Mar. 3–10: Introduction to SIMD and CUDA.
 - Mar. 13–24: More algorithms in CUDA (and a bit of Erlang)
 - Mar. 27–Apr. 6: Map-Reduce, Mutual Exclusion, & More Fun.
- Note: I'll make adjustments to this schedule as we go.

Administrative Stuff – Who

- The instructors
 - Mark Greenstreet, mrg@cs.ubc.ca
 - ICCS 323, (604) 822-3065
 - Office hours: Tuesdays, 1pm – 2:30pm, ICCS 323
 - Ian Mitchell, mitchell@cs.ubc.ca
 - ICCS 217, (604) 822-2317
 - Office hours: Fridays, 12noon – 1pm, ICCS 217
- The TAs
 - Devon Graham, dgraham@cs.ubc.ca
 - Chenxi Liu, chenxi@cs.ubc.ca
 - Carolyn Shen, shen.carolyn@gmail.com
 - Brenda Xiong, krx.sky@gmail.com
- Course webpage: <http://www.ugrad.cs.ubc.ca/~cs418>.
- Online discussion group: on [piazza](#).

Textbook(s)

- For Erlang: *Learn You Some Erlang For Great Good*, Fred Hébert,
 - Free! On-line at <http://learnyousomeerlang.com>.
 - You can buy the dead-tree edition at the same web-site if you like.
- For CUDA: *Programming Massively Parallel Processors: A Hands-on Approach* (2nd or 3rd ed.), D.B. Kirk and W.M.W. Hwu.
 - Please get a copy by late February – I'll assign readings starting after the midterm. It's available at [amazon.ca](#) and many other places.
- I'll hand-out copies of some book chapters:
 - Principles of Parallel Programming* (chap. 5), C. Lin & L. Snyder – for the reduce and scan algorithms.
 - An Introduction to Parallel Programming* (chap. 2), P.S. Pacheco – for a survey of parallel architectures.
 - Probably a few journal, magazine, or conference papers.

Why so many texts?

- There isn't one, dominant parallel architecture or programming paradigm.
- The Lin & Snyder book is a great, paradigm independent introduction,
- But, I've found that descriptions of real programming frameworks lack the details that help you write real code.
- So, I'm using several texts, but
 - You only have to buy one! ☺

Grades

Homework: 35% roughly one assignment every two weeks
Midterm: 25% March 1, in class
Final: 40%
Mini-Assignments: see description on [slide 19](#)
Bug Bounties: see description on [slide 20](#)

Homework

- Collaboration policy
 - You are welcome and encouraged to discuss the homework problems with other students in the class, with the TAs and me, and find relevant material in the text books, other book, on the web, etc.
 - You are expected to work out your own solutions and write your own code. Discussions as described above are to help understand the material. Your solutions must be your own.
 - You must properly cite your collaborators and any outside sources that you used. You don't need to cite material from class, the textbooks, or meeting with the TAs or instructor. See [slide 22](#) for more on the plagiarism policy.
- Late policy
 - Each assignment has an “early bird” date before the main date. Turn in your assignment by the early-bird date to get a 5% bonus.
 - No late homework accepted.

Exams

- Midterm, in class, on March 1.
- Final exam will be scheduled by the registrar.
- Both exams are open book, open notes, open homework and solutions – open anything printed on paper.
 - You can bring a calculator.
 - No communication devices: laptops, tablets, cell-phones, etc.

Mini-Assignments

- Mini-assignments
 - Worth 20% of points missed from HW and exams.
 - If your raw grade is 90%, you can get at most 2% from the minis. Missing one or two isn't a big deal.
 - If your raw grade is 70%, you can get 6% from the minis. This can move your letter grade up a notch (e.g. C+ to B-).
 - If your raw grade is 45%, you can get up to 11% from the minis. Do the mini-assignments – I hate turning in failing grades.
 - The first is at <http://www.ugrad.cs.ubc.ca/~cs418/2016-2/mini/1/mini1.pdf>, and due Jan. 9.
 - If you are on the course waitlist, we will select from the students who submit acceptable solutions to **Mini Assignment 1** to fill any slots that open up.

Bug Bounties

- If I make a mistake when stating a homework problem, then the first person to report the error gets extra credit.
 - If the error would have prevented solving the problem, then the extra credit is the same as the value of the problem.
 - Smaller errors get extra credit in proportion to their severity.
- Likewise, bug bounties are awarded (as homework extra credit) for finding errors in mini-assignments, lecture slides, the course web-pages, code I provide, etc.
- The midterm and final have bug bounties awarded in midterm and final exam points respectively.
- If you find an error, report it.
 - Suspected errors in homework, lecture notes, and other course materials should be posted to piazza.
 - The first person to post a bug gets the bounty.
 - Bug-bounties reward you for looking at the HW when it first comes out, and not waiting until the day before it is due.

Grades: the big picture

$$\begin{aligned} \text{RawGrade} &= 0.35 * HW + 0.25 * \text{MidTerm} + 0.40 * \text{Final} \\ \text{MiniBonus} &= 0.20 * (1 - \min(\text{RawGrade}, 1)) * \text{Mini} \\ BB &= 0.35 * BB_{HW} + 0.25 * BB_{MT} + 0.40 * BB_{FX} \\ \text{CourseGrade} &= \min(\text{RawGrade} + \text{MiniBonus} + BB, 1) * 100\% \end{aligned}$$

Plagiarism

- I have a very simple criterion for plagiarism:
Submitting the work of another person, whether that be another student, something from a book, or something off the web and representing it as your own is plagiarism and constitutes academic misconduct.
- If the source is clearly cited, then it is not academic misconduct.
If you tell me “This is copied word for word from Jane Foo's solution” that is not academic misconduct. It will be graded as one solution for two people and each will get half credit. I guess that you could try telling me how much credit each of you should get, but I've never had anyone try this before.
- I encourage you to discuss the homework problems with each other.
If you're brainstorming with some friends and the key idea for a solution comes up, that's OK. In this case, add a note to your solution that lists who you collaborated with.
- More details at:
 - <http://www.ugrad.cs.ubc.ca/~cs418/plagiarism.html>
 - <http://www.ugrad.cs.ubc.ca/~cs418/plagiarism/faq.html>

Learning Objectives (1/2)

- Parallel Algorithms
 - Familiar with parallel patterns such as reduce, scan, and tiling and can apply them to common parallel programming problems.
 - Can describe parallel algorithms for matrix operations, sorting, dynamic programming, and process coordination.
- Parallel Architectures
 - Can describe shared-memory, message-passing, and SIMD architectures.
 - Can describe a simple cache-coherence protocol.
 - Can identify how communication latency and bandwidth are limited by physical constraints in these architectures.
 - Can describe the difference between bandwidth and inverse latency, and how these impact parallel architectures.

Learning Objectives (2/2)

- Parallel Performance
 - Understands the concept of “speed-up”: can calculate it from simple execution models or measured execution times.
 - Can identify key bottlenecks for parallel program performance including communication latency and bandwidth, synchronization overhead, and intrinsically sequential code.
- Parallel Programming Frameworks
 - Can implement simple parallel programs in Erlang and CUDA.
 - Can describe the differences between these paradigms.
 - Can identify when one of these paradigms is particularly well-suited (or badly suited) for a particular application.

Lecture Outline

- Why Does Parallel Computation Matter?
- Course Overview
- Our First Parallel Program
 - Erlang quick start
 - Count 3s
 - Counting 3's in parallel
 - The root process
 - Spawning worker processes
 - The worker processes
 - Running the code

Erlang Intro – very abbreviated!

- Erlang is a functional language:
 - Variables are given values when declared, and the value never changes.
 - The main data structures are lists, [Head | Tail], and tuples (covered later).
 - Extensive use of pattern matching.
- The source code for the examples in this lecture is available at:
<http://www.ugrad.cs.ubc.ca/~cs418/2016-2/lecture/01-06/code.html>

Lists

- [1, 4, 9, 16, 25, 36, 49, 64, 81, 100] is a list of 10 elements.
- If L1 is a list, then [0 | L1] is the list obtained by prepending the element 0 to the list L1. In more detail:

```
1> L1 = [1, 4, 9, 16, 25, 36, 49, 64, 81, 100].
[1, 4, 9, 16, 25, 36, 49, 64, 81, 100]
2> L2 = [0 | L1].
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81, 100]
3> L3 = [0, L1].
[0, [1, 4, 9, 16, 25, 36, 49, 64, 81, 100]]
```
- Of course, we traverse a list by using recursive functions:

```
sum(List) ->
  if (length(List) == 0) -> 0;
  (length(List) > 0) -> hd(List) + sum(tl(List))
end.

length(L) returns the number of elements in list L.
hd(L) returns the first element of list L (the head), and throws an exception if L is the empty list.
hd([1, 2, 3]) = 1. hd([1]) = 1 as well.
tl(L) returns the list of all elements after the first (the tail).
tl([1, 2, 3]) = [2, 3]. tl([1]) = [].
```

- See `sum_wo_pm` (“sum without pattern matching”) in [simple.erl](#)

We can use Erlang's pattern matching instead of the `if` expression:


- ```
sum([]) -> 0;
sum([Head | Tail]) -> Head + sum(Tail).
```
- `sum([Head | Tail])` matches any non-empty list with `Head` being bound to the value of the first element of the list, and `Tail` being bound to the list of all the other elements.
  - More generally, we can use patterns to identify the different cases for a function.
  - This can lead to very simple code where function definitions follow the structure of their arguments.
  - See `sum` in [simple.erl](#)

Given an array (or list) with `N` items, return the number of those elements that have the value 3.

- ```
count3s([]) -> 0;
count3s([3 | Tail]) -> 1 + count3s(Tail);
count3s([_Other | Tail]) -> count3s(Tail).
```
- We'll need to put the code in an erlang module. See [count3s.erl](#) for the details.
 - To generate a list of random integers, [count3s.erl](#) uses the function `rlist(N, M)` from [course Erlang library](#) that returns a list of `N` integers randomly chosen from 1..M.

```
bash-3.2$ erl
Erlang/OTP 18 [erts-7.0] [source] ...
Eshell V7.0 (abort with ^G)

1> c(count3s).
{ok, count3s}
2> L20 = count3s:rlist(20,5).
[3,4,5,3,2,3,5,4,3,3,1,2,4,1,3,2,3,3,1,3]
3> count3s:count3s(L20).
9
4> count3s:count3s(count3s:rlist(1000000,10)).
99961
5> q().
ok
6> bash-3.2$
```



The code is in <http://www.ugrad.cs.ubc.ca/~cs418/2016-2/lecture/01-04/src/count3s.erl>

Preview of the next month

- January 6: Introduction to Erlang Programming**
- Reading: [Learn You Some Erlang](#), the first eight sections – [Introduction](#) through [Security](#). Feel free to skip the stuff on [bit syntax](#) and [binary comprehensions](#).
- January 9: Processes and Messages**
- Reading: [Learn You Some Erlang](#), Higher Order Functions and [The Trampoline Guide](#) – through [More on Mailbox Processing](#)
- Homework: [Homework 1 goes out \(due Jan. 18\)](#) – Erlang programming
- Mini-Assignment: [Mini-Assignment 1 due 10:00am](#)
- Mini-Assignment 2 goes out (due Jan. 13)
- January 11: Reduce**
- Reading: [Learn You Some Erlang](#), Errors and Exceptions through [A Short Visit to Common Data Structures](#)
- January 13: Scan**
- Reading: Lin & Snyder, chapter 5, pp. 112–125
- Mini-Assignment: [Mini-Assignment 2 due 10:00am](#)
- January 16: Generalized Reduce and Scan**
- Homework: [Homework 1 deadline for early-bird bonus \(11:59pm\)](#)
- [Homework 2 goes out \(due Feb. 1\)](#) – Reduce and Scan
- January 18: Reduce and Scan Examples**
- Homework: [Homework 1 due 11:59pm](#)
- January 20–27: Parallel Architecture**
- January 29–February 6: Parallel Performance**

Review Questions

- Name one, or a few, key reasons that parallel programming is moving into mainstream applications.
- How does the impact of your mini assignment total on your final grade depend on how you did on the other parts of the class?
- What are bug-bounties?
- What is the count 3's problem?
- How did we measure running times to compute speed up?
- Why did one approach show a speed-up greater than the number of cores used?
- Why did the other approach show that the parallel version was **slower** than the sequential one?

Supplementary Material

- [Erlang Resources](#)
- [Bibliography](#)
- [Table of Contents](#) – at the end!!!

Erlang Resources

- Learn You Some Erlang
<http://learnyousomeerlang.com>
An on-line book that gives a very good introduction to Erlang. It has great answers to the “Why is Erlang this way?” kinds of questions, and it gives realistic assessments of both the strengths and limitations of Erlang.
- Erlang Examples:
<http://www.ugrad.cs.ubc.ca/~cs418/2012-1/lecture/09-08.pdf>
My lecture notes that walk through the main features of Erlang with examples for each. Try it with an Erlang interpreter running in another window so you can try the examples and make up your own as you go. This will cover everything you'll need to make it through all (or most) of what we'll do in class, but it doesn't explain how to think in Erlang as well as “Learn You Some Erlang” or Armstrong's Erlang book (next slide).

More Erlang Resources

- The erlang.org tutorial
http://www.erlang.org/doc/getting_started/users_guide.html
Somewhere between my “Erlang Examples” and “Learn You Some Erlang”.
- Erlang Language Manual
http://www.erlang.org/doc/reference_manual/users_guide.html
My go-to place when looking up details of Erlang operators, etc.
- On-line API documentation:
<http://www.erlang.org/erldoc>
- The book: *Programming Erlang: Software for a Concurrent World*, Joe Armstrong, 2007.
<http://pragprog.com/book/iserlang/programming-erlang>
Very well written, with lots of great examples. More than you'll need for this class, but great if you find yourself using Erlang for a big project.
- More resources listed at <http://www.erlang.org/doc.html>.

Getting Erlang

- You can run Erlang by giving the command `erl` on any departmental machine. For example:
 - Linux: `bowen, thetis, lin01, ..., lin25, ...`, all machines above are `.ugrad.cs.ubc.ca`, e.g. `bowen.ugrad.cs.ubc.ca`, etc.
- You can install Erlang on your computer
 - Erlang solutions provides packages for Windows, OSX, and the most common linux distros
 - <https://www.erlang-solutions.com/resources/download.html>
Note: some linux distros come with Erlang pre-installed, but it might be an old version. You should probably install from the link above.

Starting Erlang

- Start the Erlang interpreter.

```
thetis % erl
Erlang/OTP 18 [erts-7.0] [source] ...
Eshell V7.0 (abort with ^G)

1> 2+3.
5
2>
```
- The Erlang interpreter evaluates expressions that you type.
- Expressions end with a “.” (period).

Bibliography

- Krste Asanovic, Ras Bodik, et al.
The landscape of parallel computing research: A view from Berkeley.
Technical Report UCB/ECS-2006-183, Electrical Engineering and Computer Science Department, University of California, Berkeley, December 2006.
<http://www.eecs.berkeley.edu/Pubs/TechRpts/2006/ECS-2006-183.pdf>.
- Microprocessor quick reference guide.
<http://www.intel.com/pressroom/kits/quickrefyr.htm>, June 2013.
accessed 29 August 2013.
- List of CPU power dissipation.
http://en.wikipedia.org/wiki/List_of_CPU_power_dissipation, April 2011.
accessed 26 July 2011.

Table of Contents (1/2)

- [Motivation](#)
- [Course Overview](#)
 - [Topics](#)
 - [Computer Architecture](#)
 - [Performance Analysis](#)
 - [Algorithms](#)
 - [Languages, Paradigms, and Frameworks](#)
 - [Syllabus](#)
 - [Course Administration – who's who](#)
 - [The Textbook\(s\)](#)
 - [Grades](#)
 - [Homework](#)
 - [Midterm and Final Exams](#)
 - [Mini-Assignments](#)
 - [Bug Bounties](#)
 - [Plagiarism Policy](#)
 - [Learning Objectives](#)

Table of Contents (2/2)

- [Our First Parallel Program](#)
 - [Introduction to Erlang](#)
 - [The Count 3s Example](#)
- [Preview of the next month](#)
- [Review of this lecture](#)
- [Supplementary Material](#)
 - [Erlang Resources](#)
 - [Bibliography](#)
 - [Table of Contents](#)

Introduction to Erlang

Mark Greenstreet

CpSc 418 – January 6, 2016

Outline:

- [Erlang Basics](#)
- [Functional programming](#)
- [Example, sorting a list](#)
- [Functions](#)
- [Supplementary Material](#)
- [Table of Contents](#)

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license <https://creativecommons.org/licenses/by/4.0/>

Objectives

- Learn/review key concepts of functional programming:
 - Referential transparency.
 - Structuring code with functions.
- Introduction to Erlang
 - Basic data types and operations.
 - Program design by structural decomposition.
 - Writing and compiling an Erlang module.

Erlang Basics

- Numbers:
 - Numerical Constants: `1, 8#31, 1.5, 1.5e3,` but not: `1, or .5`
 - Arithmetic: `+, -, *, /, div, band, bor, bnot, bsl, bsr, bxor`
- Booleans:
 - Comparisons: `=, <, >, <=, >=, <=, >=`
 - Boolean operations (strict): `and, or, not, xor`
 - Boolean operations (short-circuit): `andalso, or_else`
- Atoms:
 - Constants: `x, 'big bog-2'`
 - Operations: tests for equality and inequality. Therefore pattern matching.

Lists and Tuples

- Lists:
 - Construction: `[1, 2, 3], [Element1, Element2, ..., ElementN | Tail]`
 - Operations: `hd, tl, length, ++, --`
 - Erlang's list library: <http://erlang.org/doc/man/lists.html>: `all, any, filter, foldl, foldr, map, nth, nthtail, seq, sort, split, zipwith`, and **many** more.
- Tuples:
 - Construction: `{1, dog, "called Rover"}`
 - Operations: `element, setelement, tuple_size`.
 - Lists vs. Tuples:
 - Lists are typically used for an **arbitrary** number of elements of the same “type” – like arrays in C, Java, ...
 - Tuples are typically used for a **fixed** number of elements of the varying “types” – like a `struct` in C or an object in Java.

Strings

- What happened to strings?!
- Well, they're lists of integers.
 - This can be annoying. For example,

```
1> [102, 111, 111, 32, 98, 97, 114].
"foo bar"
```
 - By default, Erlang prints lists of integers as strings if every integer in the list is the ASCII code for a “printable” character.
 - [Learn You Some Erlang](#) discusses strings in the “Don't drink too much Kool-Aid” box for [lists](#).

Functional Programming

- Imperative programming (C, Java, Python, ...) is a programming model that corresponds to the von Neumann computer:
 - A program is a sequence of statements.
 - In other words, a program is a recipe that gives a step-by-step description of what to do to produce the desired result.
 - Typically, the operations of imperative languages correspond to common machine instructions.
 - Control-flow (`if, for, while`, function calls, etc.)
 - Each control-flow construct can be implemented using branch, jump, and call instructions.
 - This correspondence between program operations and machine instructions simplifies implementing a good compiler.
- Functional programming (Erlang, lisp, scheme, Haskell, ML, ...) is a programming model that corresponds to mathematical definitions.
 - A program is a collection of definitions.
 - These include definitions of expressions.
 - Expressions can be evaluated to produce results.
- See also: [the LYSE explanation](#).

Erlang Makes Parallel Programming Easier

- Erlang is functional
 - Each variable gets its value when it's declared – it **never** changes.
 - Erlang eliminates many kinds of races – another process can't change the value of a variable while you're using it, because the values of variables never change.
- Erlang uses message passing
 - Interactions between processes are under explicit control of the programmer.
 - Fewer races, synchronization errors, etc.
- Erlang has simple mechanisms for process creation and communication
 - The structure of the program is not buried in a large number of calls to a complicated API.

Big picture: Erlang makes the issues of parallelism in parallel programs more apparent and makes it easier to avoid many common pitfalls in parallel programming.

Referential Transparency

- This notion that a variable gets a value when it is declared and that the value of the variable never changes is called **referential transparency**.
 - You'll here me use the term many times in class – I thought it would be a good idea to let you know what it means. ☺
- We say that the value of the variable is **bound** to the variable.
- Variables in functional programming are much like those in mathematical formulas:
 - If a variable appears multiple places in a mathematical formula, we assume that it has the same value everywhere.
 - This is the same in a functional program.
 - This is **not** the case in an imperative program. We can declare `x` on line 17, assign it a value on line 20; and assign it another value on line 42.
 - The value of `x` when executing line 21 is different than when executing line 43.

Loops violate referential transparency

- ```
// vector dot-product
sum = 0.0;
for(i = 0; i < a.length; i++)
 sum += a[i] * b[i];

// merge as in merge-sort
while(a != null && b != null) {
 if(a.key <= b.key) {
 last->next = a;
 last = a;
 a = a->next;
 last->next = null;
 } else {
 ...
 }
}
```
- Loops rely on changing the values of variables.
  - Functional programs use recursion instead.
  - See also [the LYSE explanation](#).

## Life without loops

- Use recursive functions instead of loops.
- ```
dotProd([], []) -> 0;
dotProd([A | A1], [B | B1]) -> A*B + dotProd(A1, B1).
```
- Functional programs use recursion instead of iteration:

```
dotProd([], []) -> 0;
dotProd([A | A1], [B | B1]) -> A*B + dotProd(A1, B1).
```
 - Anything you can do with iteration can be done with recursion.
 - But the converse is not true (without dynamically allocating data structures).
 - Example: tree traversal.

Example: Sorting a List

- The simple cases:
 - Sorting an empty list: `sort([]) -> []`
 - Sorting a singleton list: `sort([A]) -> [A]`
- How about a list with more than two elements?
 - Merge sort?
 - Quick sort?
 - Bubble sort (**NO WAY! Bubble sort is DISGUSTING!!!**).
- Let's figure it out.

Merge sort: Erlang code

- If a list has more than one element:
 - Divide the elements of the list into two lists of roughly equal length.
 - Sort each of the lists.
 - Merge the sorted list.
- In Erlang:

```
sort([]) -> [];
sort([A]) -> [A];
sort([A | Tail]) ->
  ([L1, L2] = split([A | Tail]),
  L1sorted = sort(L1),
  L2sorted = sort(L2),
  merge(L1sorted, L2sorted)).
```
- Now, we just need to write `split`, and `merge`.

split (L)

- Identify the cases and their return values according to the shape of `L`:
- ```
% If L is empty (recall that split returns a tuple of two lists):
split([]) -> {}

% If L
split([A | Tail]) ->

% If L
split([A | Tail]) ->
```

## Finishing merge sort

- An exercise for the reader – see [slide 29](#).
- Sketch:
  - Write `merge(List1, List2) -> List12` – see [slide 30](#)
  - Write an Erlang module with the `sort`, `split`, and `merge` functions – see [slide 31](#)
  - Run the code – see [slide 33](#)

Fun with Functions

- Programming with patterns
  - often, the code just matches the shape of the data
  - like CPSC 110, but pattern matching makes it obvious
  - see [slide 16](#)
- Fun expressions
  - in-line function definitions
  - see [slide 17](#)
- Higher-order functions
  - encode common control-flow patterns
  - see [slide 18](#)
- List comprehensions
  - common operations on lists
  - see [slide 19](#)
- Tail call elimination
  - makes recursion as fast as iteration (in simple cases)
  - see [slide 20](#)

```
% leafCount: count the number of leaves of a tree represented by a nested list
leafCount([]) -> 0; % base case - an empty list has no leaves

leafCount([Head | Tail]) -> % recursive case
 leafCount(Head) + leafCount(Tail);
leafCount(.Leaf) -> 1; % the other base case - .Leaf is not a list
```

- Let's try it
  - examples: leafCount([1, 2, [3, 4, []], [5, [6, banana]])].
- Notice how we used **patterns** to show how the recursive structure of leafCount follows the shape of the tree.
- See **Pattern Matching** in *Learn You Some Erlang* for more explanation and examples.
- Style guideline: if you're writing code with lots of `if`'s `hd`'s, and `tl`'s, you should think about it and see if using patterns will make your code simpler and clearer.

```
fun(X, Y) -> X*X + Y*Y end. % fun ... end creates an anonymous function"
#funCrlEval,12,52032458> % ok, guess, but what can I do with it?
6> F = fun(X, Y) -> X*X + Y*Y end.
#funCrlEval,12,52032458>
F(3, 4).
25
6> Factorial = % We can even write recursive fun expressions!
fun Fact(0) -> 1;
 Fact(N) when is_integer(N), N > 0 -> N*Fact(N-1)
end.
7> Factorial(3).
6
8> Fact(3).
% 1: variable 'Fact' is unbound
9> Factorial(-2).
** exception error: no function clause matching
 erl_eval:1-inside-an-interpreted-fun*(-2)
10> Factorial(frog).
** exception error: no function clause matching
 erl_eval:1-inside-an-interpreted-fun*(-frog)
```

See **Anonymous Functions** in *Learn You Some Erlang* for more explanation and examples.

Higher-Order Functions

- Lists: `map(Fun, List)` apply `Fun` to each element of `List` and return the resulting list.
  - `11> lists:map(fun(X) -> 2*X+1 end, [1, 2, 3]).`  
`[3, 5, 7]`
- `lists:fold(Fun, Acc0, List)` use `Fun` to combine all of the elements of `List` in left-to-right order, starting with `Acc0`.
  - `12> lists:fold1(fun(X, Y) -> X*Y end, 100, [1, 2, 3]).`  
`106`
- For more explanation and examples:
  - See **Higher Order Functions** in *Learn You Some Erlang*.
  - See the **lists** module in the Erlang standard library. Examples include
    - `all(Pred, List)`: true iff `Pred` evaluates to true for **every** element of `List`.
    - `any(Pred, List)`: true iff `Pred` evaluates to true for **any** element of `List`.
    - `fold(Fun, Acc0, List)`: like `fold1` but combines elements in right-to-left order.

List Comprehensions

- Map and filter are such common operations, that Erlang has a simple syntax for such operations.
- It's called a **List Comprehension**:
  - `[Expr | Var <- List, Cond, ...]`.
  - `Expr` is evaluated with `Var` set to each element of `List` that satisfies `Cond`.
  - Example:

```
13> > count3s::list(5, 1000).
[444,724,946,502,312].
14>[X*X | X <- R, X rem 3 == 0].
[197136,97344].
```
- See also **List Comprehensions** in *LYSE*.

Head vs. Tail Recursion

- I wrote two versions of computing the sum of the first `N` natural numbers:

```
sumN(0) -> 0; % "head recursive"
sumN(N) -> N + sumN(N-1).

sumN(0) -> sumN(N, 0).
sumN(0, Acc0) -> Acc0; % "tail recursive"
sumN(N, Acc) -> sumN(N-1, N+Acc).
```
- Here are some run times that I measured:

| N    | head  | tail  | N    | head    | tail  |
|------|-------|-------|------|---------|-------|
| 1K   | 21µs  | 19µs  | 1M   | 21ms    | 11ms  |
| 10K  | 178µs | 114µs | 10M  | 1.7s    | 115ms |
| 100K | 1.7ms | 1.1ms | 100M | 28s     | 1.16s |
|      |       |       | TG   | > 8 min | 11.6s |

Summary

- Why Erlang?
  - Functional – avoid complications of side-effects when dealing with concurrency.
  - But, we can't use imperative control flow constructions (e.g. loops).
    - Design by declaration: look at the structure of the data.
    - More techniques coming in upcoming lectures.
- Sequential Erlang
  - Lists, tuple, atoms, expressions
  - Using structural design to write functions: example sorting.
  - Functions: patterns, higher-order functions, head vs. tail recursion.

Review

January 9: Processes and Messages

Reading: [Learn You Some Erlang, Higher Order Functions and The Hitchhiker's Guide](#), through [Klue on Why Erlang](#)

Homework: **Homework 1 goes out (due Jan. 18)** – Erlang programming

Mini-Assignment: **Mini-Assignment 1 due 10:00am**

Mini-Assignment 2 goes out (due Jan. 13)

January 11: Reduce

Reading: [Learn You Some Erlang, Errors and Exceptions](#) through [A Short Visit to Common Data Structures](#)

January 13: Scan

Reading: Lin & Snyder, chapter 5, pp. 112–125

Mini-Assignment: **Homework 2 due 10:00am**

January 16: Generalized Reduce and Scan

Homework: **Homework 1** deadline for early-bird bonus (11:59pm)

Homework 2 goes out (due Feb. 7) – Reduce and Scan

January 18: Reduce and Scan Examples

Homework: **Homework 1** due 11:59pm

January 20-27: Parallel Architecture

January 29-February 6: Parallel Performance

February 8-17: Parallel Sorting

Review Questions

- What is the difference between `==` and `:=` ?
- What is an atom?
- Which of the following are valid Erlang variables, atoms, both, or neither?  
`'Foo'`, `foo`, `25`, `'25'`, `'Foo foo'`,  
`"4 score and 7 years ago"`, `X2`,  
`"4 score and 7 years ago"`.
- Draw the tree corresponding to the nested list  
`[X, [Y, Z], 2, [A, B+C, [1, 23]], 54, [![]]]`.
- What is referential transparency?
- Why don't functional languages have loops?
- Use an anonymous function and `lists:filter` to implement the body of `GetEven` below:

```
% GetEven(List) -> Events, where Events is a list consisting of all
% elements of List that are integers and divisible by two.
% Example: GetEven([1, 2, frog, 1000]) -> [2, 1000]
GetEven(List) ->
 you write this part.
```

Supplementary Material

The remaining material is included in the web-version of these slides: <http://www.upgrad.cs.ubc.ca/~cs418/2016-2/lecture/01-06/slides.pdf>

I'm omitting it from the printed handout to save a few trees.

- [Erlang resources](#).
- [Finishing the merge sort example](#).
- [Common mistakes with lists](#) and how to avoid them.
- [A few remarks about atoms](#).
- [Suppressing verbose output](#) when using the Erlang shell.
- [Forgetting variable bindings](#) (only in the Erlang shell).
- [Table of Contents](#).

Erlang Resources

- [LYSE](#) – you should be reading this already!
- Install Erlang on your computer
  - Erlang solutions provides packages for Windows, OSX, and the most common linux distros <http://www.erlang-solutions.com/resources/download.html>
  - Note: some linux distros come with Erlang pre-installed, but it might be an old version. You should probably install from the link above.
- <http://www.erlang.org>
  - Searchable documentation [http://erlang.org/doc/reference\\_manual/users\\_guide.html](http://erlang.org/doc/reference_manual/users_guide.html)
  - Language reference [http://erlang.org/doc/reference\\_manual/users\\_guide.html](http://erlang.org/doc/reference_manual/users_guide.html)
  - Documentation for the standard Erlang library [http://erlang.org/doc/man\\_index.html](http://erlang.org/doc/man_index.html)
- The CPSC 418 Erlang Library
  - Documentation <http://www.upgrad.cs.ubc.ca/~cs418/resources/erl/doc/index.html>
  - Test (source, and pre-compiled .beam) <http://www.upgrad.cs.ubc.ca/~cs418/resources/erl/erl.tgz>

Modules

- To compile our code, we need to put it into a **module**.
- A module is a file (with the extension `.erl`) that contains
  - Attributes: declarations of the module itself and the functions it exports.
    - The module declaration is a line of the form:  
`-module(moduleName).`  
where `moduleName` is the name of the module.
    - Function exports are written as:  
`-export([FunctionName1/arity1,`  
`FunctionName2/arity2, ...]).`
  - The list of functions may span multiple lines and there may be more than one `-export` attribute.
  - `arity` is the number of arguments that the function has. For example, if we define  
`foo(A, B) -> A*A + B.`  
Then we could export `foo` with  
`-export([foo/2, ...]).`
  - There are many other attributes that a module can have. We'll skip the details. If you really want to know, it's all described [here](#).
  - Function declarations (and other stuff) – see the next slide

A module for sort

```
-module(sort).
-export([sort/1]).
% The next -export is for debugging. We'll comment it out later
-export([split/1, merge/2]).
sort([]) -> [];
...

The list of functions may span multiple lines and there may be more than one -export attribute.
```

Misconstructing Lists

Working with `divisible_drop` from the previous slide...

- Now, change the second alternative in the `if` to  
`A rem N /= 0 -> [A, divisible_filter(N, Tail)]`
- Trying the previous test case:

```
> examples:divisible_drop(3, [0, 1, 4, 17, 42, 100]).
[1, 4, [17, [100, []]]]
```
- Moral: If you see a list that is nesting way too much, check to see if you wrote a comma where you should have used a `.`
- Restore the code and then change the second alternative for `divisible_drop` to `divisible_drop(N, [A, Tail]) -> Trying our previous test:`

```
> examples:divisible_drop(3, [0, 1, 4, 17, 42, 100]).
** exception error: no function clause matching...
```

Punctuation

- Erlang has lots of punctuation: commas, semicolons, periods, and `end`.
- It's easy to get syntax errors or non-working code by using the wrong punctuation somewhere.
- Rules of Erlang punctuation:
  - Erlang declarations end with a period: `.`
  - A declaration can consist of several alternatives.
    - Alternatives are separated by a semicolon: `;`
    - Note that many Erlang constructions such as `case`, `fun`, `if`, and `receive` can have multiple alternatives as well.
  - A declaration or alternative can be a block expression
    - Expressions in a block are separated by a comma: `,`
    - The value of a block expression is the last expression of the block.
  - Expressions that begin with a keyword end with `end`
    - `case` Alternatives `end`
    - `fun` Alternatives `end`
    - `if` Alternatives `end`
    - `receive` Alternatives `end`

Forgetting Bindings

- Referential transparency means that bindings are forever.
  - This can be nuisance when using the Erlang shell.
  - Sometimes we assign a value to a variable for debugging purposes.
  - We'd like to overwrite that value later so we don't have to keep coming up with more names
- In the Erlang shell, `f(Variable)` makes the shell "forget" the binding for the variable.

```
12> X = 2+3.
5.
13> X = 2+3.
** exception error: no match of right hand side value 6.
14> f(X).
ok
15> X = 2+3.
6
16>
```

Table of Contents

- [Erlang Basics](#) – basic types and their operations.
- [Functional Programming](#) – referential transparency, recursion instead of loops.
- [Example: Merge Sort](#)
- [Fun with Functions](#) – patterns, anonymous functions, higher-order functions, list comprehensions, head vs. tail recursion
- [Preview of upcoming lectures](#)
- [Review of this lecture](#)
- [Supplementary Material](#)

Processes and Messages

Mark Greenstreet

CpSc 418 – Jan. 9, 2017

Outline:

- [Processes](#)
- [Messages](#)
- [Timing Measurements](#)
- [Preview, Review, etc.](#)
- [Table of Contents](#)

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license <http://creativecommons.org/licenses/by/4.0/>

Higher-Order Functions

- Introduce Erlang's features for concurrency and parallelism
  - Spawning processes.
  - Sending and receiving messages.
- Describe timing measurements for these operations and the implications for writing efficient parallel programs.
  - Communication often dominates the runtime of parallel programs.
- The source code for the examples in this lecture is available here: [procs.erl](#).



- The built-in function `spawn` creates a new process.
- Each process has a process-id, pid.
  - The built-in function `self()` returns the pid of the calling process.
  - `spawn` returns the pid of the process that it creates.
  - The simplest form is `spawn(Fun)`.
    - A new process is created – ‘the child’.
    - The pid of the new process is returned to the caller of `spawn`.
    - The function `Fun` is invoked with no arguments in that process.
    - The parent process and the child process are both running.
    - When `Fun` returns, the child process terminates.

```
hello(N) ->
[
 spawn(fun() -> io:format(
 "hello world from process ~b~n", [N]
) end)
 || I <- lists:seq(1,N)
],
].
```

Running the code:

```
1> c(procs).
{ok,procs}
2> procs:hello(3).
hello world from process 1
hello world from process 3
hello world from process 3
[<0.40.0>,<0.41.0>,<0.42.0>]
```

- To solve tasks in parallel, the processes need to communicate.
- Sending a message: `Pid ! Expr`.
  - `Expr` is evaluated, and the result is sent to process `Pid`.
  - We can send **any** Erlang term: integers, atoms, lists, tuples, ...
- Receiving a message:

```
receive
 Pattern1 -> Expr1;
 Pattern2 -> Expr2;
 ...
 PatternN -> ExprN
end
```

- If there is a pending message for this process that matches one of the patterns,
  - The message is delivered, and the value of the `receive` expression is the value of the corresponding `Expr`.
  - Otherwise, the process blocks until such a message is received.
- Message passing is asynchronous: the sending process can continue its execution before the receiver gets the message.

• The plan:

- We'll spawn a process in the shell for adding two numbers.
- This child process receives two numbers, computes the sum, and sends the result back to the parent.

```
addProc(PFid) ->
receive
 A -> receive
 B ->
 PFid ! A+B
 end
 end.

add() ->
MyPid = self(),
spawn(fun() ->
 addProc(MyPid)
end).
```

```
> Apid = procs:add(1).
<0.44.0>
4> Apid ! 2.
2
5> Apid ! 3.
3
6> receive Sum -> Sum end.
5
```

## Reactive Processes and Tail Recursion

- Often, we want processes that do more than add two numbers together.
- We want processes that wait, receive a message, process the message, and then wait for the next message.
- In Erlang, we do this with recursive functions for the child process:

```
accProc(Tally) ->
 <0.33.0>
 N when is_integer(N) ->
 8> BPid ! 1.
 accProc(Tally+N);
 {Pid, total} ->
 9> BPid ! 2.
 Pid ! Tally,
 accProc(Tally);
 10> BPid ! 3.
 11> BPid ! {self(), total}.
 <<0.33.0>, total>
 spawn(fun() ->
 12> receive T1 -> T1 end.
 accProc(0)
end).
```

## Reactive Processes and Tail Recursion

- Often, we want processes that do more than add two numbers together.
- We want processes that wait, receive a message, process the message, and then wait for the next message.
- In Erlang, we do this with recursive functions for the child process:

```
accProc(Tally) ->
 receive
 N when is_integer(N) ->
 accProc(Tally+N);
 {Pid, total} ->
 Pid ! Tally,
 accProc(Tally);
 end.

accumulator() ->
 <<0.33.0>, total>
 spawn(fun() ->
 accProc(0)
 end).
```

## Message Ordering

- Given two processes, `Proc1` and `Proc2`, messages sent from `Proc1` to `Proc2` are received at `Proc2` in the order in which they were sent.
- Message delivery is reliable: if a process doesn't terminate, any message sent to it will eventually be delivered.
- Other than that, Erlang makes no ordering guarantees.
  - In particular, the triangle inequality is not guaranteed.
  - For example, process `Proc1` can send message `M1` to process `Proc2` and after that send message `M2` to `Proc3`.
  - Process `Proc3` can receive the message `M2`, and then send message `M3` to process `Proc2`.
  - Process `Proc2` can receive messages `M1` and `M3` in either order.
  - Draw a picture to see why this is violates the spirit of the triangle inequality.


## Tagging Messages

- It's a very good idea to include "tags" with messages.
- This prevents your process from receiving an unintended message:  
*"Oh, I forgot that another process was going to send me that. I thought it would happen later."*
- For example, my accumulator might be better if instead of just receiving an integer, it received `{2, add}`


## Timing Measurements

- We write parallel code to solve problems that would take too long on a single CPU.
- To understand performance trade-offs, I'll measure the time for some common operations in Erlang programs:
  - The time to make `N` recursive tail calls.
  - The time to spawn an Erlang process.
  - The time to send and receive messages:
    - Short messages.
    - Messages consisting of lists of varying lengths.

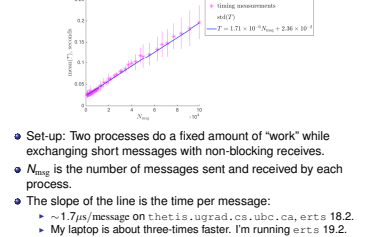
## Tail Call Time




## Process Spawning Time



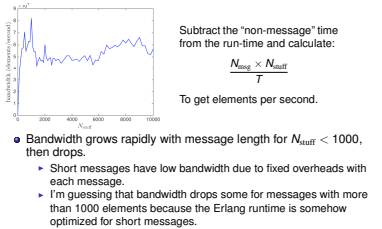
## Send+Receive Time



## Message Time vs. Message Size



## Bandwidth vs. Message Size



## Summarizing the numbers

- Interprocess operations such as `spawn`, `send`, and `receive` are **much** slower than operations within a single process such as + or a function call.
- An Erlang tail call is about 4.7ns, roughly 10 machine instructions.
- An Erlang tail call and add is about 4.7ns, roughly 10 machine instructions.
- Spawning a process is about 200× the cost of a tail call.
- For short messages, send and receive are about 350× the cost of a tail call.
  - The send/receive overhead can be amortized by sending longer message.
  - Each additional list element is about 3× the cost of a tail call.
  - **Beware** of any model that just counts the overhead and ignores the length, or just considers bandwidth and ignores the overhead.
- We will often refer to the ratio of the relationship between the time for interprocess operations and local operations as **big-O**.
  - In practice, big is 100 to 10000 for shared-memory computers.
  - **Big** can be even bigger for other architectures.

## How to Write Efficient Parallel Code

- Think about **communication costs**
  - Message passing is good – it makes communication explicit.
  - Pay attention to both the number of messages and their size.
  - Combining small messages into larger ones often helps.
- Think globally, but **compute locally**
  - Move the computation to the data, not the other way around.
  - Keep the data distributed across the parallel processes.
- Think about **big-O**
  - If `N` is the problem size, you want the computation time to grow faster with `N` than the communication costs.
  - Then, your solution becomes more efficient for larger values of `N`.

## Summary

- Processes are easy to create in Erlang.
  - The `spawn` mechanism can be used to start other processors on the same CPU or on machines spread around the internet.
- Processes communicate through messages
  - Message passing is asynchronous.
  - The receiver can use patterns to select a desired message.
- Reactive processes are implemented with tail-recursive functions.
- Interprocess operations are much slower than local ones
  - This is a key consideration in designing parallel programs.
  - We'll learn **why** when we look at parallel architectures later this month.

## Preview

January 11: Reduce  
Reading: [Learn You Some Erlang, Errors and Exceptions through A Short Visit to Common Data Structures](#)

January 13: Scan  
Reading: Lin & Snyder, chapter 5, pp. 112–125  
Mini-Assignment: [Mini-Assignment 2 due 16:00am](#)

January 16: Generalized Reduce and Scan  
Homework: [Homework 1 deadline for early-bird bonus \(11:59pm\)](#)  
[Homework 2 goes out \(due Feb. 1\) – Reduce and Scan](#)

January 18: Reduce and Scan Examples  
Homework: [Homework 1 due 11:59pm](#)

January 20-27: Parallel Architecture  
January 29-February 6: Parallel Performance  
February 9-17: Parallel Sorting

## Review Questions

- How do you spawn a new process in Erlang?
- What guarantees does Erlang provide (or not) for message ordering?
- Give an example of using patterns to select messages.
- Why is it important to use a tail-recursive function for a reactive process?
  - In other words, why is it a bad idea to use a head-recursive function for a reactive process.
  - The answer isn't explicitly on the slides, but you should be able to figure it out from what we've covered.
- Modify one of the examples in this lecture to use a time-out with one or more `receive` operations. Try it and show that it works.
- Implement the message flushing described in [LYSE](#) to show pending messages on a time-out. Demonstrate how it works.

## Supplementary material

- Debugging concurrent Erlang Code.
- Table of contents.

## Tracing Processes

When you implement a reactive process, it can be handy to trace the execution. Here's a simple approach:

- Add an `io:format` call when entering the function and after matching each receive pattern.
- Example:

```
accProc(Tally) ->
 io:format("~p: accProc(~b)~n", [self(), Tally]),
 receive
 N when is_integer(N) ->
 io:format("~p: received ~b~n", [self(), N]),
 accProc(Tally+N);
 Msg = {Pid, total}
 io:format("~p: received ~p~n", [self(), Msg]),
 Pid ! Tally,
 accProc(Tally)
end.
```
- Try it (e.g. with the example from [slide 7](#)).
- Don't forget to delete (or comment out) such debugging output before releasing your code.

## Time Outs

- If your process is waiting for a message that never arrives, e.g. because
  - You misspelled a tag for a message, or
  - The receive pattern is slightly different than the message that was sent, or
  - Something went wrong in the sending process, and it died before sending the message, or
  - You got the message ordering slightly wrong, and there's a cycle of processes waiting for each other to send something, or
  - ...
- Then your process can wait forever, your Erlang shell can hang, and it's a very unhappy time in life.
- Time-outs can handle these problems more gracefully.
  - See [Time Out in LYSE](#).
  - Note: time-outs are great for debugging. They should be used with great caution elsewhere because they are sensitive to changes in hardware, changes in the scale of the system, and so on.

## Table of Contents

- [Objectives](#)
- [Processes](#)
- [Messages](#)
- [Timing Measurements](#)
- [Summary](#)
- [Preview of upcoming lectures](#)
- [Review of this lecture](#)
- [Supplementary material \(debugging tips\)](#)
- [Table of Contents](#)

## Reduce

Mark Greenstreet

CpSc 418 – Jan. 11, 2017

Outline:

- [Problem Statement](#)
- [Design Guidelines](#)
- [Timing Measurements](#)
- [Preview, Review, etc.](#)
- [Table of Contents](#)

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

## How to Write Efficient Parallel Code

- Think about **communication costs**
  - Message passing is good – it makes communication explicit.
  - Pay attention to both the number of messages and their size.
  - Combining small messages into larger ones often helps.
- Think globally, but **compute locally**
  - Move the computation to the data, not the other way around.
  - Keep the data distributed across the parallel processes.
- Think about **big-O**
  - If `N` is the problem size, you want the computation time to grow faster with `N` than the communication costs.
  - Then, your solution becomes more efficient for larger values of `N`.

## Objectives

- Understand why using a tree-structure for communication improves efficiency.
- Learn how to implement reduce using Erlang processes and messages.
- Learn how to use the `reduce` function in the course Erlang library.

## Problem Statement

Given a list, `L` of `N` values, how can we use `P` processors to efficiently compute the sum of the values of the elements?

Possible in-class exercise:

- Divide class into groups of five or six.
- Hand each group a sheet of numbers. We could arrange the numbers in blocks, or perhaps hand each group a stack of five or six sheets, each of which has around 10 small integers (one to three digits) to add.
- Give them the task that the team that computes the sum of the numbers first wins (perhaps have a bag of M&M's or similar as a prize).
- See what they do.

Now, go back to the observations we made from the previous lecture.

## Summarizing the numbers

- Interprocess operations such as `spawn`, `send`, and `receive` are **much** slower than operations within a single process such as + or a function call.
  - Let's use a tail-call as the cost of an operation within a process.
  - Spawning a process is about 200× the cost of a tail call.
  - For short messages, send and receive are about 350× the cost of a tail call.
  - For longer messages, the time grows with message length. Sending 100 numbers takes about twice as long as sending 1.
- My guess is that many of the groups had a "team captain" who handed out the sheets of paper. Each team member would compute their local sum and report it to the team captain. The team captain computed the final sum and reported it (FTW). We can do a bit of front-of-class theatre, you, Devon, me (and if we could get one or two others that would be great). Act it out with slow communication actions, e.g. "tai chi" style. Two problems should become apparent: starting where the team captain has all the data and distributes it is a bottleneck. Having everyone communicate send their result directly to the captain is a bottleneck.

- Design an efficient way to add  $N$  numbers using  $P$  processes.
- Should plan to start with each process having  $\sim N/P$  values – this is the “Keep the data distributed across the parallel processes concept.”
- Should “discover” the tree structure for communication
  - Point to bring out: we are not using a tree to get more parallelism in the final  $P - 1$  additions. These adds don't take long enough to matter. The  $P - 1$  communication actions **do** matter.
  - Reducing the depth of the communication actions from  $P - 1$  (when the team captain handles all of them) to  $\log P$  is what matters.

- Let  $N = 2^K$ .
- We can have one process that creates a binary tree with  $N$  leaves.
- Each leaf process:
  - waits to receive a task with a tag,
  - does the task,
  - sends the result to its parent (with a parent provided tag).
- Each intermediate node:
  - Waits to receive a two functions and a tag,
  - Call the two functions *LeafTask* and *Combine*.
  - The node sends the two functions with a *left* or *right* tag to its children.
  - The node receives results from its children, combines them with *Combine* and sends the result with the parent provided tag to its parent.
  - We now discover that the leaf node probably receives *(LeafTask, Combine, LeftRightTag, Ppid)* just like the intermediate nodes. The leaves just ignore the *Combine*.

- The function for the process is pretty much like the intermediate nodes except
  - There's one function to create the tree.
  - Another function takes a process tree and the *LeafTask* and *CombineTask* functions and sends the result home.

- We want the leaf nodes to generate their arrays as one task, and compute the sums as another task.
- This means that the children need to “remember” state between the two tasks.
- There are a several possible solutions:
  - The *LeafTask* function could take an argument of *ProcState* and return a tuple of *(ValueForCombine, NewProcState)*.
  - The leaf process makes its recursive call with *NewProcState*.
  - Or, we could use the Erlang process dictionary with *erlang:put* and *erlang:get*, but that's not very functional. I prefer the *ProcState* approach.
- And, we need to deal with end-of-life issues.
  - Use the atom *exit* instead of the *(LeafTask, Combine)* tuple.

|                  |        |                        |        |                  |        |                        |        |                  |        |                        |        |                  |        |                        |        |
|------------------|--------|------------------------|--------|------------------|--------|------------------------|--------|------------------|--------|------------------------|--------|------------------|--------|------------------------|--------|
| Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 6 / 14 | Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 7 / 14 | Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 8 / 14 | Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 9 / 14 |
|------------------|--------|------------------------|--------|------------------|--------|------------------------|--------|------------------|--------|------------------------|--------|------------------|--------|------------------------|--------|

We can do better

But we don't need to code the better version in class

Summary

Preview

- Note:** I'm not sure how far we can make it through this material with the various in-class activities. If we make it through the simple implementation of reduce on the previous slide, I'm happy. The rest of this is optional – equivalently, it's material we could move into the Jan. 13 or Jan. 16 lecture.
- With the design above, half of the processes sit around idle, while waiting for the leaves to do their work.
- We can make a tree where each process forwards messages to its right subtree(s) and then does its own *LeafTask*.
- I'm sure I've got a figure from some previous year, I'll find it.

- The better version is implemented in the course Erlang library.
- We now show how to do the reduce example with `wtree:reduce`.

After we make it through reduce, we'll cover

- scan – I want to have one lecture on scan by the end of the Jan. 16 lecture to have the students at a place that they can start on HW2.
- generalized scan and reduce.
- Then, we transition to ~4 lectures on parallel architectures.

|                  |        |                        |         |                  |        |                        |         |                  |        |                        |         |                  |        |                        |         |
|------------------|--------|------------------------|---------|------------------|--------|------------------------|---------|------------------|--------|------------------------|---------|------------------|--------|------------------------|---------|
| Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 10 / 14 | Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 11 / 14 | Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 12 / 14 | Mark Greenstreet | Reduce | CS 418 – Jan. 11, 2017 | 13 / 14 |
|------------------|--------|------------------------|---------|------------------|--------|------------------------|---------|------------------|--------|------------------------|---------|------------------|--------|------------------------|---------|

Review Questions

Scan

Reduce Redux

Reduce Redux

Mark Greenstreet

Problem statement:

Given  $P$  processes that each hold part of an array of numbers, compute the sum of all the numbers in the combined array.

Accumulate step:

Each process computes the total of the elements in its local part of the array.

Outline:

- Reduce Redux**
  - The basic algorithm.
  - Performance model.
  - Implementation considerations.
- Scan**
  - Understand how reduce generalizes to a method that produces all  $N$  values for a “cumulative” operation in  $O(\log N)$  time.
- A few implementation notes**
  - Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license

[1,2,3] [-8,42] [5,12] [4,5,6] [7,-25] [] [1,1,1,1] [96]

6 34 17 15 -18 0 4 96  
+ + + + + + + +  
[1,2,3] [-8,42] [5,12] [4,5,6] [7,-25] [] [1,1,1,1] [96]

|                  |        |                        |        |                  |      |                        |        |                  |      |                        |        |                  |      |                        |        |
|------------------|--------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|
| Mark Greenstreet | Reduce | CS 418 – Jan. 13, 2017 | 1 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 1 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 2 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 3 / 13 |
|------------------|--------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|



Reduce Redux

Reduce Redux

Reduce Notes

Scan Problem Statement

Combine step:  
Each process sends its result to a combiner process. The combiners compute the sums of the values from adjacent pairs of processes.



- For simplicity, I drew the tree as if we used separate processes for accumulating the local arrays and doing the combining.
  - In practice, we use the same processes for both accumulating and combining.
  - Note that 1/2 of the processes are active in the first level of combine; 1/4 of the processes are active in the second level, and so on.
- Simple time model:

$$T \in O\left(\frac{N}{P} + \lambda \log P\right)$$

where  $\lambda$  is **big** – i.e. the communication time.

- Given an array,  $A$ , with  $N$  elements.

Let  $B = \text{scan}_+(A)$ :

$$B_i = \sum_{k=0}^i A_k$$

Example:

$A = [1, 2, 3, -8, 42, 5, 12, 4, 5, 6, 7, -25, 1, 1, 1, 1, 96]$   
 $B = [1, 3, 6, -2, 40, 45, 57, 61, 66, 72, 79, 54, 55, 56, 57, 58, 154]$

- Is there an efficient parallel algorithm for computing  $\text{scan}_+(A)$ ?
  - I wrote *scan*+, because our solution works for any associative operator.

|                  |      |                        |        |                  |      |                        |        |                  |      |                        |        |                  |      |                        |        |
|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|
| Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 2 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 2 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 3 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 4 / 13 |
|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|


Scan Example: Monthly Bank Statement

Brute force Scan


Brute force Scan

Brute force Scan


- Assumptions:
  - You make **lots** of transactions; so, the bank needs to use a parallel algorithm just for your account.
  - Months have 32 days – the power-of-two version of the algorithm is simpler. It generalizes to any number of processors.
  - Each processes has the transaction data for one day.
- Using parallel scan:
  - Each process computes the total of the transactions for its day.
  - Using parallel scan, we determine the balance at each day for each process.
  - The process can use its start-of-day balance, and compute the balance after each transaction for that day.



- Use a reduce tree to compute  $B_7$ .



- Use a reduce tree to compute  $B_7$ .
- Use another reduce tree to compute  $B_6$ .



- Use a reduce tree to compute  $B_7$ .
- Use another reduce tree to compute  $B_6$ .
- Use 6 more reduce trees to compute  $B_{6,0}$ .
- It works. It's  $O(\log P)$  time! But it's not very efficient.


|                  |      |                        |        |                  |      |                        |        |                  |      |                        |        |                  |      |                        |        |
|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|
| Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 5 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 6 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 7 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 8 / 13 |
|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|

Reuse trees


Reuse trees

Scan


Scan: explained




- Key idea: we don't need the trees to be balanced.
- We just want them to be  $O(\log P)$  in height.
- If we need a tree for  $2^k$  nodes, we'll make a balanced tree.
- Otherwise:
  - Make the largest balanced tree we can on the left.
  - Repeat this process for what's left on the right.



- Key idea: we don't need the trees to be balanced.
- We just want them to be  $O(\log P)$  in height.
- If we need a tree for  $2^k$  nodes, we'll make a balanced tree.
- Otherwise:
  - Make the largest balanced tree we can on the left.
  - Repeat this process for what's left on the right.
- Notice that while computing  $B_{10}$ , we produced many other of the  $B$ s as intermediate results.



- See the next slide for an explanation of the notation, etc.



- Notation
  - $A_{-1}$  is initializer for the sum.
  - $A_0, A_1, \dots, A_{15}$  is the initial array.
  - $B_{-1}, B_0, \dots, B_{15}$  is the result of the scan.

$$B_i = \sum_{k=-1}^i A_k \text{ Include the initializer } A_{-1}$$

$\Sigma i : j$  is shorthand for  $\sum_{i=j}^i A_k$

- Each process needs to compute its local part of the scan at the end, starting from the value it receives from the tree.

|                  |      |                        |        |                  |      |                        |        |                  |      |                        |        |                  |      |                        |        |
|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|
| Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 7 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 7 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 8 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 9 / 13 |
|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|------------------|------|------------------------|--------|

A few implementation notes

Reduce & Scan

Preview

Review Questions

- On [slide 3](#) I pointed out that for efficiency, it is better to use the same processes for the leaves and the combine.

```
reduce:
treeLevels = ceil(log2(NProcs));
tally = localAccumulate(...);
for(k = 0; k < treeLevels; k++) {
 if((myPid % (1 << k)) != 0) {
 send(myPid - (1 << k), myPid, tally);
 break;
 } else {
 tally += receive(myPid + (1 << k));
 }
} // Process 0 now has the grand total.
// We can use another loop to broadcast the result.
```

- I'll provide an Erlang version on Monday.

Scan is very similar to reduce. We just change the downward tree.

- For reduce, each process just forwards the grand total to its descendants.
- For scan:
  - Each process records the tallies from its left subtree(s) during the upward sweep.
  - During the downward sweep, each process receives the tally for everything to the left of the subtree for this process.
    - The process adds the tally from its own left subtree to the value from its parent, and sends this to its own right subtree.
    - The process continues the downward sweep for its own left subtree.
    - When we reach a leaf, the process does the final accumulate.

|                                                              |
|--------------------------------------------------------------|
| January 16: Generalized Reduce and Scan                      |
| Homework: Homework 1 deadline for early-bird bonus (11:59pm) |
| Homework 2 goes out (due Feb. 1) – Reduce and Scan           |
| January 18: Reduce and Scan Examples                         |
| Homework: Homework 1 due 11:59pm                             |
| January 20: Architecture Review                              |
| Reading: Pacheco, Chapter 2, through section 2.2             |
| January 23: Shared Memory Architectures                      |
| Reading: Pacheco, Chapter 2, through section 2.3             |
| Homework: Homework 2 deadline for early-bird bonus (11:59pm) |
| Homework 3 goes out (due Feb. 17)                            |
| Homework 2 due 11:59pm                                       |
| January 27-February 6: Parallel Performance                  |
| February 9-17: Parallel Sorting                              |

|                  |      |                        |         |                  |      |                        |         |                  |      |                        |         |                  |      |                        |         |
|------------------|------|------------------------|---------|------------------|------|------------------------|---------|------------------|------|------------------------|---------|------------------|------|------------------------|---------|
| Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 10 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 11 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 12 / 13 | Mark Greenstreet | Scan | CS 418 – Jan. 13, 2017 | 13 / 13 |
|------------------|------|------------------------|---------|------------------|------|------------------------|---------|------------------|------|------------------------|---------|------------------|------|------------------------|---------|

### Generalize Reduce and Scan

Mark Greenstreet

CpSc 418 – Jan. 16, 2017

Outline:

- Reduce in Erlang
- Scan in Erlang

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are hereby released under the terms of the Creative Commons Attribution 4.0 International License <https://creativecommons.org/licenses/by/4.0/>

- Understand relationship between reduce and scan
  - Both are tree walks.
  - The initial combination of values from leaves is identical.
  - Reduce propagates the grand total down the tree.
  - Scan propagates the total "everything to the left" down the tree.
- Generalized Reduce and Scan
  - Understand the role of the *Leaf*, *Combine*, and *Root* functions.
  - Understand the use of higher-order functions to implement reduce and scan.
- The CS418 class library
  - Ability to create a tree of processes.
  - Ability to distribute data and tasks to those processes.
  - Ability to use the *reduce* and *scan* functions from the library.
  - Know where to find more information.

<http://www.ugrad.cs.ubc.ca/~cs418/2016-2/lecture/01-16/code/reduce.erl>

### The Reduce Pattern

- It's a parallel version of *fold*, e.g. `lists:foldl/1`.
- Reduce is described by three functions:
  - Leaf*(/): What to do at the leaves, e.g.

```
fun() -> count3s(Data) end.
```
  - Combine*(/): What to do at the root, e.g.

```
fun(Left, Right) -> Left+Right end.
```
  - Root*(/): What to do with the final result. For count 3s, this is just the identity function.

### The *wtree* module

- Part of the [course Erlang library](#).
- Operations on worker trees

```
wtree:create(NProcs) -> {pid(), Create a list of NProcs processes, organized as a tree. wtree:broadcast(W, Task, Arg) -> ok. Execute the function Task on each process in W. Note: W means "worker pool". wtree:reduce(P, Leaf, Combine, Root) -> term(). A generalized reduce. wtree:reduce(P, Leaf, Combine) -> term(). A generalized reduce where Root defaults to the identity function.
```

### Store Locally

- Communication is expensive – each process should store its own data whenever possible.
- How do we store data in a functional language?
  - Our processes are implemented as Erlang functions that receive messages, process the message, and make a tail-call to be ready to receive the next message.
  - We add a parameter to these functions, *State*, that is a mapping from *Keys* to *Values*.
- What this means when we write code:
  - Functions such as *Leaf* for *wtree:reduce* or *Task* for *wtree:broadcast* have a parameter for *State*.
  - `worker:put(State, Key, Value) -> NewState.` Create a new version of *State* that associates *Value* with *Key*.
  - `worker:get(State, Key, Default) -> Value.` Return the value associated with *Key* in *State*. If no such value is found, *Default* is returned. Note: *Default* can be a function in which case it is called to determine a default value – see the documentation.

### Count3s using *wtree*

```
count3s_par(N, P) -> W = wtree:create(P), wtree:scan(W, list, 'Data'), wtree:reduce(W, fun(ProcState) -> count3s(workers:get(ProcState, 'Data')) end, fun(Left, Right) -> Left+Right end).
```

### Reduce and Scan

- The root node:
  - Reduce: `count3s_reduce(none, [], Total3s) -> Total3s;`
  - Scan: `count3s_scan(none, [], Total3s) -> 0;`
- Internal nodes:
  - Reduce

```
count3s_reduce(Parent, [Child | MoreKids], ThreesInLeftSubtree) -> ThreesInRightSubtree = count3s_wait(Child), ThreesInTree = ThreesInLeftSubtree + ThreesInRightSubtree, Totals = count3s_reduce(Parent, MoreKids, ThreesInTree), count3s_notify(Child, Totals).
```
  - Scan

```
count3s_scan(Parent, [Child | MoreKids], ThreesInLeftSubtree) -> ThreesInRightSubtree = count3s_wait(Child), ThreesInTree = ThreesInLeftSubtree + ThreesInRightSubtree, ThreesOnlyLeft = count3s_scan(Parent, MoreKids, ThreesInTree), count3s_notify(Child, ThreesOnlyLeft + ThreesInLeftSubtree), ThreesOnlyLeft.
```

### Scan in Erlang

- Remarkably like reduce.
- Reduce has
  - an upward pass to compute the grand total
  - a downward pass to broadcast the grand total.
- Scan has
  - an upward pass where the grand total – **just like reduce**
  - On the downward pass, we compute the total of all elements to the left of each subtree.
- Get the code at <http://www.ugrad.cs.ubc.ca/~cs418/2016-2/lecture/01-16/code/scan.erl>

### The Scan Pattern

- It's a parallel version of *mapfold*, e.g. `lists:mapfoldl/1` and `lists:mapfoldr/1`.
- `wtree:scan(Leaf1, Leaf2, Combine, Acc0)`
  - Leaf1*(*ProcState*) -> *Value* Each worker process computes its *Value* based on its *ProcState*.
  - Combine*(*Left*, *Right*) -> *Value* Combine values from sub-trees.
  - Leaf2*(*ProcState*, *AccIn*) -> *ProcState* Each worker updates its state using the *AccIn* value – i.e. the accumulated value of everything to the worker's "left".
  - Acc0*: The value to use for *AccIn* for the leftmost nodes in the tree.

### Scan example: prefix sum

```
prefix_sum_par(W, Key1, Key2) -> wtree:scan(W, fun(ProcState) -> % Leaf1 lists:sum(wtree:get(ProcState, Key1)) end, fun(ProcState, AccIn) -> % Leaf2 wtree:put(ProcState, Key2, prefix_sum(wtree:get(ProcState, Key1), AccIn)) end, fun(Left, Right) -> % Combine Left + Right end, 0 & Acc0).
```

```
prefix_sum(L, Acc0) -> element(1, lists:mapfoldl(fun(X, Y) -> Sum = X+Y, {Sum, Sum} end, Acc0, L)).
```

### More Examples of scan

- Account balance with interest:
  - Input: a list of transactions, where each transaction can be a deposit (add an amount to the balance), a withdrawal (subtract an amount from the balance), or interest (multiply the balance by an amount). For example:

```
{deposit, 100.00}, {withdraw, 5.43}, {withdraw, 27.75}.
```
  - Output: the account balance after each transaction. For example, if we assume a starting balance of \$1000.00 in the previous example, we get

```
[1000.00, 1094.57, 1066.82, 1067.40, ...]
```
- Delete 3s
  - Given a list that is distributed across *NProc* processes, delete all 3s, and rebalance the list so each process has roughly the same length sublists.
  - Solution (sketch):
    - Using scan, each process determines how many 3s precede its segment, the total list length preceding it, and the total list length after deleting 3s.
    - Each process deletes its 3s and send portions of its lists

### Preview

January 18: Reduce and Scan Examples  
Homework: **Homework 1 due 11:59pm**  
January 20: Finish Reduce and Scan  
Mini-assignments: Mini assignment 3 goes out.  
January 23: Architecture Review  
Reading: Pacheco, Chapter 2, through section 2.2  
January 27: Shared Memory Architectures  
Reading: Pacheco, Chapter 2, through section 2.3  
Mini-assignments: Mini assignment 3 due, 10am.  
January 27: Message Passing Architectures  
January 27-February 6: Parallel Performance  
January 30: HW 2 Earlybird due (11:59pm), HW 3 goes out.  
February 1: HW 2 due (11:59pm).  
February 8-17: Parallel Sorting  
February 15: HW 3 Earlybird (11:59pm).  
February 17: HW 3 due (11:59pm).  
February 27: TBD  
March 1: Midterm

### Scan

Mark Greenstreet

CpSc 418 – Jan. 20, 2016

### Objectives

- Prefix sum
  - Spawning processes.
  - Sending and receiving messages.
- The source code for the examples in this lecture is available here: [procs.erl](#).

### Prefix Sum

- Scan is similar to reduce, but every process calculates its cumulative total.
- Example:

```
% prefix_sum: compute prefix sum. prefix_sum(L) when is_list(L) -> prefix_sum_tr(L, 0). prefix_sum_tr([], Acc0) -> []; prefix_sum_tr([E | T], Acc0) -> MySum = B+Acc0, [MySum | prefix_sum_tr(T, MySum)].
```
- Let's try it:

```
> examples:prefix_sum([1, 13, 2, -5, 17, 0, 33]). [1,14,16,11,28,28,61]
```
- How can we do this in parallel?

### Parallel Prefix Sum

### Parallel Prefix Sum

### Parallel Prefix Sum

### Parallel Prefix Sum

### The Scan Pattern

- It's a parallel version of *mapfold*, e.g. `lists:mapfoldl/1` and `lists:mapfoldr/1`.
- `wtree:scan(Leaf1, Leaf2, Combine, Acc0)`
  - Leaf1*(*ProcState*) -> *Value* Each worker process computes its *Value* based on its *ProcState*.
  - Combine*(*Left*, *Right*) -> *Value* Combine values from sub-trees.
  - Leaf2*(*ProcState*, *AccIn*) -> *ProcState* Each worker updates its state using the *AccIn* value – i.e. the accumulated value of everything to the worker's "left".
  - Acc0*: The value to use for *AccIn* for the leftmost nodes in the tree.

### Scan example: prefix sum

```
prefix_sum_par(W, Key1, Key2) -> wtree:scan(W, fun(ProcState) -> % Leaf1 lists:sum(wtree:get(ProcState, Key1)) end, fun(ProcState, AccIn) -> % Leaf2 wtree:put(ProcState, Key2, prefix_sum(wtree:get(ProcState, Key1), AccIn)) end, fun(Left, Right) -> % Combine Left + Right end, 0 & Acc0).
```

```
prefix_sum(L, Acc0) -> element(1, lists:mapfoldl(fun(X, Y) -> Sum = X+Y, {Sum, Sum} end, Acc0, L)).
```

### Prefix Sum Using Scan, example (part 1 of 4)

- Consider the example from [slide 4](#).
- We'll assume that the original lists for each processes are associated with the key *rawdata*.
- We'll store the cumulative sum using the key *cookeddata*.
- Leaf1*: each worker computes the sum of the elements in its list:
  - Worker 0:

```
Leaf1(ProcState) -> lists:sum(wtree:get(ProcState, rawdata)) -> lists:sum([1,3,8]) -> 12.
```
  - Worker 1:

```
Leaf1(ProcState) -> lists:sum([-5,11,2]) -> 8.
```
  - Worker 2:

```
Leaf1(ProcState) -> lists:sum([17,0,-3]) -> 14.
```
  - Workers 3-6: ...
  - Worker 7:

```
Leaf1(ProcState) -> lists:sum([14,15]) -> 29.
```

### Prefix Sum Using Scan, example (part 2 of 4)

- Combine* (upward, first round):
  - Worker 0: `Combine(12, 8) -> 20.`
  - Worker 2: `Combine(14, 104) -> 118.`
  - Worker 4: `Combine(24, -155) -> -131.`
  - Worker 6: `Combine(6, 29) -> 35.`
- Combine* (upward, second round):
  - Worker 0: `Combine(20, 118) -> 138.`
  - Worker 4: `Combine(-131, 35) -> -96.`
- Combine* (upward, final round):
  - Worker 0: `Combine(138, -96) -> 42.`
  - This value is returned to the caller of `wtree:scan`.

### Prefix Sum Using Scan, example (part 3 of 4)

- Combine* (downward)
- The root sends *AccIn*, 0, to the left subtree.
- Each worker that did a combine remembers the arguments from the upward combines, and uses them in the downward sweep. In the code, each upward step is a recursive function call, and each downward step is a return.
- Combine* (downward, first round)
  - Worker 0: `Combine(0, 138) -> 138.`
  - The 0 is *AccIn* from the root.
  - The 138 is the stored value from the left subtree.
  - Worker 0 sends this result to its right subtree, worker 4.
- Combine* (downward, second round)
  - Worker 0: `Combine(0, 20) -> 20.` Send to worker 2.
  - Worker 4: `Combine(138, -131) -> 7.` Send to worker 6.
- Combine* (downward, third round)
  - Worker 0: `Combine(0, 12) -> 12.` Send to worker 1.
  - Worker 2: `Combine(20, 14) -> 34.` Send to worker 3.
  - Worker 4: `Combine(138, 24) -> 162.` Send to worker 5.
  - Worker 6: `Combine(7, 6) -> 13.` Send to worker 7.

### Prefix Sum Using Scan, example (part 4 of 4)

- Leaf2* (update worker state)
  - Worker 0:

```
Leaf2(ProcState) -> wtree:put(ProcState, Key2, prefix_sum(wtree:get(ProcState, Key1), 0)) -> wtree:put(ProcState, Key2, prefix_sum([1, 3, 8], 0)) -> wtree:put(ProcState, Key2, [1, 4, 12]).
```
  - Worker 1:

```
Leaf2(ProcState) -> wtree:put(ProcState, Key2, prefix_sum(wtree:get(ProcState, Key1), 0)) -> wtree:put(ProcState, Key2, prefix_sum([-5, 11, 2], 12)) -> wtree:put(ProcState, Key2, [7, 18, 20]).
```
  - Workers 2-7: ...

### Let's Try It

```
>> W = wtree:create(8). >> [0.65,0,<0.66,0>,<0.67,0>,<0.68,0> <0.69,0>,<0.70,0>,<0.71,0>,<0.72,0>] >> workers:update(W, rawdata, [[1,3,8], [-5,11,2], [17,0,-3], [100,-8,12], [4,15,1], [5,-168,7], [1,2,3], [14,15]]). ok >> examples:prefix_sum_par(W, rawdata, cookeddata). 42 >> workers:retrieve(W, cookeddata). [[1,4,12], [7,18,20], *8**, [134,126,138], [142,161,162], [168,0,7], "\b\n[c", "***** > 537
```

- Likewise, `* = 34, $ = 8, $n = 10, $r = 13, $e = 27, and $* = 42.`
- All is well.

### More Examples of scan

- Account balance with interest:
  - Input: a list of transactions, where each transaction can be a deposit (add an amount to the balance), a withdrawal (subtract an amount from the balance), or interest (multiply the balance by an amount). For example:

```
{deposit, 100.00}, {withdraw, 5.43}, {withdraw, 27.75}.
```
  - Output: the account balance after each transaction. For example, if we assume a starting balance of \$1000.00 in the previous example, we get

```
[1000.00, 1094.57, 1066.82, 1067.40, ...]
```
- Delete 3s
  - Given a list that is distributed across *NProc* processes, delete all 3s, and rebalance the list so each process has roughly the same length sublists.
  - Solution (sketch):
    - Using scan, each process determines how many 3s precede its segment, the total list length preceding it, and the total list length after deleting 3s.
    - Each process deletes its 3s and send portions of its lists and/or receives list portions to rebalance.



## More? Examples of scan

- Carry-Lookahead Addition:
  - Given two large integers as a list of bits (or machine words), compute their sum.
    - Note that the "pencil-and-paper" approach works from the least significant bit (or digit, or machine word) and works sequentially to the most-significant bit. This takes  $O(N)$  time where  $N$  is the number of bits in the work.
  - Carries can be computed using scan.
    - This allows a parallel implementation that adds two integers in  $O(\log N)$  time.
    - This is how the hardware in your CPU does addition – the adder takes  $O(\log N)$  gate delays to add two, machine words, where  $N$  is the number of bits in a word.
- See *Principles of Parallel Programming*, pp. 119f.
- See homework 2 (later today, I hope).

## Preview

|                                             |                                                              |
|---------------------------------------------|--------------------------------------------------------------|
| January 23: Architecture Review             | Reading: Pacheco, Chapter 2, Sections 2.1 and 2.2            |
| January 25: Shared-Memory Machines          | Reading: Pacheco, Chapter 2, Section 2.3                     |
| January 27: Distributed-Memory Machines     | Reading: Pacheco, Chapter 2, Sections 2.4 and 2.5            |
| January 30: Parallel Performance: Speed-up  | Reading: Pacheco, Chapter 2, Section 2.6                     |
| February 1: Parallel Performance: Overheads | Homework: HW 2 early/bird (11:59pm). HW 3 goes out           |
| February 3: Parallel Performance: Models    | Mini Assignments: Mini 3 due (10am)                          |
| February 6: Parallel Performance: Wrap Up   | January 8-February 15: Parallel Sorting                      |
| January 8-February 15: Parallel Sorting     | Homework (Feb. 15): HW 3 early/bird (11:59pm). HW 4 goes out |
| February 17: Map-Reduce                     | Homework: HW 3 due (11:59pm)                                 |
| February 27: TBD                            | March 1: Midterm                                             |

## Review

- What is scan? Give an example.
- Compare scan with `lists:mapfold1`?
- What property must an operator have to be amenable use with scan?
- What are the components of a generalized scan?  
As an example, what functions do you need to define to use `wtree:scan`?
- Consider the following variations on the bank account problem:
  - Add a transaction `[reset, Balance]`, where `Balance` is a number. The account balance is set to this amount. For example, this can be used to open an account with an initial balance. We'll also assume that a `reset` can be done at any point in a sequence of transactions.
  - Change interest computations so that the bank charges a daily interest of  $X\%$  for negative balances, neither charges nor pays interest for positive balances less than \$1000, and pays a daily interest of  $Y\%$  for positive balances greater than \$1000.
  - For each of these:
    - Can the account balance still be computed using scan?
    - If yes, explain how to do. If no, explain why it's not possible.

## Computer Architecture Review

Mark Greenstreet

CpSc 418 – Jan. 23, 2017


- [A microcoded machine](#)
- [A pipelined machine: RISC](#)
- [Let's write some code](#)
- [Superscalars and the memory bottleneck](#)

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

## Objectives

- Review classical, sequential architectures
  - a simple microcoded, machine
  - a pipelined, one-instruction per clock cycle machine
- Pipelining is parallel execution
  - the machine is supposed to appear (nearly) sequential
  - introduce the ideas of hazards and dependencies.

## Microcoded machines



A simple, microcoded machine


- The microcode ( $\mu$ code) ROM specifies the sequence of operations necessary to carry out an instruction.
- For simplicity, I'm assuming that the op-code bits of the instruction form the most significant bits of the  $\mu$ code ROM address, and that the value of the micro-PC ( $\mu$ PC) form the lower half of the address.

## Microcode: summary

- Separates hardware from instruction set.
  - Different hardware can run the same software.
  - Enabled IBM to sell machines with a wide range of performance that were all compatible
    - I.e. IBM built an empire and made a fortune on the IBM 360 and its successors.
    - Intel has done the same with the x86.
- But, as implemented on slide 3, it's very sequential.

```
while(true) {
 fetch an instruction;
 perform the instruction
}
```
- Instruction fetch is "overhead"
  - Motivates coming up with complicated instructions that perform lots of operations per instruction fetch.
  - But these are hard for compilers to use.
  - Can we do better?

## Pipelined instruction execution



A Pipelined (RISC) CPU

- Successive instructions in each stage
  - When instruction `i` in `i` fetch, instruction `i-1` in decode, ...
- Allows throughput of one instruction per cycle.
- Favors simple instructions that execute on a single pass through the pipeline.
  - This is known as RISC: "Reduced Instruction Set Computer"
  - A modern x86 is CISC on the outside, but RISC on the inside.

## What about Dependencies?

- Multiple-instructions are in the pipeline at the same time.
- An instruction starts before all of its predecessors have completed.
- Data hazards occur if
  - an instruction can read a different value than would have been read with a sequential execution of instructions,
  - or if a register or memory location is left holding a different value than it would have had in a sequential execution.
- Control hazards occur if
  - an instruction is executed that would not have been executed in a sequential execution.
  - This is because the instruction "depends" on a jump or branch that hasn't finished in time.

## Handling Hazards

- Bypass: If an instruction has a result that a later instruction needs, the earlier instruction can provide that result directly without waiting to go through the register file.
- Move common operations early:
  - Decide branches in decode stage
  - ALU operations in the stage after decode
  - Memory reads take longer, but they happen less often.
- Let the compiler deal with it
- If nothing else helps, stall.

## Break for Live Coding

## Back to Architecture


- the microcoded machine takes 5+ clock-cycles per instruction.
- the RISC machine takes 1 clock-cycle per instruction – in the best case:
  - There can be stalls due to cache misses,
  - unfilled delay slots, or
  - multi-cycle operations.
- Can we break the one-cycle-per instruction barrier?

## The Memory Bottleneck

- A CPU core can execute roughly one instruction per clock-cycle.
  - With a 3GHz clock, that's roughly 0.3ns per instruction.
- Main memory accesses take 60-200ns (or longer)
  - That's 200-600 instructions per main memory access.
- Why?
  - CPUs designed for speed.
  - Memory designed for capacity:
    - fast memories are small
    - large memories are slow

## Superscalar Processors

A Superscalar CPU



## Superscalar Execution

- Fetch several,  $W$ , instructions each cycle.
- Decode them in parallel, and send them to issue queues for the appropriate functional unit.
- But what about dependencies?
  - We need to make sure that data and control dependencies are properly observed.
  - Code should execute on a superscalar as if it were executing on sequential, one-instruction-at-a-time machine.
  - Data dependencies can be handled by "register renaming" – this uses register indices to dynamically create the dependency graph as the program runs.
  - Control dependencies can be handled by "branch speculation" – guess the branch outcome, and rollback if wrong.
- The opportunity to execute instructions in parallel is called **Instruction Level Parallelism, ILP**.

## What superscalars are good at

- Scientific computing:
  - often successive loop iterations are independent
  - the superscalar **pipelines** the loop
  - Perform memory reads for loop `i`, while doing multiplications for loop `i-2`, while doing additions for loop `i-4`, while storing the results for loop `i-5`.
- Commercial computing (databases, web servers, ...)
  - often have large data sets and high cache miss rates.
  - the superscalar can find executable instructions after a cache miss.
  - if it encounters more misses, the CPU benefits from **pipelined** memory accesses.
- Burning lots of power
  - many operations in a superscalar require hardware that grows quadratically with  $W$ .
  - basically, all instructions in a batch of  $W$  have to compare their register indices with all of the other ones.

## Superscalar Reality

- Most general purpose CPUs (x86, Arm, Power, SPARC) are superscalar.
- Register renaming works **very** well:
- Branch prediction is also very good, often > 90% accuracy.
  - But, data dependent branches can cause very poor performance.
- Superscalar designs make multi-threading possible
  - The features for executing multiple instruction in parallel work well for mixing instructions from several threads or processes – this is called "multithreading" (or "hyperthreading", if you're from Intel).
  - In practice, superscalars are often **better at multithreading** than they are at extracting ILP from a sequential program.

## Preview

|                                             |                                                              |
|---------------------------------------------|--------------------------------------------------------------|
| January 25: Shared-Memory Machines          | Reading: Pacheco, Chapter 2, Section 2.3                     |
| January 27: Distributed-Memory Machines     | Reading: Pacheco, Chapter 2, Sections 2.4 and 2.5            |
| January 30: Parallel Performance: Speed-up  | Reading: Pacheco, Chapter 2, Section 2.6                     |
| February 1: Parallel Performance: Overheads | Homework: HW 2 early/bird (11:59pm). HW 3 goes out           |
| February 3: Parallel Performance: Models    | Mini Assignments: Mini 3 due (10am)                          |
| February 6: Parallel Performance: Wrap Up   | January 8-February 15: Parallel Sorting                      |
| January 8-February 15: Parallel Sorting     | Homework (Feb. 15): HW 3 early/bird (11:59pm). HW 4 goes out |
| February 17: Map-Reduce                     | Homework: HW 3 due (11:59pm)                                 |
| February 27: TBD                            | March 1: Midterm                                             |

## Review

- How does a pipelined architecture execute instruction in parallel?
- What are hazards?
- What are dependencies?
- What is multithreading.
  - For further reading on RISC: "Instruction Sets and Beyond: Computers, Complexity, and Controversy" R.P. Colwell, et al., IEEE Computer, vol. 18, no. 3.
  - You can download the paper for free if your machine is on the UBC network.
  - If you are off-campus, you can use [the library's proxy](#).

## Shared Memory Multiprocessors

Mark Greenstreet

CpSc 418 – Jan. 25, 2017

## Outline:


- Shared-Memory Architectures
- Memory Consistency
- Coding Break
- Weak Consistency

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

## Objectives


- Understand how processors can communicate by sharing memory.
- Able to explain the term "sequential consistency"
  - Describe a simple cache-coherence protocol, MESI
  - Describe how the protocol can be implemented by snooping.
  - Describe "sequential consistency".
  - Be aware that real machines make guarantees that are weaker than sequential consistency.

## An Ancient Shared-Memory Machine



- Multiple CPUs (typically two) shared a memory
- If both attempted a memory read or write at the same time
  - One is chosen to go first.
  - Then the other does its operation.
  - That's the role of the switch in the figure.
- By using multiple memory units (partitioned by address), and a switching network, the memory could keep up with the processors.
- But, now that processors are 100's of times faster than memory, this isn't practical.

## A Shared-Memory Machine with Caches



- Caches reduce the number of main memory reads and writes.
- But, what happens when a processor does a write?


## Cache Inconsistency

- Assume caches are write-back:
  - write-back**: writes only update the cache. Main memory updated when the cache block is evicted.
  - write-through**: writes update cache and main memory.
  - Modern processors have to use write-back for performance: Main memory is way too slow for write-through.
- Step 0: CPU 0 and CPU 1 have both read memory location `addr0` and `addr1` and have copies in their cache.
- Step 1: CPU 0 writes to `addr0` and CPU 1 writes to `addr1`.
- Step 2: CPU 0 reads from `addr1` and CPU 1 reads from `addr0`.
  - Both CPUs see the **old** value.
  - The writers only updated the writer's cache.
  - The readers got the old values.

## Cache Coherence Protocols

- Big idea**: caches communicate with each other so that:
  - Multiple CPUs can have read-only copies for the same memory location.
  - If a cache has a dirty block, then no other cache has a copy of that block.

## The MESI protocol




- Caches can **share read-only** copies of a cache block.
- When a processor writes a cache block, the first write goes to main memory.
  - The other caches are notified and invalidate their copies.
  - This ensures that **writable blocks are exclusive**.

## How caches work

- Caching rhymes with hashing and the two ideas are similar.
  - Caches store data in "blocks" – the block size is a small power-of-two times the machine word size.
  - A cache has one or more "ways" – each way holds a power-of-two number number of blocks.
  - A hash-value is computed from the address.
    - `blockAddr = addr / blockSize; % right shift`
    - `blockIndex = blockAddr % (BlocksPerWay-1); % bit masking`
- Read:
  - The **blockIndex** is used to look up one entry in each "way".
  - Each block has a tag that includes the full-address for the data stored in that block.
  - The tags from each way are compared with the tag of the address:
    - If any tag matches, that way provides the data.
    - If no tags match, then a cache miss occurs.
    - Some current block is evicted from the cache to make room for the incoming block.
- Writes are similar to reads.

## A typical cache



- Only the read-path is shown. Writing is similar.
- This is a 16K-byte, 4-way set-associative cache, with 16 byte cache blocks.

- Caches read and write main memory over a shared memory bus.
- Each cache has two copies of the tags: one for the CPU, the other for the bus.
- If the cache sees another CPU reading or writing a block that is in this cache, it takes the action specified by the MESI protocol.

- Main memory keeps a copy of the data and
  - a bit-vector that records which processors have copies, and
  - a bit to indicate that one processor has a copy and it may be modified.
- A processor accesses main memory as required by the MESI protocol.
  - The memory unit sends messages to the other CPUs to direct them to take actions as needed by the protocol.
  - The ordering of these messages ensures that memory stays consistent.
- Comparison:
  - Snooping is simple for machines with a small number of processors.
  - Directory methods scale better to large numbers of processors.

- Memory is said to be **sequentially consistent** if
- All memory reads and writes from all processors can be arranged into a single, sequential order, such that:
    - The operations for each processor occur in the global ordering in the same order as they did on the processor.
    - Every read gets the value of the preceding write to the same address.
  - Sequential consistency corresponds to what programmers think "ought" to happen.
    - Very similar to "serializability" for database transactions.
  - MESI guarantees sequential consistency

Weak Consistency

- CPUs typically have "write-buffers" because memory writes often come in bursts.
- Typically, reads can move ahead of writes to maximize program performance.
- Why?
  - Because there may be instructions waiting for the data from a load.
  - A transition from "shared" to "modified" requires notifying all processors – this can take a long time.
  - Memory writes don't happen until the instruction commits.
- This means that real computers don't guarantee sequential consistency.
  - **Warning:** classical algorithms for locks and shared buffers fail when run on a real machines!

Programming Shared Memory Machines

- Shared memory make parallel programming "easier" because:
  - One thread can pass an entire data structure to another thread just by giving a pointer.
  - No need to pack-up trees, graphs, or other data structures as messages and unpack them at the receiving end.
- Shared memory make parallel programming **harder** because:
  - It's easy to overlook synchronization (control to shared data structures). Then, we get data races, corrupted data structures, and other hard-to-track-down bugs.
  - A defensive reaction is to wrap every shared reference with a lock. But locks are slow (that's a factor for communication), and this often results in slow code, or even deadlock.
- In practice, shared memory code that works often has a message-passing structure.
- Finally, beware of weak consistency
  - Use a thread library.
  - There are elegant algorithms that avoid locking overhead, even with weak consistency, but they are beyond the scope of this class.

Shared Memory and Performance

- Shared memory can offer better performance than message passing because
  - High bandwidth: the buses that connect the caches can be very wide, especially if the caches are on a single chip.
  - Low latency: the hardware handles moving the data – no operating system calls and context-switch overheads.
- But, shared memory doesn't scale as well as message passing
  - For large machines, the latency of directory accesses can severely degrade performance.
    - In a message passing machine, each CPU has its own memory, nearby and fast.
    - For shared memory, each CPU has part of the shared main memory – accessing a directory may require accessing the memory of a distant CPU.
  - Shared memory moves the data after the cache miss
    - this stalls a thread
    - message passing can send data in advance and avoid these stalls

Summary

- Shared-Memory Architectures
  - Use cache-coherence protocols to allow each processor to have its own cache while maintaining (almost) the appearance of having one shared memory for all processors.
    - A typical protocol: MESI
    - The protocol can be implemented by snooping or directories.
  - Using cache-memory interconnect for interprocessor communication provides:
    - High-bandwidth
    - Low-latency, but watch out for fences, etc.
    - High cost for large scale machines.
- Shared-Memory Programming
  - Need to avoid interference between threads.
  - Assertion reasoning (e.g. invariants) are crucial, much more so than in sequential programming.
  - There are too many possible interleavings to handle intuitively.
  - In practice, we don't formally prove complete programs, but we use the ideas of formal reasoning.
  - Real computers don't provide sequential consistency.
    - Use a thread library.

Preview

|                                             |                                                                                   |                   |
|---------------------------------------------|-----------------------------------------------------------------------------------|-------------------|
| January 27: Distributed-Memory Machines     | Reading: Padhoo, Chapter 2, Sections 2.4 and 2.5. Mini Assignments                | Mini 4 goes out   |
| January 30: Parallel Performance: Speed-up  | Padhoo, Chapter 2, Section 2.6. Homework: HW 2 earlybird (11:59pm). HW 3 goes out |                   |
| February 1: Parallel Performance: Overheads | Homework: HW 2 due (11:59pm).                                                     |                   |
| February 3: Parallel Performance: Models    | Mini Assignments                                                                  | Mini 4 due (10am) |
| February 6: Parallel Performance: Wrap Up   | Homework: HW 3 earlybird (11:59pm). HW 4 goes out                                 |                   |
| January 8-February 15: Parallel Sorting     | Homework (Feb. 15): HW 3 earlybird (11:59pm). HW 4 goes out                       |                   |
| February 17: Map-Reduce                     | Homework: HW 3 due (11:59pm).                                                     |                   |
| February 27: TBD                            |                                                                                   |                   |
| March 1: Midterm                            |                                                                                   |                   |

Review

- What is sequential consistency?
- Using the MESI protocol, can multiple processors simultaneously have entries in their caches for the same memory address?
- Using the MESI protocol, can multiple processors simultaneously modify entries in their caches for the same memory address?
- How can a cache-coherence protocol be implemented by snooping?
- How can a cache-coherence protocol be implemented using directories?
- What is false sharing (in the reading, but not covered in these slides)?
- Do real machines provide sequential consistency?
- How do these issues influence good software design practice?

Classifying Cache Misses

- **Compulsory:** The first reference to a cache block will cause a miss.
  - Note that the first access should be a write – otherwise the location is uninitialized.
  - A cache can avoid stalling the processor by using "allocate on write".
  - If a miss is a write, assign a block for the line, start the main memory read, track which bytes have been written, and merge with the data from memory when it arrives.
- **Capacity:** The cache is not big enough to hold all of the data used by the program.
- **Conflict:** Many active memory locations map to the same cache index.
  - If there are more such references than the associativity of the cache, these will cause conflict misses.
- **Coherence:** A cache block was evicted because another CPU was writing to it.
  - A subsequent read incurs a cache miss.

Cache Design Trade-Offs (1 of 2)

- **Capacity:** Larger caches have lower miss rates, but longer access times. This motivates using multiple levels of caches.
  - L1: closest to the CPU, smallest capacity (16-64Kbytes), fastest access (1-3 clock cycles).
  - L2: typically 128Kbytes to 1Mbyte, 5-10 cycle access time.
  - L3: becoming common, several Mbytes of capacity.
- **Block Size:**
  - Larger blocks can lower miss rate by exploiting spatial locality.
  - Larger blocks can raise miss rate due to conflict and coherence misses.
  - Larger blocks increase miss penalty by requiring more time to transfer all that data.
  - Typical block sizes are 16 to 256 bytes – sometimes block size changes with cache level.

Cache Design Trade-Offs (2 of 2)

- **Associativity:**
  - Increasing associativity generally reduces the number of conflict misses.
  - Increasing associativity makes the cache hardware more complicated.
  - Typical caches are direct mapped to four- or eight-way associative.
  - Associativity doesn't need to be a power of two!
- **Other stuff**
  - cache inclusion: is everything in the L1 also in the L2?
  - interaction with virtual memory: are cache addresses virtual or physical?
  - coherence protocol details:
    - Example, Intel uses MESI: the "F" stands for "forwarding". If a processor has a read miss, and another cache has a copy, one of the caches with a copy will be the "forwarding cache". The forwarding cache provides the data because it's much faster than main memory.
    - error detection and creation – caches + cosmic rays = flipped bits.
    - and all kinds of other optimizations that are beyond the scope of this class.

False Sharing

- False sharing occurs when two CPUs are actively writing different words in the same cache block.
  - Each write forces the other CPU to invalidate its cache block.
  - Each read forces the other CPU to change its cache block from *modified* or *exclusive* to *shared*.
- Example: count 3s
  - Here's an implementation with awful performance.
  - We create a global array of `ints` to hold the accumulators for each process.
  - Each time a process finds a 3, it writes to its element in the array.
  - This forces the other CPUs whose accumulators are in the same block to invalidate their cache entry.
  - This turns accumulator accesses into main memory accesses.
  - And these accesses are serialized: one CPU at a time.

Message Passing Computers

Mark Greenstreet

CpSc 418 – Jan. 27, 2017

Outline:

- [Network Topologies](#)
- [Performance Considerations](#)
- [Examples](#)

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license <http://creativecommons.org/licenses/by/4.0/>

Objectives

- Familiar with typical network topologies: rings, meshes, crossbars, tori, hypercubes, trees, fat-trees.
- Understand implications for programming
  - bandwidth bottlenecks
  - latency considerations
  - location matters
  - heterogeneous computers.

Message Passing Computers

- Multiple CPUs
- Communication through a network:
  - Commodity networks for small clusters.
  - Special high-performance networks for super-computers
- Programming model:
  - Explicit message passing between processes (like Erlang)
  - No shared memory or variables.

Some simple message-passing clusters

- 25 linux workstations (e.g. lin01...lin25.ugrad.cs.ubc.ca) and standard network routers.
  - A good platform for learning to use a message-passing cluster.
  - But, we'll figure out that network bandwidth and latency are key bottlenecks.
- A "blade" based cluster, for example:
  - 16 "blades" each with 4 6-core CPU chips, and 32G of DRAM.
  - An "infiniband" or similar router for about 10-100 times the bandwidth of typical ethernet.
  - The price tag is ~\$300K.
    - Great if you need the computer power.
    - But, we won't be using one in this class.

The Sunway TaihuLight

- The world's fastest (Linpack) super-computer (as of June 2016)
- 40,960 multicore CPUs
  - 256 cores per CPU chip.
  - 1.45GHz clock frequency, 8 flops/core/cycle.
- Total of 10,485,760 cores
- LINPACK performance: 93 PFlops
- Power consumption 15MW (computer) + cooling (unspecified)
- Tree-like
  - Five levels of hierarchy.
  - Each level has a high-bandwidth switch.
  - Some levels (all?) are fully-connected for that level.
- Programming model: A version linux with MPI tuned for this machine.
- For more information, see [Report on the Sunway TaihuLight System](#), J. Dongarra, June 2016.

The Westgrid Clusters

- Clusters at various Western Canadian Universities (including UBC).
- Up to 9600 cores.
- Available for research use.

Network Topologies

- Network topologies are to the message-passing community what cache-coherence protocols are to the shared-memory people:
  - Lots of papers have been published.
  - Machine designers are always looking for better networks.
  - Network topology has a strong impact on performance, the programming model, and the cost of building the machine.
- A message-passing machine may have multiple networks:
  - A general purpose network for sending messages between machines.
  - Dedicated networks for reduce, scan, and synchronization:
    - The reduce and scan networks can include ALUs (integer and/or floating point) to perform common operations such as sums, max, product, all, any, etc. in the networking hardware.
    - A synchronization network only needs to carry a few bits and can be designed to minimize latency.


Ring-Networks

- Advantages: simple.
- Disadvantages:
  - Worst-case latency grows as  $O(P)$  where  $P$  is the number of processors.
  - Easily congested – limited bandwidth.

Star Networks

- Advantages:
  - Low-latency – single hop between any two nodes
  - High-bandwidth – no contention for connections with different sources and destinations.
- Disadvantages:
  - Amount of routing hardware grows as  $O(P^2)$ .
  - Requires lots of wires, and from switch – Imagine trying to build a switch that connects to 1000 nodes!
- Summary
  - Surprisingly practical for 10-50 ports.
  - Hierarchies of cross-bars are often used for larger networks.

A crossbar switch



Meshes

- Advantages:
  - Easy to implement: chips and circuit boards are effectively two-dimensional.
  - Cross-section bandwidth grow with number of processors – more specifically, bandwidth grows as  $\sqrt{P}$ .
- Disadvantages:
  - Worst-case latency grows as  $\sqrt{P}$ .
  - Edges of mesh are "special cases."

Tori

- Advantages:
  - Has the good features of a mesh, and
  - No special cases at the edges.
- Disadvantages:
  - Worst-case latency grows as  $\sqrt{P}$ .

Hypercubes

A 0-dimensional (1 node), radix-2 hypercube

Hypercubes

A 1-dimensional (2 node), radix-2 hypercube




## A 2-dimensional (4 node), radix-2 hypercube




## A 3-dimensional (8 node), radix-2 hypercube



## A 4-dimensional (16 node), radix-2 hypercube




## A 5-dimensional (32 node), radix-2 hypercube



## Hypercubes

A 5-dimensional (32 node), radix-2 hypercube




- Advantages
  - Small diameter ( $\log N$ )
  - Lots of bandwidth
  - Easy to partition.
  - Simple model for algorithm design.
- Disadvantages
  - Needs to be squeezed into a three-dimensional universe.
  - Lots of long wires to connect nodes.
  - Design of a node depends on the size of the machine.

## Dimension Routing


```
% Send a message, msg, from node src to node dst
for i = 1:d
 if (bit(i, src) != bit(i, dst)) % d is dimension of the hypercube
 send(msg, link(i)) % if different for dimension i
 end % then send msg to our i-neighbour
```

## Trees



- Simple network: number of routing nodes = number of processors - 1.
- Wiring:  $O(\log N)$  extra height ( $O(N \log N)$ ) extra area.
  - Wiring:  $O(\sqrt{N \log N})$  extra area for H-tree.
- Low-latency:  $O(\log N)$  + wire delay.
- Low-bandwidth: bottleneck at root.

## Fat-Trees



- Use  $M^2$  parallel links to connect subtrees with  $M$  leaves.
- $0 \leq \alpha \leq 1$ 
  - $\alpha = 0$ : simple tree
  - $\alpha = 1$ : strange crossbar
- Fat-trees are "universal"
  - For  $\frac{2}{3} < \alpha < 1$  a fat-tree interconnect with volume  $V$  can simulate any interconnect that occupies the same volume with a time overhead that is poly-log factor of  $N$ .

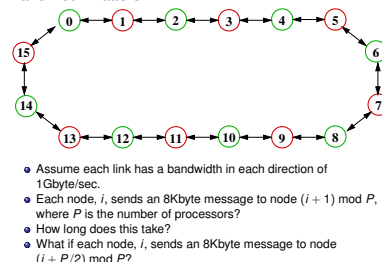
## Performance Considerations

- Bandwidth
  - How many bytes per-second can we send between two processors?
    - May depend on which two processors: neighbours may have faster links than spanning the whole machine.
  - Bisection bandwidth: find the **worst** way to divide the processors into two sets of  $P/2$  processors each.
    - How many bytes per-second can we send between the two partitions?
    - If we divide this by the number of processors, we typically get a much smaller value than the peak between two processors.
- Latency
  - How long does it take to send a message from one processor to another?
    - Typically matters the most for short messages.
    - Round-trip time is often a good way to measure latency.
- Cost
  - How expensive is the interconnect - it may dominate the total machine cost.
    - Cost of the network interface hardware.
    - Cost of the cables.

## Real-life networks

- InfiniBand is becoming increasingly prevalent
- Peak bandwidths  $\geq 6$  GBytes/sec.
  - achieved bandwidths of 2-3GB/s.
- Support for RDMA and "one-sided" communication
  - CPU A can read or write a block of memory residing with CPU B.
- Often, networks include trees for synchronization (e.g. barriers), and common reduce and scan operations.
- The MPI (message-passing interface) evolves to track the capabilities of the hardware.

## Bandwidth Matters



- Assume each link has a bandwidth in each direction of 1 Gbyte/sec.
- Each node,  $i$ , sends an 8kbyte message to node  $(i + 1) \bmod P$ , where  $P$  is the number of processors?
- How long does this take?
- What if each node,  $i$ , sends an 8kbyte message to node  $(i + P/2) \bmod P$ ?

## What this means for programmers

- Location matters.
  - The meaning of location depends on the machine.
  - Getting a good programming model is hard.
  - Challenges of heterogeneous machines.
- What it means for different kinds of computers
  - Supercomputers
  - Clouds
  - PCs of the future(?)

## Summary

- Message passing machines have an architecture that corresponds to the message-passing programming paradigm.
- Message passing machines can range from
  - Clusters of PC's with a commodity switch.
  - Clouds: lots of computers with a general purpose network.
  - Super-computers: lots of compute nodes tightly connected with high-performance interconnect.
- Many network topologies have been proposed:
  - Performance and cost are often dominated by network bandwidth and latency.
  - The network can be more expensive than the CPUs.
  - Peta-flops or other instruction counting measures are an indirect measure of performance.
- Implications for programmers
  - Location matters
  - Communication costs of algorithms is very important
  - Heterogeneous computing is likely in your future.

## Preview

|                                            |                                          |
|--------------------------------------------|------------------------------------------|
| January 30: Parallel Performance: Speed-up | Pacheco, Chapter 2, Section 2.6.         |
| Reading:                                   | HW 2 earlybird (11:59pm). HW 3 goes out. |
| Homework:                                  | HW 2 due (11:59pm).                      |
| February 1: Parallel Performance: Models   | Mini Assignments                         |
| Mini Assignments                           | Mini 4 due (10am)                        |
| February 6: Parallel Performance: Wrap Up  |                                          |
| January 8-February 15: Parallel Sorting    | HW 3 earlybird (11:59pm). HW 4 goes out. |
| Homework (Feb. 15):                        | HW 3 due (11:59pm).                      |
| February 17: Map-Reduce                    |                                          |
| Homework:                                  | HW 3 due (11:59pm).                      |
| February 27: TBD                           |                                          |
| March 17: Midterm                          |                                          |

- There will be readings assigned from [Programming Massively Parallel Processors](#) starting after the midterm.
- Make sure you have a copy. Note: this year, we'll make sure the course works with either the 2<sup>nd</sup> or 3<sup>rd</sup> edition.


## Review

- Consider a machine with 4096 processors.
- What is the maximum latency for sending a message between two processors (measured in network hops) if the network is
  - A ring?
  - A crossbar?
  - A 2-D mesh?
  - A 3-D mesh?
  - A hypercube?
  - A binary tree?
  - A radix-4 tree?

## Supplementary Material


- [Message-passing origami](#): how to fold a mesh into a torus.
- [How big is a hypercube](#): it's all about the wires.

## From a mesh to a torus (1/2)



- Fold left-to-right, and make connections where the left and right edges meet.
- Now, we've got a cylinder.
- Note that there are no "long" horizontal wires: the longest wires jump across one processor.

## From a mesh to a torus (2/2)



- Fold top-to-bottom, and make connections where the top and bottom edges meet.
- Now, we've got a torus.
- Again there are no "long" wires.

## How big is a hypercube?

- Consider a hypercube with  $N = 2^d$  nodes.
- Assume each link can transfer one message in each direction in one time unit. The analysis here easily generalizes for links of higher or lower bandwidths.
- Let each node send a message to each of the other nodes.
- Using dimension routing,
  - Each node will send  $N/2$  messages for each of the  $d$  dimensions. This takes time  $N/2$ .
  - As soon as one batch of messages finishes the dimension-0 route, that batch can continue with the dimension-1 route, and the next batch can start the dimension 0 route.
  - So, we can route with a throughput of  $\left(\frac{N}{2}\right)$  messages per  $N/2$  time.

## How big is a hypercube?

- Consider a hypercube with  $N = 2^d$  nodes.
- Assume each link can transfer one message in each direction in one time unit. The analysis here easily generalizes for links of higher or lower bandwidths.
- Let each node send a message to each of the other nodes.
- Using dimension routing, we can route with a throughput of  $\left(\frac{N}{2}\right)$  messages per  $N/2$  time.
- Consider any plane such that  $N/2$  nodes are on each side of the plane.
  - $\frac{1}{2} \left(\frac{N}{2}\right)$  messages must cross this plane in  $N/2$  time.
  - This means that at least  $N - 1$  links must cross the plane.
  - The plane has area  $O(N)$ .

## How big is a hypercube?

- Consider a hypercube with  $N = 2^d$  nodes.
- Assume each link can transfer one message in each direction in one time unit. The analysis here easily generalizes for links of higher or lower bandwidths.
- Let each node send a message to each of the other nodes.
- Using dimension routing, we can route with a throughput of  $\left(\frac{N}{2}\right)$  messages per  $N/2$  time.
- Consider any plane such that  $N/2$  nodes are on each side of the plane.
  - The plane has area  $O(N)$ .
- Because the argument applies for *any* plane, we conclude that the hypercube has diameter  $O(\sqrt{N})$  and thus volume  $O(N^3)$ .
- Asymptotically, the hypercube is all wire.

## Speed-Up

Mark Greenstreet

CpSc 418 - Jan. 30, 2017

## Outline:

- [Measuring Performance](#)
- [Speed-Up](#)
- [Amdahl's Law](#)
- [The law of modest returns](#)
- [Superlinear speed-up](#)
- [Embarrassingly parallel problems](#)

## But first, USRA

## Summer Undergraduate Research Opportunities

- Natural Sciences and Engineering Research Council (NSERC) Undergraduate Student Research Awards (USRAs)
  - Same process to apply for Science Undergraduate Research Experience (SURE) and Work Learn International Undergraduate Research Awards
- See what academic research really looks like
- Many research areas:
  - Google "robotics usra" for full list of projects seeking students
- I have several project proposals:
  - Collaborative control of smart wheelchairs for older adults
  - Numerical software for demonstrating correctness of robots and cyber-physical systems
  - 15 weeks, flexible schedule
  - You get paid!
  - Email potential sponsor ASAP (full applications due by Feb 10)

January 2017

Ian M. Mitchell - USC Computer Science

1

## Objectives

- Understand key measures of performance
  - Time: [latency vs. throughput](#)
  - Time: [wall-clock vs. operation count](#)
  - Speed-up: [slide 5](#)
- Understand common observations about parallel performance
  - Amdahl's law: limitations on parallel performance (and how to evade them)
  - The law of modest returns: high complexity problems are bad, and worse on a parallel machine.
  - Superlinear speed-up: more CPUs  $\Rightarrow$  more, fast memory - and sometimes you win.
  - Embarrassingly parallel problems: sometimes you win, without even trying.



Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International License (<http://creativecommons.org/licenses/by/4.0/>)

## Measuring Performance

- The main motivation for parallel programming is performance
  - Time: make a program run faster.
  - Space: allow a program to run with more memory.
- To make a program run faster, we need to know how fast it is running.
- There are many possible measures:
  - Latency: time from starting a task until it completes.
  - Throughput: the rate at which tasks are completed.
  - Key observation:

$$\text{throughput} = \frac{1}{\text{latency}} \quad \text{sequential programming}$$

$$\text{throughput} \geq \frac{1}{\text{latency}} \quad \text{parallel programming}$$

## Speed-Up

- Simple definition:

$$\text{speed\_up} = \frac{\text{time}(\text{sequential\_execution})}{\text{time}(\text{parallel\_execution})}$$

- We can also describe speed-up as how many percent faster:

$$\% \text{faster} = (\text{speed\_up} - 1) \times 100\%$$

- But beware of the spin:
  - Is "time" latency or throughput?
  - How big is the problem?
  - What is the sequential version:
    - The parallel code run on one processor?
    - The fastest possible sequential implementation?
    - Something else?
- More practically, how do we measure time?

## Speed-Up - Example


- Let's say that count 3s of a million items takes 10ms on a single processor.
- If I run count 3s with four processes on a four CPU machine, and it takes 3.2ms, what is the speed-up?
- If I run count 3s with 16 processes on a four CPU machine, and it takes 1.8ms, what is the speed-up?
- If I run count 3s with 128 processes on a 32 CPU machine, and it takes 0.28ms, what is the speed-up?

## Time complexity

- What is the time complexity of sorting?
  - What are you counting?
  - Why do you care?
- What is the time complexity of matrix multiplication?
  - What are you counting?
  - Why do you care?

- In our algorithms classes, we count "operations" because we have some belief that they have something to do with how long the actual program will take to execute.
  - Or maybe not. Some would argue that we count "operations" because it allows us to use nifty techniques from discrete math.
  - I'll take the position that the discrete math is nifty *because* it tells us something useful about what our software will do.
- In our architecture classes, we got the formula:
$$\text{time} = \frac{(\# \text{inst. executed}) \times (\text{cycles/instruction})}{\text{clock frequency}}$$
- The approach in algorithms class of counting comparisons or multiplications, etc., is based on the idea that everything else is done in proportion to these operations.
- BUT**, in parallel programming, we can find that a communication between processes can take 1000 times longer than a comparison or multiplication.
  - This may not matter if you're willing to ignore "constant factors."
  - In practice, factors of 1000 are too big to ignore.

- Given a sequential program where
  - fraction  $s$  of the execution time is inherently sequential.
  - fraction  $1 - s$  of the execution time benefits perfectly from speed-up.
- The run-time on  $P$  processors is:
$$T_{\text{parallel}} = T_{\text{sequential}} * (s + \frac{1-s}{P})$$
- Consequences:
  - Define
$$\text{speed\_up} = \frac{T_{\text{sequential}}}{T_{\text{parallel}}}$$
  - Speed-up on  $P$  processors is at most  $\frac{1}{s}$ .
  - Gene Amdahl argued in 1967 that this limit means that parallel computers are only useful for a few special applications where  $s$  is very small.



Amdahl's law is not a physical law.

- Amdahl's law is mathematical theorem:
  - If  $T_{\text{parallel}} = (s + \frac{1-s}{P}) T_{\text{sequential}}$
  - and  $\text{speed\_up} = T_{\text{sequential}} / T_{\text{parallel}}$
  - then for  $0 < s \leq 1$ ,  $\text{speed\_up} \leq \frac{1}{s}$ .
- Amdahl's law is also an **economic** law:
  - Amdahl's law was formulated when CPUs were expensive.
  - Today, CPUs are cheap
    - The cost of fabricating eight cores on a die is very little more than the cost of fabricating one.
    - Computer cost is dominated by the rest of the system: memory, disk, network, monitor, ...
- Amdahl's law assumes a fixed problem size.


## Amdahl's Law, 49 years later

- Amdahl's law is an **economic** law, not a **physical** law.
  - Amdahl's law was formulated when CPUs were expensive.
  - Today, CPUs are cheap (see previous slide)
- Amdahl's law assumes a fixed problem size
  - Many computations have  $s$  (sequential fraction) that decreases as  $N$  (problem size) increases.
  - Having lots of cheap CPUs available will
    - Change our ideas of what computations are easy and which are hard.
    - Determine what the "killer-apps" will be in the next ten years.
    - Ten years from now, people will just take it for granted that most new computer applications will be parallel.
- Examples: see next slide

## Amdahl's Law, 49 years later

- Amdahl's law is an **economic** law, not a **physical** law.
- Amdahl's law assumes a fixed problem size
  - Ten years from now, people will just take it for granted that most new computer applications will be parallel.
  - Examples:
    - Managing/searching/mining massive data sets.
    - Scientific computation.
      - Note that most of the computation for animation and rendering resembles scientific computation. Computer games benefit tremendously from parallelism.
      - Likewise for multimedia computing.

## Amdahl's Law, one more try



- We can have problems where the parallel work grows faster than the sequential part.
- Example: parallel work grows as  $N^{3/2}$  and the sequential part grows as  $\log P$ .

## The Law of Modest Returns

More bad news. ☹

- Let's say we have an algorithm with a sequential run-time  $T = (12s)N^4$ .
  - If we're willing to wait for one hour for it to run, what's the largest value of  $N$  we can use?
  - If we have 10000 machines, and perfect speed-up (i.e.  $\text{speed\_up} = 10000$ ), now what is the largest value of  $N$  we can use?
  - What if the run-time is  $(5ns)1.2^{N^2}$ ?
- The law of modest returns
  - Parallelism offers modest returns, unless the problem is of fairly low complexity.
  - Sometimes, modest returns are good enough: weather forecasting, climate models.
  - Sometimes, problems have huge  $N$  and low complexity: data mining, graphics, machine learning.

## Super-Linear Speed-up

Sometimes,  $\text{speed\_up} > P$ . ☹

- How does this happen?
  - Impossibility "proof": just simulate the  $P$  parallel processors with one processor, time-sharing  $P$  ways.
- Memory: a common explanation
  - $P$  machines have more main memory (DRAM)
  - and more cache memory and registers (total)
  - and more I/O bandwidth, ...
- Multi-threading: another common explanation
  - The sequential algorithm underutilizes the parallel capabilities of the CPU.
  - A parallel algorithm can make better use.
- Algorithmic advantages: once in a while, you win!
  - Simulation as described above has overhead.
  - If the problem is naturally parallel, the parallel version can be more efficient.
- BUT**: be very skeptical of super-linear claims, especially if  $\text{speed\_up} \gg P$ .

## Embarrassingly Parallel Problems

Problems that can be solved by a large number of processors with very little communication or coordination.

- Rendering images for computer-animation: each frame is independent of all the others.
- Brute-force searches for cryptography.
- Analyzing large collections of images: astronomy surveys, facial recognition.
- Monte-Carlo simulations: same model, run with different random values.
- Don't be ashamed if your code is embarrassingly parallel:**
  - Embarrassingly parallel problems are great: you can get excellent performance without heroic efforts.
  - The only thing to be embarrassed about is if you *don't* take advantage of easy parallelism when it's available.

## Lecture Summary

Parallel Performance

- Speed-up: [slide 5](#)
- Limits
  - Amdahl's law, [slide 9](#).
  - Modest gains, [slide 15](#).
- Sometimes, we win
  - Super-linear speedup, [slide 16](#).
  - Embarrassingly Parallel Problems, [slide 17](#).

## Preview

- February 1: Parallel Performance: Overheads
  - Homework: HW 2 due (11:59pm).
- February 3: Parallel Performance: Models
  - Mini Assignments: Mini 4 due (10am)
- February 6: Parallel Performance: Wrap Up
- February 6: Parallel Sorting – The Zero-One Principle
  - Homework (Feb. 15): HW 3 earlybird (11:59pm), HW 4 goes out.
- February 10: Bitonic Sorting (part 1)
  - February 15: Family Day – no class
- February 13: Bitonic Sorting (part 2)
  - Homework (Feb. 15): HW 3 earlybird (11:59pm), HW 4 goes out.
- February 17: Map-Reduce
  - Homework: HW 3 due (11:59pm).
- February 27: TBD
- March 1: Midterm
- Reading from "Programming Massively Parallel Computers" (D.B. Kirk & W.-M. Hwu) start right after the midterm. Make sure you have a copy.
- You can use either the 2<sup>nd</sup> or 3<sup>rd</sup> edition.

## Review Questions

- What is speed-up? Give an intuitive, English answer [and](#) a mathematical formula.
- Why can it be difficult to determine the sequential time for a program when measuring speed-up?
- What is Amdahl's law? Give a mathematical formula. Why is Amdahl's law a concern when developing parallel applications? Why in many cases is it not a show-stopper?
- Is parallelism an effective solution to problems with high big-O complexity? Why or why not?
- What is super-linear speed-up? Describe two causes.
- What is an embarrassingly parallel problem. Give an example.

## Performance-Loss

Mark Greenstreet

CpSc 418 – Feb. 1, 2017

Outline:

- Overhead**: work the parallel code has to do that the sequential version avoids.
  - Communication and Synchronization
  - Extra computation, extra memory
- Limited parallelism**
  - Code that is inherently sequential or has limited parallelism
  - Idle processors
  - Resource contention

## Objectives

- Learn about main causes of performance loss:
  - Overhead
    - Non-parallelizable code
    - Idle processors
    - Resource contention
- See how these arise in message-passing, and shared-memory code.

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International License <http://creativecommons.org/licenses/by/4.0/>

## Causes of Performance Loss


- Ideally, we would like a parallel program to run  $P$  times faster than the sequential version when run on  $P$  processors.
- In practice, this rarely happens because of:
  - Overhead**: work that the parallel program has to do that isn't needed in the sequential program.
  - Non-parallelizable code**: something that has to be done sequentially.
  - Idle processors**: There's work to do, but some processor are waiting for something before they can work on it.
  - Resource contention**: Too many processors overloading a limited resource.

## Overhead

Overhead: work that the parallel program has to do that isn't needed in the sequential program.

- Communication**:
  - The processes (or threads) of a parallel program need to communicate.
  - A sequential program has no interprocess communication.
- Synchronization**:
  - The processes (or threads) of a parallel program need to coordinate.
  - This can be to avoid interference, or to ensure that a result is ready before it's used, etc.
  - Sequential programs have a completely specified order of execution: no synchronization needed.
- Computation**.
  - Recomputing a result is often cheaper than sending it.
- Memory Overhead**.
  - Each process may have its own copy of a data structure.

## Communication Overhead



- In a parallel program, data must be sent between processors.
- This isn't a part of the sequential program.
- The time to send and receive data is overhead.
- Communication overhead occurs with both shared-memory and message passing machines and programs.
- Example: Reduce (e.g. Count 3s):
  - Communication between processes adds time to execution.
  - The sequential program doesn't have this overhead.

## Communication with shared-memory

- In a shared memory architecture:
  - Each core has its own cache.
  - The caches communicate to make sure that all references from different cores to the same address look like there is one, common memory.
  - It takes longer to access data from a remote cache than from the local cache. This creates overhead.
- False sharing** can create communication overhead even when there is no logical sharing of data.
  - This occurs if two processors repeatedly modify different locations on the same cache line.

## Communication overhead: example

- The *Principles of Parallel Programming* book considered an example of Count 3s (in C, with threads), where there was a global array, `int count[P]` where  $P$  is the number of threads.
  - Each thread (e.g. thread  $i$ ) initially sets its count, `count[i]` to 0.
  - Each time a thread encounters a 3, it increments its element in the array.
- The parallel version ran much slower than the sequential one.
  - Cache lines are much bigger than a single `int`. Thus, many entries for the `count` array are on the same cache line.
  - A processor has to get exclusive access to update the count for its thread.
  - This invalidates the copies held by the other processors.
  - This produces lots of cache misses and a slow execution.
- A better solution:
  - Each thread has a local variable for its count.
  - Each thread counts its threes using this local variable and copies its final total to the entry in the global array.

## Communication overhead with message passing

- The time to transmit the message through the network.
- There is also a CPU overhead: the time set up the transmission and the time to receive the message.
- The context switches between the parallel application and the operating system adds even more time.
- Note that many of these overheads can be reduced if the sender and receiver are different threads of the same process running on the same CPU.
  - This has led to SMP implementations of Erlang, MPI, and other message passing parallel programming frameworks.
  - The overheads for message passing on an SMP can be very close to those of a program that explicitly uses shared memory.
  - This allows the programmer to have one parallel programming model for both threads on a multi-core processor and for multiple processes on different machines in a cluster.

## Synchronization Overhead

- Parallel processes must coordinate their operations.
  - Example: access to shared data structures.
  - Example: writing to a file.
- For shared-memory programs (e.g. **pthread**s or **Java threads**), there are explicit locks or other synchronization mechanisms.
- For message passing (e.g. **Erlang** or **MP1**), synchronization is accomplished by **combinators**.

## Computation Overhead

A parallel program may perform computation that is not done by the sequential program.

- Redundant computation: it's faster to recompute the same thing on each processor than to broadcast.
- Algorithm: sometimes the fastest parallel algorithm is fundamentally different than the fastest sequential one, and the parallel one performs more operations.

## Sieve of Eratosthenes

To find all primes  $\leq N$ :

```
1. Let MightBePrime = {2, 3, ..., N}.
2. Let KnownPrimes = [].
3. while(MightBePrime != []) do
 % Loop invariant: KnownPrimes contains all primes less than the
 % smallest element of MightBePrime, and MightBePrime
 % is ascending order. This ensure that the first element of
 % MightBePrime is prime.
 3.1. Let p = first element of MightBePrime.
 3.2. Append p to KnownPrimes.
 3.3. Delete all multiples of p from MightBePrime.
4. end
```

See [http://en.wikipedia.org/wiki/Sieve\\_of\\_Eratosthenes](http://en.wikipedia.org/wiki/Sieve_of_Eratosthenes)

## Prime-Sieve in Erlang

```
% primes(N): return a list of all primes ≤ N.
primes(N) when is_integer(N) and (N < 2) -> [];
primes(N) when is_integer(N) ->
 do_primes([], lists:seq(2, N)).

% invariants of do_primes(Known, Maybe):
% All elements of Known are prime.
% No element of Maybe is divisible by any element of Known.
% lists:reverse(Known) ++ Maybe is an ascending list.
% Known ++ Maybe contains all primes ≤ N, where N is from p(N).
do_primes(KnownPrimes, []) -> lists:reverse(KnownPrimes);
do_primes(KnownPrimes, [P | Etc]) ->
 do_primes([P | KnownPrimes],
 lists:filter(fun(E) -> (E rem P) /= 0 end, Etc)).
```

## A More Efficient Sieve

- If  $N$  is composite, then it has at least one prime factor that is at most  $\sqrt{N}$ .
- This means that once we've found a prime that is  $\geq \sqrt{N}$ , all remaining elements of `Maybe` must be prime.
- Revised code:

```
% primes(N): return a list of all primes ≤ N.
primes(N) when is_integer(N) and (N < 2) -> [];
primes(N) when is_integer(N) ->
 do_primes([], lists:seq(2, N), trunc(math:sqrt(N))).

do_primes(KnownPrimes, [P | Etc], RootN)
 when (P == RootN) ->
 do_primes([P | KnownPrimes],
 lists:filter(fun(E) -> (E rem P) /= 0 end, Etc), RootN);
do_primes(KnownPrimes, Maybe, _RootN) ->
 lists:reverse(KnownPrimes, Maybe).
```

## Prime-Sieve: Parallel Version

- Main idea
  - Find primes from  $1 \dots \sqrt{N}$ .
  - Divide  $\sqrt{N} + 1 \dots N$  evenly between processors.
  - Have each processor find primes in its interval.
- We can speed up this program by having each processor compute the primes from  $1 \dots \sqrt{N}$ .
  - Why does doing extra computation make the code faster?

## Memory Overhead

The total memory needed for  $P$  processes may be greater than that needed by one process due to replicated data structures and code.

- Example: the parallel sieve: each process had its own copy of the first  $\sqrt{N}$  primes.

**Overhead:** Summary  
Overhead is loss of performance due to extra work that the parallel program does that is not performed by the sequential version. This includes:

- **Communication:** parallel processes need to exchange data. A sequential program only has one process; so it doesn't have this overhead.
- **Synchronization:** Parallel processes may need to synchronize to guarantee that some operations (e.g. file writes) are performed in a particular order. For a sequential program, this ordering is provided by the program itself.
- **Extra Computation:**
  - Sometimes it is more efficient to repeat a computation in several different processes to avoid communication overhead.
  - Sometimes the best parallel algorithm is a different algorithm than the sequential version and the parallel one performs more operations.
- **Extra Memory:** Data structures may be replicated in several different processes.

Limited Parallelism

Sometimes, we can't keep all of the processors busy doing useful work.

- **Non-parallelizable code**
  - The dependency graph for operations is narrow and deep.
- **Idle processors**
  - There is work to do, but it hasn't been assigned to an idle processor.
- **Resource contention**
  - Several processes need exclusive access to the same resource.

Unshared State

- Finding the length of a linked list:

```
int length=0;
for(List p = listHead; p != null; p = p->next)
 length++;
```

- Must dereference each `p->next` before it can dereference the next one.
- Could make more parallel by using a different data structure to represent lists (some kind of skiplist, or tree, etc.)
- Searching a binary tree
  - Requires  $2^k$  processes to get factor of  $k$  speed-up.
  - Not practical in most cases.
  - Again, could consider using another data structure.
- Interpreting a sequential program.
- Finite state machines.

Idle Processors

- There is work to do, but processors are idle.
- Start-up and completion costs.
- Work imbalance.
- Communication delays.

Resource Contention

- Processors waiting for a limited resource.
- It's easy to change a compute-bound task into an I/O bound one by using parallel programming.
- Or, we run-into memory bandwidth limitations:
  - Processing cache-misses.
  - Communication between CPUs and co-processors.
- Network bandwidth.

Lecture Summary

Causes of Performance Loss in Parallel Programs

- **Overhead**
  - Communication, [slide 5](#).
  - Synchronization, [slide 9](#).
  - Computation, [slide 10](#).
  - Extra Memory, [slide 15](#).
- Other sources of performance loss
  - Non-parallelizable code, [slide 18](#)
  - Idle Processors, [slide 19](#).
  - Resource Contention, [slide 20](#).

Review Questions

- What is overhead? Give several examples of how a parallel program may need to do more work or use more memory than a sequential program.
- Do programs running on a shared-memory computer have communication overhead? Why or why not?
- Do message passing programs have synchronization overhead? Why or why not?
- Why might a parallel program have idle processes even when there is work to be done?

Mark Greenstreet

CpSc 418 – Feb. 6, 2017

- The RAM Model of Sequential Computation
- Models of Parallel Computation
- An entertaining proof

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

Objectives

- Learn about models of computation
  - Sequential: **Random Access Machine (PRAM)**
  - Parallel
    - **Parallel Random Access Machine (PRAM)**
    - **Candidate Type Architecture (CTA)**
    - **Latency-Overhead-Bandwidth-Processors (LogP)**
- **An entertaining algorithm and its analysis**
  - If a model has invalid assumptions, then we can show that algorithm 1 is faster than algorithm 2, but in real life algorithm 2 is faster.
  - Valiant's algorithm also provides some mathematical entertainment.

The RAM Model

RAM = Random Access Machine

- Axioms of the model
  - Machines work on words of a "reasonable" size.
  - A machine can perform a "reasonable" operation on a word as a single step.
    - such operations include addition, subtraction, multiplication, division, comparisons, bitwise logical operations, bitwise shifts and rotates.
- The machine has an unbounded amount of memory.
  - A memory address is a "word" as described above.
  - Reading or writing a word of memory can be done in a single step.

The Relevance of the RAM Model

- If a single step of a RAM corresponds to (within a factor close to 1) to a single step of a real machine.
- Then algorithms that are efficient on a RAM will also be efficient on a real machine.
- Historically, this assumption has held up pretty well.
  - For example, **mergesort** and **quicksort** are better than **bubblesort** on a RAM and on real machines, and the RAM model predicts the advantage quite accurately.
  - Likewise, for many other algorithms
    - graph algorithms, matrix computations, dynamic programming, ....
    - hard on a RAM generally means hard on a real machine as well: NP complete problems, undecidable problems, ....

The Irrelevance of the RAM Model

The RAM model is based on assumptions that don't correspond to physical reality:

- Memory access time is highly non-uniform.
  - Architects make heroic efforts to preserve the illusion of uniform access time fast memory –
    - caches, out-of-order execution, prefetching, ...
  - – but the illusion is getting harder and harder to maintain.
  - Algorithms that randomly access large data sets run **much** slower than more localized algorithms.
  - Growing memory size and processor speeds means that more and more algorithms have performance that is sensitive to the memory hierarchy.
- The RAM model does not account for energy:
  - Energy is the critical factor in determining the performance of a computation.
  - The energy to perform an operation drops rapidly with the amount of time allowed to perform the operation.

The PRAM Model

PRAM = **Parallel** Random Access Machine

- Axioms of the model
  - A computer is composed of multiple processors and a shared memory.
  - The processors are like those from the RAM model.
    - The processors operate in lockstep.
    - I.e. for each  $k > 0$ , all processors perform their  $k^{\text{th}}$  step at the same time.
  - The memory allows each processor to perform a read or write in a single step.
    - Multiple reads and writes can be performed in the same cycle.
    - If each processor accesses a different word, the model is simple.
    - If two or more processors try to access the same word on the same step, then we get a bunch of possible models:
      - **EREW:** Exclusive-Read, Exclusive-Write
      - **CREW:** Concurrent-Read, Exclusive-Write
      - **CRCW:** Concurrent-Read, Concurrent-Write
    - See [slide 25](#) for more details.

The Irrelevance of the PRAM Model

The PRAM model is based on assumptions that don't correspond to physical reality:

- Connecting  $N$  processors with memory requires a switching network.
  - Logic gates have bounded fan-in and fan-out.
  - $\Rightarrow$  any switch fabric with  $N$  inputs (and/or  $N$  outputs) must have depth of at least  $\log N$ .
- This gives a lower bound on memory access time of  $\Omega(\log N)$ .
- Processors exist in physical space
  - $N$  processors take up  $\Omega(N)$  volume.
  - The processor has a diameter of  $\Omega(N^{1/3})$ .
  - Signals travel at a speed of at most  $c$  (the speed of light).
  - This gives a lower bound on memory access time of  $\Omega(N^{1/3})$ .

The CTA Model

CTA = Candidate Type Architecture

- Axioms of the model
  - A computer is composed of multiple processors.
  - Each processor has
    - Local memory that can be accessed in a single processor step (like the RAM model).
    - A small number of connections to a communications network.
  - There is a communication network connecting the processors.
    - The general model:
      - The communication network is a graph where all vertices (processors and switches) have bounded degree.
      - Each edge has an associated bandwidth and latency.
    - The simplified model:
      - Global actions have a cost of  $\lambda$  times the cost of local actions.
      - $\lambda$  is assumed to be "large".
      - The exact communication mechanism is not specified.

The (Ir)Relevance of the CTA Model

- Recognizing that communication is expensive is the one, most important point to grasp to understand parallel performance.
  - CTA highlights the central role of communication.
  - PRAM ignores it.
- The general model is parameterized by the communication network
  - Can we apply results from analysing a machine with a 3-D toroidal mesh to a machine with fat trees?
  - PRAM ignores it.
- The simple model neglects bandwidth issues
  - Messages are assumed to be "small".
  - But, bigger messages often lead to better performance.
  - If we talk about bandwidth, do we mean the bandwidth of each link?
  - Or, do we mean the bisection bandwidth?

The LogP Model

- **Motivation (1993): convergence of parallel architectures**

- Individual nodes have microprocessors and memory of a workstation or PC.
- A large parallel machine had at most 2000 such nodes.
- Point-to-point interconnect –
  - Network bandwidth much lower than memory bandwidth.
  - Network latency much higher than memory latency.
  - Relatively small network diameter: 5 to 20 "hops" for a 1000 node machine.

- **The model parameters:**

- $L$ : the latency of the communication network fabric
- $o$ : the overhead of a communication action
- $g$ : the bandwidth of the communication network
- $P$ : the number of processors

Why does **g** stand for "bandwidth"?

Marketing!

- What if we used **b** for "bandwidth"?
- Need a catchy acronym with 'l', 'o', 'b', and 'p' ...
  - got it: **BLOP**
  - but the marketing department vetoed it.


logP in practice

- The authors got some surprisingly good performance prediction for a few machines and a few algorithms by finding the "right" values for  $l$ ,  $o$ ,  $g$ , and  $P$  for each architecture.
- It's rare to get a model that comes to within 10-20% on several examples. So, this looked very promising.
- Since then, logP seems to be a model with more parameters than simplified CTA, but not particularly better accuracy.
- Good to know about, because if you meet an algorithms expert, they'll probably know that PRAM is unrealistic.
  - Then, you'll often hear "What about logP?" – the paper has lots of citations.
  - In practice, it's a slightly fancier way of saying "communication costs matter".

Fun with the PRAM Model

Finding the maximum element of an array of  $N$  elements.

- The obvious approach
  - Do a reduce.
  - Use  $N/2$  processors to compute the result in  $\Theta(\log_2 N)$  time.



A Valiant Solution

L. Valiant, 1975

- Use  $N$  processors.
- The big picture:
  - Initially, we can use clumps of three processors to find the largest of three elements in  $O(1)$  time – just do all three comparisons.
  - Now, we have  $N/3$  elements but we still have  $N$  processors. We can perform all of the comparisons for larger clusters of elements in  $O(1)$  time in a single step because we have more processors per element.
  - Valiant showed that the size of a cluster for which we can do all of the pairwise comparisons in a single step grows as  $2^k$  where  $k$  is the number of steps.
  - This leads to a  $\log \log N$  time bound for finding the max.
- I'll sketch the proof.
- Then we'll look at why this shows that you can't actually build a PRAM.


Valiant's algorithm, step 1

- Step 1:
  - Divide the  $N$  elements into  $N/3$  sets of size 3.
  - Assign 3 processors to each set, and perform all three pairwise comparisons in parallel.
  - Mark all the "losers" (requires a CRCW PRAM) and move the max of each set of three to a fixed location.
- The PRAM operations in a bit more detail.
  - Initially, every element has a flag set to 1 that says "might be the max".
  - When  $\binom{k}{2}$  processors perform all of the pairwise comparisons of  $k$  values.
    - Each processor sets the flag for the smaller value to 0.
    - Note that several processors may write 0 to the same location, but the CRCW allows this because they are all writing the same value.
  - One processor for each value checks if its flag is still set to 1.
    - The flag for that location is set to 1.
    - And now we're ready for subsequent rounds.

Valiant's algorithm, step 2

- We now have  $N/3$  elements left and still have  $N$  processors.
- We can make groups of 7 elements, and have 21 processors per group, which is enough to perform all  $\binom{7}{2} = 21$  pairwise comparisons in a single step.
- Thus, in  $O(1)$  time we move the max of each set to a fixed location. We now have  $N/21$  elements left to consider.

Visualizing Valiant



Valiant's Algorithm, the remaining steps

- On step  $k$ , we have  $N/m_k$  elements left.
- On step  $m_k$  is the "sparsity" of the problem – i.e. the number of processors per remaining element.
- We can make groups of  $2m_k + 1$  elements, and have

$$m_k(2m_k + 1) = \frac{(2m_k + 1)(2m_k + 1 - 1)}{2} = \binom{2m_k + 1}{2}$$

- processors per group, which is enough to perform all pairwise comparisons in a single step.
- We now have  $N/(m_k(2m_k + 1))$  elements to consider.
- Therefore,  $m_{k+1} = 2m_k^2 + m_k$ 
  - The sparsity is squared at each step.
  - It follows that the algorithm requires  $O(\log \log N)$ .
- Valiant showed a matching lower bound and extended the results to show merging is  $\Theta(\log \log N)$  and sorting is  $\Theta(\log N)$  on a CRCW PRAM.
- See [slide 26](#) to see the details of the first few rounds.

Valiant's Algorithm, run-time

- The sparsity is roughly squared at each step.
- It follows that the algorithm requires  $O(\log \log N)$ .
- Valiant showed a matching lower bound and extended the results to show merging is  $\Theta(\log \log N)$  and sorting is  $\Theta(\log N)$  on a CRCW PRAM.
- See [slide 27](#) for the details.

Take-home message from Valiant's algorithm

- The PRAM model is simple, and elegant, and many clever algorithms have been designed based on the PRAM model.
- It is also physically unrealistic.
  - As shown on [slide 7](#), logic gates have bounded fan-in and fan-out.
  - Implementing the processor to memory interconnect requires a logic network of depth  $\Omega(\log P)$ .
  - Therefore, access time must be  $\Omega(\log P)$ .
  - Each step of the PRAM must take  $\Omega(\log P)$  physical time.
- Valiant's  $O(\log \log N)$  algorithms takes  $O(\log N \log \log N)$  physical time
  - It's **slower** than doing a simple reduce.
  - And it uses lots of communication – think of all those  $\lambda$  penalties!
  - But it's very clever. ☺
- Valiant understood this and pointed these issues in his paper.
  - But there has still been extensive research on PRAM algorithms.
  - It's an elegant model, what can I say?

Summary

- Simplified CTA reminds us that communication is expensive, but it doesn't explicitly charge for bandwidth.
- LogP accounts for bandwidth, but doesn't recognize that all bandwidth is not the same:
  - Communicating with an immediate neighbour is generally much cheaper than communicating with a distant machine.
  - Otherwise stated, the bisection bandwidth for real machines is generally much less than the per-machine bandwidth times the number of machines.
  - We can't have everyone talk at once at full bandwidth.
  - LogP uses the bisection bandwidth – this is conservative, but it doesn't recognize the advantages of local communication.
- Both are based on a 10-20 year old machine model
  - That's ok, the papers are 18-25 years old.
  - Doesn't account for the heterogeneity of today's parallel computers:
    - multi-core on chip, faster communication between processors on the same board than across boards, etc.
  - But recognize the limitations of any of these models.
- Getting a model of parallel computation that's as all-purpose as the RAM is still a work-in-progress.





- **EREW**: Exclusive-Read, Exclusive-Write
  - ▶ If two processors access the same location on the same step,
    - then the machine fails.
- **CREW**: Concurrent-Read, Exclusive-Write
  - ▶ Multiple machines can read the same location at the same time, and they all get the same value.
  - ▶ At most one machine can try to write a particular location on any given step.
  - ▶ If one processor writes to a memory location and another tries to read or write that location on the same step,
    - then the machine fails.
- **CRW**: Concurrent-Read, Concurrent-Write
  - ▶ If two or more machines try to write the same memory word at the same time, then if they are all writing the same value, that value will be written. (Otherwise (depending on the model),
    - the machine fails, or
    - one of the writes "wins", or
    - an arbitrary value is written to that address.

Mark Greenstreet      Models of Parallel Computation      CS 418 – Feb. 6, 2017      25 / 23

We could use reduce?

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license <http://creativecommons.org/licenses/by/4.0/>



Mark Greenstreet      Sorting Networks      CS 418 – Feb. 8, 2017      3 / 20

## Sorting Networks: Definition


Structural version:



network).

## Monotonicity Lemma

$y_{t-1} \rightarrow$     $\rightarrow$



Lemma: sorting networks commute with monotonic functions.

- Let  $S$  be a sorting network with  $n$  inputs and  $N$  outputs.

- Let  $S$  be a sorting network with  $n$  inputs and  $n$  outputs.
  - I'll write  $x_0, \dots, x_{n-1}$  to denote the inputs of  $S$ .
  - I'll write  $y_0, \dots, y_{n-1}$  to denote the outputs of  $S$ .
- Let  $f$  be a monotonic function.
  - If  $x \leq y$ , then  $f(x) \leq f(y)$ .
- The monotonicity lemma says
  - applying  $S$  and then  $f$  produces the same result as
  - applying  $f$  and then  $S$ .
- Observation:  $f(X)$  when  $X < X_j \rightarrow 0; f(\infty) \rightarrow 1$ .  
is monotonic.

|                  |                  |                       |         |
|------------------|------------------|-----------------------|---------|
| Mark Greenstreet | Sorting Networks | CS 418 – Feb. 8, 2017 | 11 / 20 |
|------------------|------------------|-----------------------|---------|

## The 0-1 Principle

1s, then it correctly sorts inputs of any values.

I'll prove the contrapositive.

- If a sorting network does not correctly sort inputs of any values, then it does not correctly sort all inputs consisting only of 0s and 1s

- Let  $S$  be a sorting network, let  $x$  be an input vector, and let  $y = S(x)$ , such that there exist  $i$  and  $j$  with  $i < j$  such that  $y_i > y_j$ .

- Let  $f(x) = \begin{cases} 0, & \text{if } x < y \\ 1, & \text{if } x \geq y \end{cases}$
- $\bar{y} = S(f(x))$
- By the definition of  $f$ ,  $f(x)$  is an input consisting only of 0s and 1s.
- By the monotonicity lemma,  $\bar{y} = f(y)$ . Thus,
 
$$\bar{y}_i = f(y_i) = 1 > 0 = f(y_i) = \bar{y}_i$$
- Therefore,  $S$  does not correctly sort an input consisting only of 0s and 1s.
- 

Mark Greenstreet      Sorting Networks      CS 418 – Feb 8, 2017      15 / 20

## Review 2

- 

Consider the two sorting networks shown above. One sorts correctly; the other does not.

- Identify the network that sorts correctly, and prove it using the 0-1 principle.
- Show that the other network does not sort correctly by giving an input consisting of 0s and 1s that is not sorted correctly.


|                  |                  |                       |         |
|------------------|------------------|-----------------------|---------|
| Mark Greenstreet | Sorting Networks | CS 418 – Feb. 8, 2017 | 19 / 20 |
|------------------|------------------|-----------------------|---------|



- The first compare-and-swap operates on  $y[b_2, b_2, b_1, 0]$  and  $y[b_2, b_2, b_1, 1]$ , for all 8 choices of  $b_2, b_2$ , and  $b_1$ .
- This corresponds to a compare-and-swap of  $x[b_0, b_1, b_2, 0]$  with  $x[b_0, b_1, b_2, 1]$ .
  - I'll call the result of the compare-and-swap  $z$  where
    - $z[b_0, b_2, b_1, 0] = \min(y[b_2, b_2, b_1, 0], y[b_2, b_2, b_1, 1]);$
    - $z[b_0, b_2, b_1, 1] = \max(y[b_2, b_2, b_1, 0], y[b_2, b_2, b_1, 1]);$
  - And I'll write  $\tilde{z}$  for  $z$  with "x indexing":
    - $\tilde{z}[b_2, b_2, b_1, b_0] = z[b_0, b_1, b_2, b_0];$
    - $\tilde{z}[0, b_2, b_1, b_0] = \min(x[0, b_2, b_1, b_0], x[1, b_2, b_1, b_0]);$
    - $\tilde{z}[1, b_2, b_1, b_0] = \max(x[0, b_2, b_1, b_0], x[1, b_2, b_1, b_0]);$
    - These are comparisons with a "stride" of 8 (for  $x$ ).

- The first shuffle takes  $z$  as an input and I'll call the output  $w$ .
- The first shuffle is a `shuffle4`; so
  - $w[i] = z[\text{rot1}(i, 2)];$
  - Equivalently,  $w[b_2, b_2, b_1, b_0] = z[b_2, b_2, b_1, b_0];$
- Let
  - $\tilde{w}[b_2, b_2, b_1, b_0] = w[b_2, b_1, b_2, b_0]$
  - $\tilde{w}[b_2, b_2, b_1, b_0] = z[b_2, b_2, b_1, b_0]$
- The second stage of compare-and-swap modules operates on
  - $w[b_2, b_2, b_1, 0]$  and  $w[b_2, b_2, b_1, 1]$
  - Equivalently,  $\tilde{w}[b_1, 0, b_2, b_2]$  and  $\tilde{w}[b_1, 1, b_2, b_2]$ .
  - These are comparisons with a stride of 4 for  $\tilde{w}$  and  $\tilde{w}$ .


- In the same way, the third stage of compare-and-swap modules operates has a stride of 2 for  $x$  indices,
- And the final stage has a stride of 1.
- More generally, to merge two sequences of length  $2^k$ :
  - Flip the lower sequence
  - Or, just sort it in reverse in the first place.
  - Perform compare-and-swap operations with stride  $L$ .
  - Perform compare-and-swap operations with stride  $L/2$  – note that these operate on pairs of elements whose indices differ in the  $L/2$  bit, and all of their other index bits are the same.
  - Perform compare-and-swap operations with stride  $L/4, \dots$
  - Perform compare-and-swap operations with stride 1 – this compares the element at  $2 \times i$  with the element at  $2 \times i + 1$  for  $0 \leq i < 2^k$ .
- Done!



## Flipping Out

- What should we do about the flips?
- Push them back (right-to-left) through the network
    - Keep track of how many flips we've accumulated.
    - Sort up for an even number of flips.
    - Sort down for an odd number of flips.
  - Flip the wiring in the bottom half of each unshuffle.
  - In practice:
    - Do the one that's easier for your implementation.

## Bitonic Sort



## Bitonic Sort in practice

- Sorting networks can be used to design practical sorting algorithms.
- To sort  $N$  values with  $P$  processors:
  - Divide input into  $2P$  segments of length  $\frac{N}{2P}$ .
  - Each processor sorts its pair of segments into one long segment.
    - The sorted segments are the inputs to the sorting network.
  - Now, follow the actions of the sorting network:
    - Processor  $i$  handles rows  $2i$  and  $2i+1$  of the sorting network.
    - Each compare-and-swap is replaced with "merge two sorted sequences and split into top half and bottom half."
    - When the sorting network has a compare-and-swap between rows  $2i$  and  $2i+1$ , each processor handles it locally.
    - When the sorting network has a compare-and-swap between rows  $2i$  and  $2i+K$  for  $K > 1$ , then processor  $i$  sends the upper half of its data to processor  $i+K/2$ , and processor  $i+K/2$  sends the lower half of its data to processor  $i$ . Both perform merges.
    - Note, if the compare-and-swap was flipped, then flip "upper-half" and "lower-half".

## Practical performance

- Complexity
  - Total number of comparisons:  $O(N \log N \log^2 P)$ .
  - Time:  $O(\frac{N}{P} (\log N + \log^2 P))$ , assuming each processor sorts  $N/P$  elements in  $O(N/P \log(N/P))$  time and merges two sequences of  $N/P$  elements in  $O(N/P)$  time.
- Remarks:
  - The idea of replacing compare-and-swap modules with processors that can perform merge using an algorithm optimized for the processor, is an extremely powerful and general one. It is used in the design of many practical parallel sorting algorithms.
  - Sorting networks are cool because they avoid branches:
    - Ideal for SIMD machines that can't really branch.
    - Need to experiment some to see the trade-offs of branch-divergence vs. higher asymptotic complexity on a GPU.

## Related Algorithms

- Counting Networks
  - How to match servers to requests.
- FFT
  - The Platonic Ideal of a Divide-and-Conquer Algorithm
  - Used for speech processing, signal processing, and lots of scientific computing tasks.

## Map-Reduce

Mark Greenstreet & Ian M. Mitchell  
CPSC 418 – February 27, 2017

- Problem Domain: Large-Scale Data Analysis
- The Map-Reduce Pattern
- Implementation Issues
- Results

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet & Ian Mitchell and are made available under the terms of the Creative Commons Attribution 4.0 International license <http://creativecommons.org/licenses/by/4.0/>

## Portrait of a Data Centre

- Sketch of a big data center:
  - Tens of thousands of machines, each with its own disk(s).
  - Distributed file system—what is that?
  - Commodity networks and routers.
    - Each machine has a network interface (e.g. 10Gb ethernet)
    - Cross-section bandwidth is way smaller than the number of machines times the per-machine bandwidth.
  - Scale is so big that there will be failures: chips, cores, memory, disks, network interfaces, switches, ...
    - If the average lifetime of a machine is five years, then 10,000 machine data center will have a failure every four hours.
    - Even without failure, maintenance will take machines offline.
- Need to analyse the huge data sets stored on these machines.

## Problem Domain: Large-Scale Data Analysis

- Data analysis requires:
  - Fetching the relevant records.
  - Performing analysis of related records.
  - Summarizing the results.
- Example: word frequency in documents
- Example: core curriculum
  - How do 200-level courses impact success in 400-level courses?
  - Look at all transcripts.
  - Analyze relationships for  $\langle 2XX, 4YY \rangle$  pairs.
- Google noted that each such problem was getting its own custom code.
  - All of that code development is expensive.
  - Often required similar rewrites when underlying system changed.

## The MapReduce Pattern

Slight generalization of description from [Dean & Ghemawat 2008].

- All data is represented as collections of  $(Key, Value)$  pairs.
- **map**
  - For each  $(Key1, Value1)$  pair of the input, user code produces a collection of  $(Key2, Value2)$  pairs for the output.
- **shuffle**
  - All  $(Key2, Value2)$  pairs with the same  $Key2$  are combined into a  $(Key2, [list\ of\ Value2])$  result and sorted by  $Key2$ .
- **reduce**
  - For each  $(Key2, [list\ of\ Value2])$ , user code produces a  $(Key2, Value3)$  result (where  $Value3$  might be a list itself).

Apache Hadoop is an open-source implementation of this basic framework.

## MapReduce: Word-Count

Example from [Dean & Ghemawat 2008] revised.

- Input data:
  - $Key1$  is the document name.
  - $Value1$  is a document text.
- **map**:
  - $Key2$  is a word.
  - $Value2$  is the count of times it appears in the document.
- **shuffle**:
  - Collect counts from all documents for each word  $(Word, [list\ of\ Counts])$ .
- **reduce**:
  - $Value3$  is the sum of all counts; in other words, the total number of times the word appears in all documents.

## MapReduce: Curriculum

- Input data  $(Key1, Value1)$ :
  - $Key1$  is the student number for the transcript
  - $Value1$  is a list of  $(CourseNumber, Grade)$  pairs.
- **map**: For each 200-level course and for each 400-level course in the transcript:
  - $Key2$  is the pair  $(Course200, Course400)$ .
  - $Value2$  is the pair  $(Grade200, Grade400)$ .
- **shuffle**:
  - Collect matching course pairs from all students  $((Course200, Course400), (Grade200, Grade400), (Grade200, Grade400), \dots)$ .
- **reduce**:
  - $Value3$  is (for example) the sample Pearson correlation coefficient  $r$  for the data set  $[(Grade200, Grade400), (Grade200, Grade400), \dots]$
  - More complex analyses could be performed.

## Wait a Minute Now...

But didn't we already study "reduce"?

- The course library's `wtree.reduce()` in Erlang had `leaf()`, `combine()` and `root()`.
- MapReduce at Google has  $(Key1, Value1)$ ,  $(Key2, Value2)$ , and shuffle?

These patterns have similar names but seek to solve different problems.

- **Reduce** is a generic name for a functional programming pattern which takes a collection of data and produces some kind of summary information.

## Many flavours of Reduce

- In Erlang, `wtree.reduce()` is designed to spread the computation of a reduction across many workers.
  - Implementation maximizes parallelism for a single reduce operation.
  - Collection and combination of data occurs in `combine()` functions.
  - Note that the `leaf()` function can perform a map operation before the reduction, so no loss of generality.
- In MapReduce, many independent reductions (one for each  $Key2$ ) are spread across many workers, but each reduction is performed by a single worker.
  - Implementation emphasizes fault tolerance and disk-based data storage.
  - Collection (but not combination) of data occurs in shuffle step.
  - If reduction is too big for a single worker, user must change the intermediate  $(Key2, Value2)$  representation and/or break the problem into multiple MapReduce steps.

## Programming Model

The user creates a **MapReduce specification object** which provides:

- The **map** and **reduce** functions.
- The names of the input and output files.
- Optionally other tuning parameters; for example:
  - Number of map and reduce workers to use.
  - Custom function to combine intermediate results within a map worker to reduce size of intermediate data.
  - A custom hashing function to help with shuffle step.

The user then invokes the **MapReduce** function.

## Execution (Part 1)

- The **MapReduce** function spawns  $M$  map worker,  $R$  reduce workers, and one master.
- Each map worker:
  - Is assigned fragment(s) of the input file by the master – these fragments are called **spits**.
  - Reads a  $(Key1, Value1)$  record from an assigned spit.
  - Runs user **map** code on that record.
  - Writes the  $(Key2, Value2)$  result to a temporary file.
  - Repeats until all records in the spit are completed.
  - Repeats until all assigned spits are completed.
  - Notifies the master when it is done.
- Result is a bunch of temporary files holding  $(Key2, Value2)$  pairs.

## Execution (Part 2)

- Start from temporary files holding  $(Key2, Value2)$  pairs.
- **Shuffle**:
  - Each reduce worker is assigned  $Key2$  value(s).
  - Corresponding  $Value2$  lists are taken from map worker's temporary output files and written to reduce worker's temporary input files.
  - Reduce worker receives  $(Key2, [list\ of\ Value2])$  pairs sorted by  $Key2$ .
- Each reduce worker:
  - Reads a  $(Key2, [list\ of\ Value2])$  record from a temporary file.
  - Runs user **reduce** code on that record.
  - Writes the  $(Key2, Value3)$  result to a file.
  - Notifies the master when it is done.
- When all the reduce computations are complete, the master sends a message to wake up the user process, and the **MapReduce** function returns.

## Do the MapReduce Shuffle

How do the intermediate results get from the map workers to the reduce workers? Simple version described in [Dean & Ghemawat, 2008]:

- Map workers know the number of reduce workers  $R$ .
- Each  $(Key2, Value2)$  is written to a different file according to `hash(Key2) mod R`.
- The master tells the reduce worker which file to read from each map worker.

Later versions of MapReduce utilized more complex or even user-specified hashing; for example to:

- Better balance size of reduce problems.
- Reduce network traffic and/or simplify sorting during shuffle step.
- Cluster certain  $Key2$  tuples onto the same reduce workers.

## Fault Tolerance

Bad things happen: Failed disks, partitioned networks, power shortages, ...

- Key Idea:
  - The **map** and **reduce** operations are based on functional programming ideas: referential transparency and no side-effects.
  - If a worker crashes, it is as if it never existed.
  - The master can restart the task on another machine.
- The master periodically pings the tasks, and restarts dead ones.
  - Map tasks produce only temporary files, so if a completed map task fails before informing the master then it must be re-executed.
  - Reduce tasks produce files in the distributed file system (redundant and fault-tolerant), so no need to re-execute.

## Semantics

- If the **map** and **reduce** functions are deterministic, then the result of **MapReduce** is the same as a sequential execution.
  - This is really cool!
  - There is a sequential implementation of **MapReduce**:
    - Read all of the  $(Key1, Value1)$  pairs from the input file.
    - Write all of the  $(Key2, [list\ of\ Value2])$  tuples to an intermediate file.
    - Sort the intermediate file by the  $Key2$  values.
    - Perform the reduce operation for each  $Key2$  value and write the results to the output file.
- If the **map** and **reduce** functions are not deterministic, then
  - It's a bit more complicated, but it's still reasonable.
  - If the reduce tasks are non-deterministic, then the result for each  $Key2$  is the result from some sequential implementation.
  - The paper doesn't talk about non-determinism for **map**, but it is probably similar.

## Work Stealing

- Sometimes a worker is slow – **stragglers**.
- If the **MapReduce** task is near completion, the master assigns straggler tasks to idle processors.
- These are called **backup tasks**.
- Either the original or the backup process can complete the task.
- In practice, this **work stealing** by backup tasks:
  - Only adds a few percent to the total compute resources used.
  - Can result in substantial performance improvements: The paper reported a 44% slow-down when the sort benchmark was run without backup tasks.

## Performance Issues

- The master attempts to schedule map tasks on the processor whose local disk holds the split being processed, or are nearby (by the network connections).
- Shuffle moves data from many map tasks to many reduce tasks.
  - Easily saturates the cross-section bandwidth of the network.
- For good performance, the map tasks should be filters that output much less data than they read.
  - Often not true of the "natural" intermediate representation (such as curriculum problem above).
  - Fewer distinct  $Key2$  values means less parallelism in reduce tasks.
- Can often reduce intermediate data size by partial reduction in the map workers.
  - May change the semantics of MapReduce (but not if reduce is associative and commutative).

## MapReduce is Designed for BIG Data

- Communication between standard linux machines with generic networks takes milliseconds.
- Reading large disk files takes seconds.
- The task needs to be big enough to justify these overheads:
  - Equivalent to increasing  $\lambda$  by a few orders of magnitude.
  - MapReduce makes sense if the task is disk-limited and harnessing a few thousand disks provides the necessary disk bandwidth.
    - Think of it as "disk parallelism" instead of "CPU parallelism".
    - Note: big-data companies like Amazon, Facebook and Google are moving to using FLASH memory and DRAM instead of disks, exactly because of these I/O bottlenecks.
  - Or if you have a really huge data set the compute time may dominate all of these overheads.

## Results (Part 1)

Achieves impressive performance on massive data sets 2008–2013(?)

- Report in [Dean & Ghemawat, 2008]: Good performance on ~ 2000 machines: **grep** and **sort** work through  $10^{13}$  100-byte records (1TB) in minutes.
- Google estimates ~ 20PB / day in total MapReduce processing in January 2008.
- Google research blog reports sorting  $10^{13}$  100-byte records (1PB) on ~ 4000 machines (and ~ 48,000 disks) in six hours in November 2008, then 33 minutes for 1 PB or 6 hours for 10 PB on 8000 machines in September 2011.
- Open source implementation in Hadoop widely available as a cloud service.
- Many example algorithms documented; for example, search for "map reduce" on <http://scholar.google.ca>.

## Results (Part 2)

Big data processors are now moving away from MapReduce.

- Framework emphasizes batch processing of data residing in distributed file system, which limits flexibility and efficiency.
- "We don't really use MapReduce anymore" [Urs H olzie, senior vice president of technical infrastructure at Google [speaking at Google I/O conference in 2014](#)]
- Machine learning project Apache Mahout is shifting away from MapReduce algorithms to alternatives such as Apache Spark.
  - Worth noting: Spark also uses a functional programming model with referential transparency.



- MapReduce is a parallel programming pattern
  - Data are represented by collections of *(Key, Value)* pairs.
  - User provides *map* to take input *(Key1, Value1)* pairs to intermediate *(Key2, Value2)* pairs.
  - Shuffle step collects intermediate data into *(Key2, [list of Value2])* pairs and sorts it by *Key2*.
  - User provides *reduce* to take sorted *(Key2, [list of Value2])* pairs into *(Key2, Value3)* pairs.
- Details of the parallel implementation are handled by the MapReduce API:
  - Creating workers processes, delivering input and output files, shuffling intermediate data between map and reduce workers.
  - Handling failures and slow nodes.
- Performance is often bandwidth limited.
  - Locality matters: perform *map* on the machine with the data.
  - If *map* is an effective filter (bytes out < bytes in), then we can reduce the impact of network congestion.
  - In practice, user must choose *(Key2, Value2)* representation wisely to trade-off shuffle effort for reduce parallelism.

- How does MapReduce distribute work between map tasks?
- How does MapReduce distribute work between reduce tasks?
- How does MapReduce handle machine, network or other failures?
- How does MapReduce handle slow (i.e. straggler) machines?
- What are the requirements for the type-signatures of the *map* and *reduce* functions in a map-reduce?

I want to fly from Vancouver to Timbuktu. There are no direct flights, so I want to find the fastest route with one stop. How could I do this using map reduce?

- Input data is a table of airline flights of the form:
 

```
(DepartCity, DepartTime, ArriveCity, ArriveTime)
```

  - Hint: use the intermediate city as *Key2*.
  - For simplicity, assume that all times are GMT (no need for time-zone conversion).
  - How does *map* filter out irrelevant flights?


Mark Greenstreet

CpSc 418 – Mar. 3, 2017

- GPUs
  - Early geometry engines.
  - Adding functionality and programmability.
  - GPGPUs
- CUDA
  - Execution Model
  - Memory Model
  - A simple example

## Why is dedicated hardware so much faster?

A simple multiplier



Latency and period are 2N.

## The GPGPU

General Purpose Graphics Processing Unit

- The volume market is for graphics, and the highest profit is GPUs for high-end gamers.
  - Most of the computation is floating point.
  - Latency doesn't matter.
  - Abundant parallelism.
- Make the architecture fit the problem:
  - SIMD – single instruction, multiple (parallel) data streams.
    - Amortize control overhead over a large number of functional units.
    - They call it SMT (... multiple threads) because they allow conditional execution.
  - High-latency operations
    - Allows efficient, high-throughput, high-latency floating point units.
    - Allows high latency accesses to off-chip memory.
- This means lots of threads per processor.

## Before the first GPU

Early 1980's: bit-blit hardware for simple 2D graphics.

- Draw lines, simple curves, and text.
- Fill rectangles and triangles.
- Color used a "color map" to save memory:
  - bit-wise logical operations on color map indices!

## 1989: The SGI Geometry Engine

- Basic rendering: coordinate transformation.
  - Represent a 3D point with a 4-element vector.
  - The fourth element is 1, and allows translations.
  - Multiply vector by matrix to perform coordinate transformation.
- Dedicated hardware is much more efficient than a general purpose CPU for matrix-vector multiplication.
  - For example, a  $32 \times 32$  multiplier can be built with  $32^2 = 1024$  one-bit multiplier cells.
    - A one-bit multiplier cell is about 50 transistors.
    - That's about 50K transistors for a very simple design.
  - 30K is quite feasible using better architectures.

## 1989: The SGI Geometry Engine

- Basic rendering: coordinate transformation.
- Dedicated hardware is much more efficient than a general purpose CPU for matrix-vector multiplication.
  - For example, a  $32 \times 32$  multiplier can be built with  $32^2 = 1024$  one-bit multiplier cells.
    - A one-bit multiplier cell is about 50 transistors.
    - That's about 50K transistors for a very simple design.
  - 30K is quite feasible using better architectures.
- The 80486DX was 1.2M transistors, 16MHz, 13MIPs.
  - The 80486DX was 1.2M transistors, 16MHz, 13MIPs.
  - That's equal to 24 dedicated multipliers.
  - 16 multiply-and-accumulate units running at 50MHz (easy in the same  $1\mu$  process) produce 1.6GFlops!

## Building a better multiplier

- Simple multiplier has latency and period of 2N.
- Pipelining: add registers between rows.
  - The period is N, but the latency is  $N^2$ .
  - The bottleneck is the time for carries in each row.
- Use carry-lookahead adders (compute carries with a scan)
  - period is  $\log N$ , the latency is  $N \log N$ .
  - but the hardware is more complicated.
- Use carry-save adders and one carry-lookahead at the end
  - each adder in the multiplier forwards its carries to the next adder.
  - the final adder resolves the carries.
  - period is 1, latency is N.
  - and the hardware is *way* simpler than a carry-lookahead design
- Graphics and many scientific and machine learning computations are very tolerant of latency.

## Why is dedicated hardware so much faster?

Example: matrix-vector multiplication


- addition and multiplication are "easy".
- it's the rest of CPU that's complicated and the usual performance bottleneck
  - memory read and write
  - instruction fetch, decode, and scheduling
  - pipeline control
  - handling exceptions, hazards, and speculation
  - etc.
- GPU architectures amortize all of this overhead over a lot of execution units.

## The fundamental challenge of graphics


Human vision isn't getting any better.

- Once you can perform a graphics task at the limits of human perception (or the limits of consumer budget for monitors), then there's no point in doing it any better.
- Rapid advances in chip technology meant that coordinate transformations (the specialty of the SGI Geometry Engine) were soon as fast as anyone needed.
- Graphics processors have evolved to include more functions. For example,
  - Shading
  - Texture mapping
- This led to a change from hardwired architectures, to programmable ones.

## The Fermi Architecture




## What does a core look like?




- RISC pipeline: see Jan 23 slides (e.g. slides 5ff.)
  - Instruction fetch, decode and other control takes much more power than actually performing ALU and other operations!
- SIMD: Single-Instruction, Multiple-Data
- What about memory?

## What does a core look like?



- RISC pipeline: see Jan 23 slides (e.g. slides 5ff.)
- SIMD: Single-Instruction, Multiple-Data
  - Multiple execution pipelines execute the same instructions.
  - Each pipeline has its own registers and operates on separate data values.
  - Commonly, pipelines access adjacent memory locations.
  - Great for operating on matrices, vectors, and other arrays.
- What about memory?

## What does a core look like?



- RISC pipeline: see Jan 23 slides (e.g. slides 5ff.)
- SIMD: Single-Instruction, Multiple-Data
- What about memory?
  - On-chip "shared memory" switched between cores: see Jan 25 slides (e.g. slide 3)
  - Off-chip references are "coalesced": the hardware detects reads from (or writes to) consecutive locations and combines them into larger, block transfers.

## More about GPU Cores

- Execution pipeline can be very deep – 20-30 stages.
  - Many operations are floating point and take multiple cycles.
  - A floating point unit that is deeply pipelined is easier to design, can provide higher throughput, and use less power than a lower latency design.
- No bypasses
  - Instructions block until instructions that they depend on have completed execution.
  - GPUs rely on extensive multi-threading to get performance.
- Branches use predicated execution.
  - Execute the then-branch code, disabling the "else-branch" threads.
  - Execute the else-branch code, disabling the "then-branch" threads.
  - The order of the two branches is unspecified.
- Why?
  - All of these choices optimize the hardware for graphics applications.
  - To get good performance, the programmer needs to understand how the GPGPU executes programs.


## Lecture Outline

- GPUs
  - been there, done that.
- CUDA – we are here!
  - Execution Model
  - Memory Model
  - Code Snippets

## Execution Model: Functions

- A CUDA program consists of three kinds of functions:
  - Host functions:
    - callable from code running on the host, but not the GPU.
    - run on the host CPU.
    - In CUDA C, these look like normal functions.
  - Device functions.
    - callable from code running on the GPU, but not the host.
    - run on the GPU.
    - In CUDA C, these are declared with a `__device__` qualifier.
  - Global functions
    - called by code running on the host CPU, they execute on the GPU.
    - In CUDA C, these are declared with a `__global__` qualifier.

## Execution Model: Memory



- Host memory: DRAM and the CPU's caches
  - Accessible to host CPU but not to GPU.
- Device memory: GDDR DRAM on the graphics card.
  - Accessible by GPU.
  - The host can initiate transfers between the host memory and device memory.
- The CUDA library includes functions to:
  - Allocate and free device memory.
  - Copy blocks between host and device memory.
  - BUT** host code can't read or write the device memory directly.

## Structure of a simple CUDA program

- A `__global__` function to be called by the host program to execute on the GPU.
  - There may be one or more `__device__` functions as well.
- One or more host functions, including `main` to run on the host CPU.
  - Allocate device memory.
  - Copy data from host memory to device memory.
  - "Launch" the device kernel by calling the `__global__` function.
  - Copy the result from device memory to host memory.
- We'll do the `saxpy` example from the paper.
  - `saxpy = "Scalar s times x plus y"`.

## saxpy: device code

```
__global__ void saxpy(uint n, float a, float *x, float *y) {
 uint i = blockIdx.x * blockDim.x + threadIdx.x; // nvec built-ins
 if(i < n)
 y[i] = a*x[i] + y[i];
}
```

- Each thread has `x` and `y` indices.
  - We'll just use `x` for this simple example.
- Note that we are creating one thread per vector element:
  - Exploits GPU hardware support for multithreading.
  - We need to keep in mind that there are a large, but limited number of threads available.

## saxpy: host code (part 1 of 5)

```
int main(int argc, char **argv) {
 uint n = atoi(argv[1]);
 float *x, *y, *yy;
 float *dev_x, *dev_y;
 int size = n*sizeof(float);
 x = (float *)malloc(size);
 y = (float *)malloc(size);
 yy = (float *)malloc(size);
 for(int i = 0; i < n; i++) {
 x[i] = i;
 y[i] = i+1;
 }
 ...
}
```

- Declare variables for the arrays on the host and device.
- Allocate and initialize values in the host array.

## saxpy: host code (part 2 of 5)

```
int main(void) {
 ...
 cudaMalloc((void**)(&dev_x), size);
 cudaMalloc((void**)(&dev_y), size);
 cudaMemcpy(dev_x, x, size, cudaMemcpyHostToDevice);
 cudaMemcpy(dev_y, y, size, cudaMemcpyHostToDevice);
 ...
}
```

- Allocate arrays on the device.
- Copy data from host to device.

## saxpy: host code (part 3 of 5)

```
int main(void) {
 ...
 float a = 3.0f;
 saxpy<<ceil(n/256.0), 256>>>(n, a, dev_x, dev_y);
 cudaMemcpy(yy, dev_y, size, cudaMemcpyDeviceToHost);
 ...
}
```

- Invoke the code on the GPU:
  - `saxpy<<ceil(n/256.0), 256>>>(n, a, dev_x, dev_y)` says to create  $\lceil n/256 \rceil$  blocks of threads.
  - Each block consists of 256 threads.
  - See slide 22 for an explanation of threads and blocks.
  - The pointers to the arrays (in device memory) and the values of `n` and `a` are passed to the threads.
- Copy the result back to the host.

## saxpy: host code (part 4 of 5)

```
...
for(int i = 0; i < n; i++) { // check the result
 if(yy[i] != a*x[i] + y[i]) {
 printf("at addr: %d, a[i]=%f, b[i]=%f, c[i]=%f\n",
 i, a[i], b[i], c[i]);
 exit(-1);
 }
}
printf("The results match!\n");
...
}
```

- Check the results.

## saxpy: host code (part 5 of 5)

```
int main(void) {
 ...
 Free(x);
 Free(y);
 Free(yy);
 cudaFree(dev_x);
 cudaFree(dev_y);
 exit(0);
}
```

- Clean up.
- We're done.

## Threads and blocks

- Our example created  $\lceil \frac{n}{256} \rceil$  blocks with 256 threads each.
- The GPU hardware has a pool of running threads.
  - Each thread has a "next instruction" pending execution.
  - If the dependencies for the next instruction are resolved, the "next instruction" can execute.
  - The hardware in each streaming multiprocessor dispatches an instruction each clock cycle if a ready instruction is available.
  - The GPU in 11n25 supports 1024 such threads.
- What if our application needs more threads?
  - Threads are grouped into "thread blocks".
  - Each thread block has up to 1024 threads (the HW limit).
  - The GPU can swap thread-block in and out of main memory
    - This is GPU system software that we don't see as user-level programmers.

- For this example, not really.
  - Execution time dominated by the memory copies.
- But, it shows the main pieces of a CUDA program.
- To get good performance:
  - We need to perform many operations for each value copied between memories.
  - We need to perform many operations in the GPU for each access to global memory.
  - We need enough threads to keep the GPU cores busy.
  - We need to watch out for **thread divergence**:
    - If different threads execute different paths on an if-then-else, then the else-threads stall while the then-threads execute, and vice-versa.
  - And many other constraints.
- GPUs are great if your problem matches the architecture.

|                                                              |         |                  |
|--------------------------------------------------------------|---------|------------------|
| March 8: Intro to CUDA                                       | Reading | Kirk & Huu Ch. 2 |
| March 8: CUDA Threads, Part 1                                | Reading | Kirk & Huu Ch. 3 |
| March 10: CUDA Threads, Part 2                               | Reading | Kirk & Huu Ch. 4 |
| March 13: CUDA Memory                                        | Reading | Kirk & Huu Ch. 5 |
| March 15: CUDA Memory: examples                              | Reading | Kirk & Huu Ch. 5 |
| March 17: CUDA Performance                                   | Reading | Kirk & Huu Ch. 5 |
| March 20: Matrix multiplication with CUDA, Part 1            | Reading | Kirk & Huu Ch. 5 |
| March 22: Matrix multiplication with CUDA, Part 2            | Reading | Kirk & Huu Ch. 5 |
| March 24 – April 3: Other Topics                             | Reading | Kirk & Huu Ch. 5 |
| April 5: Party: 50 <sup>th</sup> Anniversary of Amdahl's Law | Reading | Kirk & Huu Ch. 5 |

- What is SIMD parallelism?
- How does a CUDA GPU handle branches?
- How does a CUDA GPU handle pipeline hazards?
- What is the difference between "shared memory" and "global memory" in CUDA programming.
- Think of a modification to the *saxpy* program and try it.
  - You'll probably find you're missing programming features for many things you'd like to try.
  - What do you need?
  - Stay tuned for upcoming lectures.

Mark Greenstreet

CpSc 418 – Mar. 6, 2017

- GPU Summary: [slide 2](#)
- CUDA
  - Data parallelism: [slide 6](#)
  - Program structure: [slide 8](#)
  - Memory: [slide 10](#)
  - A simple example: [slide 12](#)
  - Launching kernels: [slide 19](#)

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

|                  |                         |                       |         |
|------------------|-------------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to GPU/CPU | CS 418 – Mar. 3, 2017 | 23 / 25 |
|------------------|-------------------------|-----------------------|---------|

## GPU Summary: architecture

- Lots of cores:
  - Up to 90 or more SIMD processors.
  - Each SIMD processors has 32 pipelines.
  - This is the nVidia architecture – other GPUs are similar.
- Deep, simple, execution pipelines
  - Optimized for floating point.
  - No bypassing: use multi-threading for performance.
  - Branches handled by predicated execution

*"When you come to a fork in the road, take it."  
(Often attributed to Yogi Berra.)*
- Limited on-chip memory.
  - 1 or 2 MBytes total. Big CPUs have 32-64MB of L3 cache.
  - The programmer manages data placement.

|                  |                         |                       |         |
|------------------|-------------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to GPU/CPU | CS 418 – Mar. 3, 2017 | 24 / 25 |
|------------------|-------------------------|-----------------------|---------|

## GPU Summary: Performance

- Today's processors are constrained by how much performance can you get using ~ 200 watts.
  - Moving bits around takes lots of energy.
  - Performing operations as quickly as possible takes lots of energy.
  - $E \sim dt^{-\alpha}$ , where  $E$  is energy,  $d$  is distance,  $t$  is time per operation, and  $1 < \alpha < 2$  depending on design details.
    - Corollary:  $P \sim d^{-\alpha}$ . Power grows someplace between quadratically and cubically with clock frequency.
- How GPUs optimize performance/power
  - SIMD: instruction fetch and decode moves lots of bits. Amortize over many cores.
  - Simple pipelines: bypassing means moving bits quickly. GPUs omit bypasses.
  - High latency: avoid pipeline stages that must do a lot in a hurry.
  - Expose the memory hierarchy: let the programmer control moving data bits around.

|                  |                         |                       |         |
|------------------|-------------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to GPU/CPU | CS 418 – Mar. 3, 2017 | 25 / 25 |
|------------------|-------------------------|-----------------------|---------|

## GPU Summary: Economics

- GPUs are designed for the high-volume, consumer graphics market.
  - Amortize high design cost over a large number of units sold.
- This means GPUs aren't really optimized for scientific computing:
  - More on-chip memory would certainly help scientific computing, but not needed for graphics rendering.
  - Comparison: An nVidia GPU has about 2 MBytes of on-chip memory, an Intel Xeon can have 40MBytes or more.
  - Cache memory is about 60-70 transistors per byte.
  - A high-end nVidia GPU has 7 billion transistors, 1 or 2% for memory.
  - What if the chip were 30-40% memory?
    - better for general purpose computing
    - little pay-off for graphics
    - smaller distinction with Intel CPUs
- Cheap is good
  - It's the economics of cheap-computing that drives Moore's Law and all the other exponential growth-rate trends that make computing a field of intense, ongoing innovation.
  - That keeps the field in transition – deal with it.

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 1 / 24 |
|------------------|----------------------|-----------------------|--------|

## Programming GPUs: CUDA

- [Data Parallelism](#)
- [CUDA program structure](#)
- [Memory](#)
- [Launching kernels](#)

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 2 / 24 |
|------------------|----------------------|-----------------------|--------|

## Data Parallelism

- When you see a `for`-loop:
  - Is the loop-index used as an array index?
  - Are the iterations independent?
  - If so, you probably have data-parallel code.
- Data-Parallel problems:
  - Run well on GPUs because each element (or segment) of the array can be handled by a different thread.
  - Data parallel problems are good candidate for most parallel techniques because the available parallelism grows with the problem size.
  - Compare with "task parallelism" where the problem is divided into the same number of tasks regardless of its size.

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 3 / 24 |
|------------------|----------------------|-----------------------|--------|

## Which of the following loops are data parallel?

```
for(int i = 0; i < N; i++)
 c[i] = a[i] + b[i];

dotprod = 0.0;
for(int i = 0; i < N; i++)
 dotprod += a[i]*b[i];

for(int i = 1; i < N; i++)
 a[i] = 0.5*(a[i-1] + a[i]);

for(int i = 1; i < N; i++)
 a[i] = sqrt(a[i-1] + a[i]);

for(int i = 0; i < N; i++) {
 for(int j = 0; j < N; j++) {
 sum = 0.0;
 for(int k = 0; k < L; k++)
 sum += a[i,k]*b[k,j];
 c[i,j] = sum;
 }
}
```

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 4 / 24 |
|------------------|----------------------|-----------------------|--------|

## CUDA Program Structure

- A CUDA program consists of three kinds of functions:
  - Host functions:
    - callable from code running on the host, but not the GPU.
    - run on the host CPU;
    - In CUDA C, these look like normal functions – they can be preceded by the `__host__`-qualifier.
  - Device functions.
    - callable from code running on the GPU, but not the host.
    - run on the GPU;
    - In CUDA C, these are declared with a `__device__`-qualifier.
  - Global functions.
    - called by code running on the host CPU,
    - they execute on the GPU.
    - In CUDA C, these are declared with a `__global__`-qualifier.


|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 5 / 24 |
|------------------|----------------------|-----------------------|--------|

## Structure of a simple CUDA program

- A `__global__` function to be called by the host program to execute on the GPU.
  - There may be one or more `__device__` functions as well.
- One or more host functions, including `main` to run on the host CPU.
  - Allocate device memory.
  - Copy data from host memory to device memory.
  - "Launch" the device kernel by calling the `__global__` function.
  - Copy the result from device memory to host memory.

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 6 / 24 |
|------------------|----------------------|-----------------------|--------|

## Execution Model: Memory



- Host memory: DRAM and the CPU's caches
  - Accessible to host CPU but not to GPU.
- Device memory: GDDR DRAM on the graphics card.
  - Accessible by GPU.
  - The host can initiate transfers between host memory and device memory.
- The CUDA library includes functions to:
  - Allocate and free device memory.
  - Copy blocks between host and device memory.
  - BUT** host code can't read or write the device memory directly.

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 7 / 24 |
|------------------|----------------------|-----------------------|--------|

## More Memory

- GPUs support fairly large off-chip memory bandwidth: 200-400GB/s.
  - But this isn't fast enough to keep 1000 processors busy at 1Gflop/s each!
- The GPU has on-chip memory to help:
  - Shared memory: 16KBytes or 48KBytes.
  - Registers: 128KBytes (256KBytes on more recent GPUs).
  - Note that we need to use each value from memory for 20 or more instructions or else the memory bandwidth will limit performance.
- GPUs also have L2 caches, around 1.5MByte in the most recent chips.
  - But I haven't found a good way to understand them from the textbook, or from other CUDA manuals.
  - The coherence/consistency guarantees seem to be pretty weak.

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 8 / 24 |
|------------------|----------------------|-----------------------|--------|

## Example: saxpy

saxpy = "Scalar  $\times$  times  $x$  plus  $y$ ."

- [The device code.](#)
- [The host code.](#)
- [The running saxpy](#)

|                  |                      |                       |        |
|------------------|----------------------|-----------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 9 / 24 |
|------------------|----------------------|-----------------------|--------|

## saxpy: device code

```
__global__ void saxpy(int n, float a, float *x, float *y) {
 uint i = blockIdx.x*blockDim.x + threadIdx.x; // nvec h/w
 if(i < n)
 y[i] = a*x[i] + y[i];
}
```

- Each thread has  $x$  and  $y$  indices.
  - We'll just use  $x$  for this simple example.
- Note that we are creating one thread per vector element:
  - Exploits GPU hardware support for multithreading.
  - We need to keep in mind that there are a large, but limited number of threads available.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 10 / 24 |
|------------------|----------------------|-----------------------|---------|

## saxpy: host code (part 1 of 5)

```
int main(int argc, char **argv) {
 uint n = atoi(argv[1]);
 float *x, *y, *yy;
 float *devx, *devy;
 int size = n*sizeof(float);
 x = (float *)malloc(size);
 y = (float *)malloc(size);
 yy = (float *)malloc(size);
 for(int i = 0; i < n; i++) {
 x[i] = i;
 y[i] = i+1;
 }
 ...
}
```

- Declare variables for the arrays on the host and device.
- Allocate and initialize values in the host array.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 11 / 24 |
|------------------|----------------------|-----------------------|---------|

## saxpy: host code (part 2 of 5)

```
int main(void) {
 ...
 cudaMalloc((void**)&devx, size);
 cudaMalloc((void**)&devy, size);
 cudaMemcpy(devx, x, size, cudaMemcpyHostToDevice);
 cudaMemcpy(devy, y, size, cudaMemcpyHostToDevice);
 ...
}
```

- Allocate arrays on the device.
- Copy data from host to device.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 12 / 24 |
|------------------|----------------------|-----------------------|---------|

## saxpy: host code (part 3 of 5)

```
int main(void) {
 ...
 float a = 3.0;
 saxpy<<<cell(n/256.0, 256)>>>(n, a, devx, devy);
 cudaMemcpy(yy, devy, size, cudaMemcpyDeviceToHost);
 ...
}
```

- Invoke the code on the GPU:
  - `add<<<cell(n/256.0, 256)>>>(...)` says to create  $\lceil n/256 \rceil$  blocks of threads.
  - Each block consists of 256 threads.
  - See [slide 20](#) for an explanation of threads and blocks.
  - The pointers to the arrays (in device memory) and the values of  $n$  and  $a$  are passed to the threads.
- Copy the result back to the host.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 13 / 24 |
|------------------|----------------------|-----------------------|---------|

## saxpy: host code (part 4 of 5)

```
...
for(int i = 0; i < n; i++) { // check the result
 if(yy[i] != a*x[i] + y[i]) {
 printf("saxder,
 *%s: i=%d, a[i]=%f, b[i]=%f, c[i]=%f\n",
 i, a[i], b[i], c[i]);
 exit(-1);
 }
}
printf("The results match!\n");
...
}
```

- Check the results.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 14 / 24 |
|------------------|----------------------|-----------------------|---------|

## saxpy: host code (part 5 of 5)

```
int main(void) {
 ...
 free(x);
 free(y);
 free(yy);
 cudaFree(devx);
 cudaFree(devy);
 exit(0);
}
```

- Clean up.
- We're done.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 15 / 24 |
|------------------|----------------------|-----------------------|---------|

## Launching Kernels

- Terminology
  - Data parallel code that runs on the GPU is called a **kernel**.
  - Invoking a GPU kernel is called **launching** the kernel.
- How to launch a kernel
  - The host CPU invokes a `__global__` function.
  - The invocation needs to specify how many threads to create.
  - Example:
    - `add<<<cell(n/256.0, 256)>>>(...)`
    - creates [256](#) blocks
    - with 256 threads each.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 16 / 24 |
|------------------|----------------------|-----------------------|---------|

## Threads and Blocks

- The GPU hardware combines threads into **warps**
  - Warps are an aspect of the hardware.
  - All of the threads of warp execute together – this is the SIMD part.
  - The functionality of a program doesn't depend on the warp details.
  - But understanding warps is critical for getting good performance.
- Each warp has a "next instruction" pending execution.
  - If the dependencies for the next instruction are resolved, it can execute for all threads of the warp.
  - The hardware in each streaming multiprocessor dispatches an instruction each clock cycle if a ready instruction is available.
  - The GPU in `lin25` supports 32 such warps of 32 threads each in a "thread block".
- What if our application needs more threads?
  - Threads are grouped into "thread blocks".
  - Each thread block has up to 1024 threads (the HW limit).
  - The GPU can swap thread-blocks in and out of main memory
    - This is GPU system software that we don't see as user-level programmers.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 17 / 24 |
|------------------|----------------------|-----------------------|---------|

## Compiling and running

```
lin255 nvcc saxpy.cu -o saxpy
lin255 ./saxpy 1000
The results match!
```

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 22 / 24 |
|------------------|----------------------|-----------------------|---------|

## But is it fast?

- For the *saxpy* example as written here, not really.
  - Execution time dominated by the memory copies.
- But, it shows the main pieces of a CUDA program.
- To get good performance:
  - We need to perform many operations for each value copied between memories.
  - We need to perform many operations in the GPU for each access to global memory.
  - We need enough threads to keep the GPU cores busy.
  - We need to watch out for **thread divergence**:
    - If different threads execute different paths on an if-then-else, then the else-threads stall while the then-threads execute, and vice-versa.
  - And many other constraints.
- GPUs are great if your problem matches the architecture.

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 23 / 24 |
|------------------|----------------------|-----------------------|---------|

## Preview

|                                                              |         |                                                                      |
|--------------------------------------------------------------|---------|----------------------------------------------------------------------|
| March 8: CUDA Threads, Part 1                                | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 3 (Ch. 4 in 2 <sup>nd</sup> ed.) |
| March 10: CUDA Threads, Part 2                               | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 3 (Ch. 4 in 2 <sup>nd</sup> ed.) |
| March 13: CUDA Memory                                        | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 4 (Ch. 5 in 2 <sup>nd</sup> ed.) |
| March 15: CUDA Memory: examples                              | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 4 (Ch. 5 in 2 <sup>nd</sup> ed.) |
| March 17: CUDA Performance                                   | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 5 (Ch. 6 in 2 <sup>nd</sup> ed.) |
| March 20: Matrix multiplication with CUDA, Part 1            | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 5 (Ch. 6 in 2 <sup>nd</sup> ed.) |
| March 22: Matrix multiplication with CUDA, Part 2            | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 5 (Ch. 6 in 2 <sup>nd</sup> ed.) |
| March 24 – April 3: Other Topics                             | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 5 (Ch. 6 in 2 <sup>nd</sup> ed.) |
| April 5: Party: 50 <sup>th</sup> Anniversary of Amdahl's Law | Reading | Kirk & Huu 3 <sup>rd</sup> ed., Ch. 5 (Ch. 6 in 2 <sup>nd</sup> ed.) |

|                  |                      |                       |         |
|------------------|----------------------|-----------------------|---------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 24 / 24 |
|------------------|----------------------|-----------------------|---------|

## Review

- What is SIMD parallelism?
- What is the difference between "shared memory" and "global memory" in CUDA programming.
- Think of a modification to the *saxpy* program and try it.
  - You'll probably find you're missing programming features for many things you'd like to try.
  - What do you need?
  - Stay tuned for upcoming lectures.

Mark Greenstreet &amp; Ian M. Mitchell

CPSC 418 – March 8 &amp; March 10, 2017

- [Kernel organization: grids, blocks & threads.](#)
- [Hardware organization: SMs, SPs & warps.](#)

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet & Ian M. Mitchell and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

|                  |                      |                       |         |                  |                      |                       |         |                  |                      |                       |         |                        |              |                              |        |
|------------------|----------------------|-----------------------|---------|------------------|----------------------|-----------------------|---------|------------------|----------------------|-----------------------|---------|------------------------|--------------|------------------------------|--------|
| Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 22 / 24 | Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 23 / 24 | Mark Greenstreet | Introduction to CUDA | CS 418 – Mar. 6, 2017 | 24 / 24 | Greenstreet & Mitchell | CUDA Threads | CPSC 418 – Mar. 8 & 10, 2017 | 1 / 20 |
|------------------|----------------------|-----------------------|---------|------------------|----------------------|-----------------------|---------|------------------|----------------------|-----------------------|---------|------------------------|--------------|------------------------------|--------|


- Lots of nVidia jargon here.
- Lots of very specific constraints on hardware capabilities.
- Values of those constraints depend on the *compute capability*: essentially a version number for the GPU hardware.
  - CS department lab (`lin01.lin02...lin05.ucsd.edu`) has GeForce GTX 550 Ti which feature compute capability 2.1.
  - Examples of recent GPUs:
    - Compute capability 3.5: GT 730 & GTX 780.
    - Compute capability 5.0: GTX 750, 840M & 980M.
    - Compute capability 5.2: GTX 960, 980M.
    - Compute capability 6.1: GTX 10xx.
  - More details at the CUDA wikipedia page.

- When a kernel is launched, it creates a collection of threads.
- This collection is called a **grid**.
  - A grid is organized as an array of **blocks**
  - Each block is an array of **threads**
  - Array sizes are fixed once a kernel is launched.
- Why so many details?
  - Switching between blocks is done (I infer) by software in the GPU.
  - Switching between threads in a block is done by hardware.
  - By distinguishing blocks from threads, the CUDA model exposes the performance issues to the programmer.

|       |       |       |       |       |       |
|-------|-------|-------|-------|-------|-------|
| (0,0) | (1,0) | (2,0) | (3,0) | (4,0) | (5,0) |
| (0,1) | (1,1) | (2,1) | (3,1) | (4,1) | (5,1) |
| (0,2) | (1,2) | (2,2) | (3,2) | (4,2) | (5,2) |
| (0,3) | (1,3) | (2,3) | (3,3) | (4,3) | (5,3) |

A grid

- Blocks are scheduled by the GPU **software**.
- Blocks can be arranged as 1D, 2D or 3D array.
  - Dimensions are called "x", "y" and "z".
- There can be **lots** of blocks:
  - Each dimension can be up to  $2^{16} - 1 = 65535$ .
  - CC 3.0+ allows x dimension up to  $2^{31} - 1$  blocks.



Where do they put all those threads?

- Threads are scheduled by the GPU **hardware**.
- Threads can be arranged as a 1D, 2D, or 3D array
  - Grid and block dimensions and sizes may be different.
- There can be a moderate number of threads in each dimension:
  - x or y up to 1024 threads.
  - z up to 64 threads.
- However, total number of threads per block (product of all dimensions) is also capped at 1024.

### Threads and blocks: launching a kernel

- Let's say we have:

```
__global__ void kernelFun(args)
```
- To launch this kernel, we execute a statement like:

```
kernelFun<<dimGrid, dimBlock>>>(actuals);
```

where
  - *dimGrid* specifies the dimension(s) of the grid (an array of blocks):
    - *dimGrid* can be an int, in which case the array is 1D.
    - *dimGrid* can be a dim3, for example:

```
dim3(6,4,1)
```
  - *dimBlock* specifies the dimension(s) of each block (an array of threads):
    - *dimBlock* can be an int or a dim3.

### Threads and Blocks within a Kernel's Grid

- Within a running kernel, CUDA-C provides four built-in variables to determine the position of a thread within the grid: *blockDim*, *blockIdx*, *threadDim*, and *threadIdx*.
- There is a naming pattern:
  - Each of these structures has three fields: x, y and z corresponding to the three possible dimensions.
  - *blockDim*.? gives the size of the grid in each dimension x, y or z.
  - *threadDim*.? gives the size of each block in each dimension.
  - *blockIdx*.? gives the indices of the thread's block within the grid.
  - *threadIdx*.? gives the indices of the thread within its block.
- For dimensions which are absent:
  - *blockDim* or *threadDim* will be 1.
  - *blockIdx* or *threadIdx* will be 0.

### Threads and Blocks: Where are We?

- Note the constraints:

```
0 ≤ blockDim.x < blockDim.x
0 ≤ blockIdx.x < blockDim.x
0 ≤ blockDim.y < blockDim.y
0 ≤ blockDim.z < blockDim.z
0 ≤ threadIdx.x < blockDim.x
0 ≤ threadIdx.y < blockDim.y
0 ≤ threadIdx.z < blockDim.z
```
- Because the size of blocks are limited, it is common to use code such as:

```
uint my_idx = blockDim.x*blockIdx.x + threadIdx.x;
```

to combine the block and thread indices into a single index.

### Bounds checking: launching kernels

- Consider executing *kernelFun* on an array of *n* elements.
- Because *n* might be large, we'll use *n/256* blocks of 256 threads.
  - **THINK**: what if *n* is not a multiple of 256?
  - We'll round up to make sure we have enough threads.
- The kernel launch looks like:

```
kernelFun<<ceil(n/256.0), 256>>>(n, myArray);
```

  - Why divide by 256.0 instead of 256?
  - Why use *ceil*?

### Bounds checking: in the kernel

- The kernel launch looks like:

```
kernelFun<<ceil(n/256.0), 256>>>(n, myArray);
```
- **THINK**: what if *n* is not a multiple of 256?
  - We'll launch more than *n* threads?
  - For example, if *n*=1000, then we'll launch 4 blocks of 256 threads. A total of 1024 threads.
  - What will the last 24 threads do?
- Add a test:

```
uint my_idx = blockDim.x*blockIdx.x + threadIdx.x;
if(my_idx < n) {
 ...
}
```

### SMs, SPs and Warps (oh my!)

- Each *streaming multiprocessor* (SM) has multiple *streaming processors* (SPs) and can be responsible for multiple groups of 32 threads called *warps*.
  - From the *New Oxford American Dictionary*: (the) "warp" is "the threads on a loom over and under which other threads (the weft) are passed to make cloth"
- Details, details...
  - These concepts are not part of the CUDA platform and API: Code is written in terms of a grid of blocks of threads.
  - You can write correct code without thinking about these details.
  - If you want to write fast code, you must take them into account.
  - The block vs grid structure exposes these details if you want to take advantage of them.

### SMs, SPs and Warps: What are They?

- Each streaming multiprocessor (SM) in the GPU executes threads in SIMD fashion.
  - All threads in a block are assigned to the same SM.
  - Each SM has a single (or small number of?) instruction fetch unit(s) and a larger number of execution units.
- Each SM has multiple streaming processors (SPs) that actually execute an instruction.
  - The SPs are specialized: ALUs, load / store, special function units.
  - A single SP can perform a single operation on a small set of threads.
- A warp is a collection of 32 threads that execute together on the same SP.

### SMs, SPs and Warps: Why do We Care?

- Fill your warps: Ensure the number of threads in a block is a multiple of the warp size to avoid idle hardware.
- Have lots of warps: If one warp is waiting on a long latency operation, the SM can find another warp to execute.
  - Provides *latency tolerance* or *latency hiding*.
- Watch out for hardware limits (per SM).
  - Maximum number of resident blocks (8 in 2.x, 32 in 6.x).
  - Maximum number of resident warps (48 in 2.x, 24 thereafter).
  - Maximum number of resident threads (1536 in 2.x, 2048 thereafter).
  - Exceeding these limits will not crash the system, but will result in slower execution.
- Watch out for thread divergence.
  - If different threads in the same warp are following different code paths, all possible paths will be executed sequentially and those threads not on the current path will be idle.
  - Execution is still correct, but much slower.

### A Warped Example: Reduce (part 1)

- Consider a reduce of an array *data* containing *n* elements using *n/2* threads (assume *n* is power of 2).
- Simple code:

```
for(int stride = 1; stride < n; stride += stride) {
 if((myIdx & (stride-1)) == 0)
 data[2*myIdx] += data[2*myIdx + stride];
}
__syncthreads();
```
- The *\_\_syncthreads()* call ensures that every thread has completed an iteration of the loop before any thread starts the next iteration.
  - More discussion on [slide 18](#).

### A Warped Example: Reduce (part 2)

- Consider *n* == 16
  - First iteration, for *i* in 0, ..., 7:

```
data[2*i] += data[2*i+1]
```

Now, all the even indexed elements have their sum with their odd counterpart.
  - Second iteration, for *i* in 0, 2, 4, 6:

```
data[2*i] += data[2*i+2]
```

All elements with indices that are multiples of four, have their sum with the next three elements.
  - Third iteration leads with *data[0]* and *data[8]* holding sums for their halves of the array.
  - The fourth iteration puts the complete sum into *data[0]*.
- There are at most 8 threads working, so everything fits within a single warp.

### A Warped Example: Reduce (part 3)

- What if *n*=1024?
  - We have 512 threads: 16 warps of 32 threads.
  - In the first iteration all threads are active.
  - In the next iteration each warp has 16 active threads, so the GPU has to execute the code for all 16 warps even though half the threads do nothing.
  - In subsequent iterations, the warps are more and more poorly utilized.
- This solution is correct, but much of the parallel hardware will sit idle much of the time.
- We would like to pack the busy threads into the minimum number of warps.

### Warp Speed!

- ```
for(int stride = n/2; stride > 0; stride >>= 1) {
    if(myIdx < stride)
        data[myIdx] += data[myIdx] + stride;
    __syncthreads();
}
```
- Consider *n* == 1024 again.
 - In the first iteration, there are 16 active warps – all threads in each warp are busy.
 - In the second iteration, there are 8 active warps – all threads in each active warp are busy.
 - Similarly, for the 3rd through 5th iterations
 - The number of active warps decreases after each iteration, but all threads in each active warp are busy.
 - The inactive warps have no pending instructions, so they will not be scheduled and will not occupy processing resources.

Synchronization

- The reduce example used *__syncthreads()*: all the threads in the block must execute this statement before any can continue beyond it.
 - Be very careful about thread divergence: All threads in the block must meet at the **same** barrier.
 - That means the **same** line of code.
 - In loops, that means the **same** iteration.
 - Executing different *__syncthreads()* commands will cause the kernel to hang.
- Also, *__syncthreads()* only synchronizes between threads within a single block.
 - Note that threads within a warp already stay synchronized because they are executed together.
 - The only way to synchronize between threads in different blocks is to finish the kernel and launch another.
- We'll cover synchronization in more detail later.

Preview

- March 10: CUDA Threads, Part 2
- March 13: CUDA Memory
 - Reading [Kirk & Hwu Ch. 4](#)
- March 15: CUDA Memory examples
- March 17: CUDA Performance
 - Reading [Kirk & Hwu Ch. 5](#)
- March 20: Matrix multiplication with CUDA, Part 1
- March 22: Matrix multiplication with CUDA, Part 2
- March 24 – April 3: Other Topics
 - more parallel algorithms, e.g. dynamic programming?
 - reasoning about concurrency, e.g. termination detection
 - other paradigms, e.g. Scala and futures?
- April 5: Party: 50th Anniversary of Amdahl's Law

Review

- In CUDA, what is a grid, a block, and thread?
- Why does CUDA allow millions of thread blocks but only 1024 threads per block?
- How does a programmer specify the number of blocks and number of threads when launching a CUDA kernel?
- How does a thread determine its position within the grid?
- Why do threads need to check their indices against array bounds?
- What is a warp? Why does it matter?


CUDA: Memory

Mark Greenstreet


CpSc 418 – March 13 & 15, 2017

- [Architecture Snapshot](#)
- [Registers](#)
- [Shared Memory](#)
- [Global Memory](#)
- [Other Memory: texture memory, constant memory, caches](#)
- [Summary, preview, review, slide chart](#)

First, GPU Architecture Review



A Streaming Multiprocessor (SM)



- Each of the pipelines is an SP (streaming processor)
- Lots of deep pipelines.
- Lots of threads: when we encounter an architectural challenge:
 - Raising throughput is **easy**, lowering latency is **hard**.
 - Solve problems by increasing latency and adding threads.
 - Make the programmer deal with it.


Why do we need a memory hierarchy

- ```
__global__ void saxpy(uint n, float a, float *x, float *y) {
 uint myId = blockDim.x*blockIdx.x + threadIdx.x;
 if(myId < n)
 y[myId] = a*x[myId] + y[myId];
}
```
- A GPU with 3584 SPs, and a 1.6GHz clock rate (see [slide 2](#) can perform over 5700 single-precision GFlops.
    - With a main memory bandwidth of 484 GB/s, and 4 bytes per float, a CUDA kernel needs to perform  $\frac{5700 \times 4}{484} \approx 48$  floating point operations per memory read or write.
    - Otherwise, memory bandwidth becomes the bottleneck.
  - Registers and shared memory let us use a value many times without going to the off-chip, GDDR memory.
    - But, we need to program carefully to make this work.
  - Is *saxpy* a good candidate for GPU execution?


### Matrix Multiplication and Memory

- ```
for(int i = 0; i < N; i++)
    for(int j = 0; j < N; j++)
        for(int k = 0; k < i; k++)
            c[i,j] += a[i,k]*b[k,j];
```
- Focus on the innermost loop: *for (k...)*
 - Why?
 - How many floating point operations per iteration?
 - How many memory reads?
 - How many memory writes?
 - What is the "Compute-to-Global-Memory-Access" ratio (CGMA)?

Registers




More Registers



Registers and Memory Bandwidth

- The GPU on [slide 2](#) has 28 SMs, each with 128 SPs.
- Each SP has access to a register file.
 - I'll guess two register reads and one write per clock cycle, per SP.
 - I'll assume 4-byte registers.
 - We get
$$28SM \times 128 \frac{SP}{SM} \times 3 \frac{RW}{SP \times cycle} \times 1.6 \times 10^9 \frac{cycle}{sec} \times 4 \frac{Byte}{RW} = 68813 \frac{GByte}{sec}$$
 - 142 times faster than main memory bandwidth!
 - ymmv: the GPUs in the linXX boxes are older.

Registers and Thread Scheduling




```
for(int i = 0; i < M; i += 2) {
    for(int j = 0; j < N; j += 2) {
        sum0 = sum01 = sum10 = sum11 = 0.0;
        for(int k = 0; k < L; k += 2) {
            a00 = a[i,k]; a01 = a[i,k+1];
            a10 = a[i+1,k]; a11 = a[i+1,k+1];
            b00 = b[k,j]; b01 = b[k,j+1];
            b10 = b[k+1,j]; b11 = b[k+1,j+1];
            sum0 += a00*b00 + a01*b10;
            sum01 += a00*b01 + a01*b11;
            sum10 += a10*b00 + a11*b10;
            sum11 += a10*b01 + a11*b11;
        }
        c[i,j] = sum00; c[i,j+1] = sum01;
        c[i+1,j] = sum10; c[i+1,j+1] = sum11;
    }
}
```

- use a register to accumulate $c[i,j]$
- hold blocks of each matrix in main memory.
- can use each value loaded from $a[i,k]$ or $b[k,j]$ three or four times.
- What is the CGMA for the example above?

- On-chip, one bank per SP.
- Banks are interleaved by:
 - Early CUDA GPUs: 4-byte word
 - Later GPUs: programmer configurable 4-byte or 8-byte words
 - Why?
- Shared memory is a limited resource: 48KBytes to 96KBytes/SM.
 - Each SM has more registers than shared-memory.
 - Shared memory demands limit how many blocks can execute concurrently on a SM.

```
__shared__ float v[512];
__device__ float f(float x) {
    return (5/2)*(x*x*x - x);
}
__device__ void compute_and_reduce(uint n, uint m, float *x) {
    uint myid = threadIdx.x;
    if(myid < n) {
        float y = x[myid];
        for(int i = 0; i < m; i++)
            y = f(y);
        v[myid] = y;
        if(myid == 0) {
            float sum = 0.0;
            for(int i = 0; i < n; i++)
                sum += v[i];
            *x = sum;
        }
    }
}
```

See notes on the next slide.

- We calculate $E^n(x[\text{myId}])$ locally using the register variable y .
- For the reduce, we use the **shared** array v
 - This avoids the penalty of off-chip, global memory access for each step of the reduce.
 - All threads in a warp can access shared memory on the same chip.
 - We can have multiple blocks running on multiple SMs
 - Each SM has its own shared memory.
 - Blocks running on different SMs in parallel can **all** access their shared memories in parallel.
 - But, threads in one block do not share shared-memory with threads in other blocks.
 - To perform a reduce across blocks:
 - Each block writes its subtotal to the **global** memory.
 - The results from the blocks are combined on the host CPU or by launching a new kernel.
- At the end, we copy our value from shared memory to the global memory so the CPU or a subsequent kernel can access it.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	10/24
------------------	--------------	-----------------------------	-------

Shared Memory: Collisions

- When one thread in a warp accesses shared memory, **all** active threads in the warp access shared memory.
- If each thread accesses a different bank, then all accesses are performed in a single cycle.
 - otherwise, the load or store can take multiple cycles.
 - Multiple accesses to the same bank are called **collisions**.
 - The **worst-case** occurs when **all threads access the same bank**.
- The programmer needs to think about the index calculations to avoid collisions.
 - When programming GPUs, the programmer needs to think about index calculations a lot.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	11/24
------------------	--------------	-----------------------------	-------

Shared Memory Example: Matrix Multiply

Running example from the textbook: $C = AB$

- Each thread-block loads a 16×16 block from A and B .
 - The threads to these loads "cooperatively":
 - Read $A_{i,K}$ and $B_{K,j}$ from global memory with "coalesced" loads.
 - Write these blocks to shared-memory in a way that avoids bank conflicts.
- Compute: $C_{i,j} += A_{i,K} B_{K,j}$.
 - This takes $16^2 = 4096$ fused multiply-adds.
 - Loading $A_{i,K}$ fetches $16^2 = 256$ floats from global memory. Likewise for $B_{K,j}$. Total of 512 floats fetched.
 - CGMA = $4096/512 = 8$.
- Note: the L2 cache may help here: A and B are read-only.
 - Need to try more experiments.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	12/24
------------------	--------------	-----------------------------	-------

Global Memory

- Off-chip DRAM
 - GDDR supports higher-bandwidth than regular DDR.
 - A GPU can have multiple memory interfaces.
 - Total bandwidth 80 to 484 Gbytes/sec
- Memory accesses can be a big bottleneck.
 - CGMA: compute to global memory access ratio

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	13/24
------------------	--------------	-----------------------------	-------

Now for a word about DRAM

The memory that you plug into your computer is mounted on DIMMs (dual-inline memory modules).

- A DIMM typically has 16 or 18 chips
- E.g. each chip of an 8Gbyte DIMM holds 512Mbytes = 4Gbits.
- Each chip consists of many "tiles".
 - a typical chip has 1Mbit/tile
 - that's 4096 tiles for a 4Gbit chip.
- Each tile is an array of capacitors.
 - each capacitor holds 1 bit.
 - a typical file could have 1024 rows and 1024 columns.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	14/24
------------------	--------------	-----------------------------	-------

Writing and reading DRAM

- Writing: easy
 - drive all 1024 column-lines to the values you want to write.
 - open up all the valves for one row.
 - the drinking cups for each column in that row get filled or emptied.
 - note: you end up writing **every** column in the row; so writes are often preceded by reads.
- Reading: hard
 - drive all 1024 column-lines to "half-way", and let them "float".
 - open up all the valves for one row.
 - if the level in the pipe goes up a tiny amount, that cup held a 1.
 - if the level in the pipe goes down a tiny amount, that cup held a 0.
 - it's a delicate measurement – it takes time to set it up.
 - This is why DRAM is slow.
- But: we just read 1024 bits, from each chip of the DIMM.
 - That's 16Kbits = 2Kbytes total.
 - Conclusion: DRAM has awful latency, but we can get very good bandwidth.
 - The bandwidth bottleneck is the wires from the DIMM to the CPU or GPU.
 - But I'm pretty sure that I can't let me give a lecture on transmission lines, phase-locked loops, equalizers, and all the other cool stuff in the DDR (or GDDR) interface.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	15/24
------------------	--------------	-----------------------------	-------

GPUs meet DRAM

- DRAM summary: terrible latency (60-200ns or more), fairly bandwidth.
- The GPU lets the program take advantage of high bandwidth.
 - If the 32 loads from a warp access 32 consecutive memory locations.
 - The GPU does **one** GDDR access, and it transfers a large block of data.
 - The same optimization is applied to stores, and to loads from the on-chip caches.
- In CUDA-speak, if the loads from a warp access consecutive locations, we say that the memory accesses are **coalesced**.
- It's a big deal to make sure that your memory accesses are coalesced.
 - Note that the memory optimizations are exposed to the programmer.
 - You can get the performance by considering the memory model.
 - But, it's not automatic.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	16/24
------------------	--------------	-----------------------------	-------

Example: Matrix Multiplication

- In C, matrices are usually stored in row-major order.
 - $A[i,k]$ and $A[i,k+1]$ are at adjacent locations, but
 - $B[k,j]$ and $B[k+1,j]$ are N words apart (for $N \times N$ matrices).
- For matrix multiplication, accesses to A are naturally coalesced, but accesses to B .
- The optimized code loads a block of B into shared memory.
 - This allows accesses to be coalesced.
 - But we need to be careful about how we store the data in the shared memory to avoid bank conflicts.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	17/24
------------------	--------------	-----------------------------	-------

Other Memory

- Constant memory: cached, read-only access of global memory.
- Texture memory: global memory with special access operations.
- L1 and L2 caches: only for memory reads.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	18/24
------------------	--------------	-----------------------------	-------

Summary

- GPUs can have thousands of execution units, but only a few off-chip memory interfaces.
 - This means that the GPU can perform 10-50 floating point operations for every memory read or write.
 - Arithmetic operations are very cheap compared with memory operations
- To mitigate the off-chip memory bottleneck
 - GPUs have, limited on-chip memory
 - Registers and the per-block, shared-memory will be our main concerns in this class.
- Moving data between different kinds of storage is the programmer's responsibility.
 - The programmer explicitly declares variables to be stored in shared memory.
 - The programmer needs to be aware of the per-thread register usage to achieve good SM utilization.
 - The only way to communicate between thread blocks is to write to global memory, and the kernel, and start a new kernel (ouch)

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	19/24
------------------	--------------	-----------------------------	-------

Preview

March 15: CUDA Memory: examples

March 17: CUDA Performance

Reading: [Kirk & Huw](#) 3rd ed., Ch. 5 (Ch. 6 in 2nd ed.)

Mini Assignment: Mini Assignment 5 due at 10am.

March 20: Matrix multiplication, Part 1

March 22: Matrix multiplication, Part 2

March 24: Complete CUDA

March 27 – April 3: **this may change**

March 27: Using Parallel Libraries

March 29 – April 3: Verification of and Parallel Programs

April 5: Party: 50th Anniversary of AMDahl's Law


Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	20/24
------------------	--------------	-----------------------------	-------

Review

- What is CGMA?
- On slide 15 we computed the CGMA for matrix-multiplication using 16×16 blocks of the A , B , and C matrices.
 - How many such thread-blocks can execute concurrently on an SM with 48Kbytes of memory?
 - How does the CGMA change if we use 32×32 blocks?
 - If we use the larger matrix-blocks, how many thread blocks can execute concurrently on an SM with 48Kbytes of memory?
 - If we use the larger matrix-blocks, how many thread blocks can execute concurrently on an SM with 96Kbytes of memory?
- What are bank conflicts?
- How can increasing the number of registers used by a thread improve performance?
- How can increasing the number of registers used by a thread degrade performance?
- What is a "coalesced memory access"?

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	21/24
------------------	--------------	-----------------------------	-------

Beaware the Tides of March



Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	22/24
------------------	--------------	-----------------------------	-------

CUDA: Performance Considerations

Mark Greenstreet & Ian M. Mitchell

CPSC 418 – March 17, 2017

- [Thread Divergence](#)
- [Floating Point Foibles](#)
- [Memory Accesses](#)
- [Occupancy](#)
- [Granularity](#)

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	23/24
------------------	--------------	-----------------------------	-------

Thread Divergence

- If threads in a warp are following different code paths, execution will be much slower.
- See "A Warped Example" from [March 13 slides](#).

Try to minimize thread divergence within warps.

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	24/24
------------------	--------------	-----------------------------	-------

Remarks about floating point

- When working on my solution to [last year's HW3, Q1](#),
 - I first wrote:

```
x = alpha*(1.0 - x);
```
 - and the performance was disappointing.
 - After many frustrating attempts to track down the problem, I added one, little `l`:

```
x = alpha*(1.0f - x);
```
 - and my code ran 5.5x faster.
- What happened?

Mark Greenstreet	CUDA: Memory	CS 418 – Mar. 13 & 15, 2017	25/24
------------------	--------------	-----------------------------	-------

Floats, doubles, and GPUs

- GPUs are optimized for single-precision floating point arithmetic.
- For the GeForce GTX 550 Ti, double precision arithmetic is way slower than single precision.
- In C, `1.0` is a **double precision** constant, and `1.0f` is single precision.
- When I wrote `x = alpha*x*(1.0-x)`, the compiler generated code that:
 - computes the product `alpha*x`.
 - both operands are single precision.
 - the computation is done using single precision arithmetic.
 - computes the difference `1.0-x`.
 - `1.0` is double precision, `x` is single precision.
 - the computation is done using double precision arithmetic
 - and the result is double precision.
 - computes the product `alpha*x*(1.0-x)`.
 - the computation is done using double precision arithmetic
 - and the result is double precision.
- When I wrote `x = alpha*x*(1.0f-x)`, everything stays in single-precision, and it's **much** faster.

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	1/22
------------------------	----------------------------------	---------------------------	------

Fused multiply adds

- Calculating $ax + b$ is very common
 - Example: dot product.
- The multiplier hardware is just a pipeline of adders.
 - When multiplying $a \times x$, the hardware can start the pipeline from b instead of from 0.
 - We get the sum for "free".
 - This is called a **fused** multiply-add.
- The marketing people like to count the fused multiply-add as **two** floating point operations.
 - This helps make some performance claims make sense.
- For the obsessive compulsive:
 - Rounding with a fused-multiply add can be slightly different than when doing two, separate operations.
 - Compilers usually let the users specify "strict" floating point (no fusing) or "fast" floating point (with fusing).
 - `nvcc` uses fused multiply-add unless you give it an option not to.

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	2/22
------------------------	----------------------------------	---------------------------	------

Memory Access

Memory is slow.

- It takes a long time to identify, access and deliver data to/from a memory address.
- Delivery rate is limited by clock rate of the memory interface.

How can we get enough data to/from our thousands of threads?

Parallelism!

- Retrieve lots of data at once.
- Use multiple memory interfaces.
- Build with lots of independent memory components.

All standard techniques in the CPU world, but

- CPU design philosophy: Try to achieve maximum performance even if the programmer uses the RAM model.
- GPU design philosophy: Expose (almost) everything and let the programmer figure it out.

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	3/22
------------------------	----------------------------------	---------------------------	------

Memory System Parallelism 1: Get Lots of Data

Memory is addressed per byte, but you retrieve a bunch of (sequential) bytes at once.

- GDDR5 DRAM: 32-bit bus per chip and transfers are in 16 word bursts (so 64 bytes per access per chip).
- GPU global memory (from GDDR DRAMs): Accessed by 32-, 64- or 128-byte transactions.
 - Transactions must be "naturally" aligned: First address must be a multiple of the transaction size.
 - CC 2.x: L1 cache (1 per SM) serviced by 128-byte transactions, L2 cache (shared by SMs) by 32-byte transactions.
 - CC 6.x: Same as 2.x, but L1 cache rules are complicated.
- GPU shared memory (on-chip SRAM): Access in 32-bit words.

Amortize addressing overhead and thereby increase bandwidth.

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	4/22
------------------------	----------------------------------	---------------------------	------

Memory System Parallelism 2: Multiple Interfaces

If one memory component cannot give you enough bandwidth, use a bunch (see [March 13 slides](#)).

- Global memory: K&H(3) calls these "channels" (March 13 slide 2).
- Shared memory: Mark called these "banks" (March 13 slide 11) and Nvidia documentation does too.
 - Do not confuse with K&H(3)'s "banks" (see next slide).

Consecutive chunks are placed into components in a round-robin fashion, where "chunk" means

- 32-bytes (64 more recently?) in global memory.
- 32- or 64-bits in shared memory.

Separate subsystems can all provide data at their native rate and thereby increase bandwidth.

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	5/22
------------------------	----------------------------------	---------------------------	------

Memory System Parallelism 3: Independent Memory Components

Even after memory address is delivered, it still takes time for the DRAM to return the data.

- Rather than let the memory bus sit idle while waiting, pipeline a bunch of memory requests to different memory components.
 - K&H(3) calls these "banks".
 - Mark called these "tiles".
- Consecutive memory chunks are assigned first to channels / banks (see previous slide).
 - These subsystems allow concurrent access because they have independent communication lines.
- Then assign next set of consecutive chunks to banks / tiles.
 - These subsystems allow sequential but pipelined access because they share communication lines

Pipelining increases throughput (although latency remains).

- Only relevant for global memory.
- Shared memory achieves dramatically lower latency with SRAM.

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	6/22
------------------------	----------------------------------	---------------------------	------

Is This on the Final?

No (sort of): This is not a course on memory system design and implementation.

- Mark and I are far from experts.
- Details depend on the particular GPU chip and card, and change regularly.
- Program correctness does not depend on getting it right.

Yes (sort of): By following some simple rules, speed can be improved dramatically.

- Design global memory access pattern to allow accesses from threads in the same warp to be **coalesced** into a single memory transaction.
- Design memory access pattern to avoid channel / bank / tile **collisions**.

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	7/22
------------------------	----------------------------------	---------------------------	------

Implications for Shared Memory

See [CUDA Toolkit Documentation C Programming Guide Figure 17 and Figure 18](#).

- Consider shared memory address bits:
 - 48KB / thread block requires 16 bits to address.
 - Bottom two bits specify the byte within a 32-bit word of data.
 - Next five bits specify which of 32 banks.
 - Top nine bits specify which word within the bank.
- Key takeaway: If two threads in a warp access a memory location in the same bank (same middle five bits of address):
 - If threads access the same location (same top nine bits), then broadcast (on read) or one value wins (on write).
 - If threads access different location, access is serialized (slower but still correct).

Greenstreet & Mitchell	CUDA: Performance Considerations	CPSC 418 – March 17, 2017	8/22
------------------------	----------------------------------	---------------------------	------

Implications for Global Memory

Try to get memory access addresses from threads in a warp to be very close together.

- Accesses to consecutive (or nearly so) addresses are coalesced into a single transaction on the off-chip memory bus.
 - You should already be doing this for your CPU designs so that your caches can take advantage of spatial locality.
- Best coalescing occurs when the set of addresses is naturally aligned.
 - For two and higher dimensional arrays, that may mean padding thread block and array width allocation in memory to be a multiple of the warp size.
- Possibility of channel / bank collisions would argue for avoiding addresses with the same "middle" bits.
 - I could not find Nvidia documentation of these details.
 - How do caches interact with channels / banks?

Comments from Mark?


- Occupancy: how many warps are available for the SM
 - Why we care: the SP pipelines have long latencies.
 - The CUDA approach is to run lots of threads simultaneously to keep the pipelines busy.
- Limits to occupancy
 - How many blocks per SM.
 - How SM can simultaneously execute up to 1536 threads (48 warps).
 - How many threads per block.
 - How many registers per thread.
- Figuring it out
 - `nvcc -O3 -c --ptxas-options -v examples.cu`
 - The nVidia occupancy calculator: [CUDA.Occupancy.calculator.xls](#)
 - But we can do it manually?

- Different GPUs at level CUDA 2.1 have differing numbers of SMS.
 - But the SMS all look the same, even for different GPUs.
- CUDA 2.1 SMS
 - An SM has warps of 32 threads
 - An SM can simultaneously execute up to 1536 threads (48 warps).
 - An SM has 32K (2¹⁵) 32-bit registers (128K/bytes, 1K registers/SP).
 - An SM has 48K bytes of shared memory.
 - An SM can simultaneously execute up to 8 blocks.
 - Each block can have up to 1024 threads.

- When designing a new generation of GPUs, the GPU architects run lots of simulations to estimate the performance for various choices of the architectural parameters.
- For example, if more warps are allowed in the scheduling pool
 - The SM will have useful instructions to dispatch more often → better performance.
 - BUT the on-chip circuitry to hold and manage the scheduling pool will be larger.
 - This means instruction scheduling will be slower → a longer clock period.
 - Instruction scheduling will use more power → a longer clock period, or fewer SMS, or more expensive chip cooling.
 - The real-estate on the chip could have been used for something else. Is this the best use of that area.
 - Note that CUDA 5 made the increase to 64 warps/SM.
- Architects explore these trade-offs to optimize performance for graphics applications, the main source of revenue.
- Architects are also risk-averse: make the chip as much like the last one that worked as you can.
- These hard-wired constraints have a large impact on program performance.

- A SM can have simultaneously execute most 8 blocks.
- All blocks have the same number of threads.
- Thus, a SM can execute at most
$$\min \left(8, \left\lfloor \frac{1536}{\text{threadsPerBlock}} \right\rfloor \right)$$
blocks.
- The ratio of the number of threads executing to the maximum possible is called the "thread occupancy":
$$\text{threadOccupancy} \leq \frac{1536}{\min \left(8, \left\lfloor \frac{1536}{\text{threadsPerBlock}} \right\rfloor \right)} \frac{\text{threadsPerBlock}}{1536}$$

SMS, blocks, and threads – the plot



- I get 100% occupancy when $\text{threadsPerBlock} \in \{192, 384, 768\}$, but the CUDA calculator doesn't.
- I'll have to try some experiments – stay tuned.
- This assumes the grid had enough blocks to keep the SMS busy.
- A grid with a single block will have poor performance.

SMS, threads, and registers

- Each SM has 32K registers – that's 1K registers per SP.
- This is another constraint:
$$\text{nbkbs} \leq \frac{1024}{\text{registersPerThread}}$$
- An SM can run 48 warps simultaneously
 - But only if each warp uses at most 21 registers.

Hitting the register constraint

- What if each thread uses 22 registers?
- $22 \times 48 = 1056 > 1024 \rightarrow$ can't run 48 warps.
 - $\left\lfloor \frac{1024}{22} \right\rfloor = \lfloor 46.54 \rfloor = 46$.
 - Can we run 46 warps?
 - One block with 46 warps would have $46 \times 32 = 1472 > 1024$ threads. Not allowed.
 - Two block with 23 warps each would each have 736 threads. That should work.
 - But, the plot with the occupancy calculator only shows warp counts that are multiples of 8.
 - Have I overlooked another architectural constraint?
 - probably
 - Let's assume that with 23 registers per thread, the SM can run at most 40 warps simultaneously.
 - Then either each thread must have enough instruction-level parallelism to keep the SPs busy.
 - Or, we'll see a drop in performance.

How many registers does my thread use?

- use the `--ptxas-options -v` option for `nvcc`

```
nvcc-pxt-as-options -v -O3 -c examples.cu
ptxas info: 0 bytes gmem
ptxas info: Compiling entry function 'shmem_2j1ip' for 'sm_20'
ptxas info: Function properties for 'shmem_2j1ip'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info: Used 17 registers, 4096 bytes smem, 56 bytes cmem[0]
ptxas info: Compiling entry function 'shmem_4j1ip' for 'sm_20'
ptxas info: Function properties for 'shmem_4j1ip'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info: Used 14 registers, 4096 bytes smem, 56 bytes cmem[0]
```
- Translation:
 - kernel `shmem_2` uses 17 registers per thread.
 - kernel `shmem_4` uses 14 registers per thread.
 - both kernels use 4024 bytes of shared memory per block.
 - neither kernel spills registers to global memory (good).

Granularity

- How much work should a kernel do?
- Do more work within a kernel: Launching each kernel takes time.
 - Do less work within a kernel: New kernels allow for changes in block and grid size, and ensure synchronization between threads even in different blocks.
 - Either way: Minimize movement of data to and from the host.
- How much work should a thread do?
- Do more work in a single thread: Fewer chances for memory collisions, easier synchronization, less register contention.
 - Do less work in a single thread: More potential parallelism, more chance for latency hiding.
 - Tradeoff will depend on GPU resources, typically SM block, thread and register limits.

Review

March 20: Matrix multiplication, Part 1
March 22: Matrix multiplication, Part 2
March 24: Complete CUDA
March 27 – April 3: this may change
March 27: Using Parallel Libraries
March 29 – April 3: Verification of And/Parallel Programs
April 5: Party: 50 th Anniversary of Amdahl's Law

CUDA: Performance Considerations

Mark Greenstreet & Ian M. Mitchell

CPSC 418 – March 17, 2017

- [Thread Divergence](#)
- [Floating Point Foibles](#)
- [Memory Accesses](#)
- [Occupancy](#)
- [Granularity](#)

Unless otherwise noted or stated, these slides are copyright 2017 by Mark Greenstreet & Ian M. Mitchell and are made available under the terms of the Creative Commons Attribution 4.0 International license (<https://creativecommons.org/licenses/by/4.0/>)

Thread Divergence

- If threads in a warp are following different code paths, execution will be much slower.
- See "A Warped Example" from [March 13 slides](#).

Try to minimize thread divergence within warps.

Remarks about floating point

- When working on my solution to [last year's HW3](#), Q1,
 - I first wrote:

```
x = alpha*x*(1.0 - x);
```
 - and the performance was disappointing.
 - After many frustrating attempts to track down the problem, I added one, little fix:

```
x = alpha*x*(1.0f - x);
```
 - and my code ran 5.5x faster.
- What happened?

Floats, doubles, and GPUs

- GPUs are optimized for single-precision floating point arithmetic.
- For the GeForce GTX 550 Ti, double precision arithmetic is way slower than single precision.
- In C, `1.0` is a **double** **precision** constant, and `1.0f` is single precision.
- When I wrote `x = alpha*x*(1.0-x)`, the compiler generated code that:
 - computes the product `alpha*x`.
 - both operands are single precision.
 - the computation is done using single precision arithmetic.
 - computes the difference `1.0-x`.
 - 1.0 is double precision, x is single precision.
 - the computation is done using double precision arithmetic
 - and the result is double precision.
 - computes the product `alpha*x*(1.0-x)`.
 - the computation is done using double precision arithmetic
 - and the result is double precision.
- When I wrote `x = alpha*x*(1.0f-x)`, everything stays in single-precision, and it's **much** faster.

Fused multiply adds

- Calculating `ax + b` is very common
 - Example: dot product.
- The multiplier hardware is just a pipeline of adders.
 - When multiplying `a*x`, the hardware can start the pipeline from `b` instead of from 0.
 - We get the sum for "free".
 - This is called a **fused** multiply-add.
- The marketing people like to count the fused multiply-add as **two** floating point operations.
 - This helps make some performance claims make sense.
- For the obsessive compulsive:
 - Rounding with a fused-multiply add can be slightly different than when doing two, separate operations.
 - Compilers usually let the users specify "strict" floating point (no fusing) or "fast" floating point (with fusing).
 - `nvcc` uses fused multiply add unless you give it an option not to.

Memory Access

- Memory is slow.
- It takes a long time to identify, access and deliver data to/from a memory address.
 - Delivery rate is limited by clock rate of the memory interface.
- How can we get enough data to/from our thousands of threads? **Parallelism!**
- Retrieve lots of data at once.
 - Use multiple memory interfaces.
 - Build with lots of independent memory components.
- All standard techniques in the CPU world, but
- CPU design philosophy: Try to achieve maximum performance even if the programmer uses the RAM model.
 - GPU design philosophy: Expose (almost) everything and let the programmer figure it out.

Memory System Parallelism 1: Get Lots of Data

- Memory is addressed per byte, but you retrieve a bunch of (sequential) bytes at once.
- GDDR5 DRAM: 32-bit bus per chip and transfers are in 16 word bursts (so 64 bytes per access per chip).
 - GPU global memory (from GDDR DRAMs): Accessed by 32-, 64- or 128-byte transactions.
 - Transactions must be "naturally" aligned: First address must be a multiple of the transaction size.
 - C2.x: L1 cache (1 per SM) serviced by 128-byte transactions, L2 cache (shared by SMS) by 32-byte transactions.
 - C6.x: Same as 2.x, but L1 cache rules are complicated.
 - GPU shared memory (on-chip SRAM): Access in 32-bit words.
- Amortize addressing overhead and thereby increase bandwidth.

Memory System Parallelism 2: Multiple Interfaces

- If one memory component cannot give you enough bandwidth, use a bunch (see [March 13 slides](#)).
- Global memory: K&H(3) calls these "channels" (March 13 slide 2).
 - Shared memory: Mark called these "banks" (March 13 slide 11) and nVidia documentation does too.
 - Do not confuse with K&H(3) "banks" (see next slide).
- Consecutive chunks are placed into components in a round-robin fashion, where "chunk" means
- 32-bytes (64 more recently?) in global memory.
 - 32- or 64-bits in shared memory.
- Separate subsystems can all provide data at their native rate and thereby increase bandwidth.

Memory System Parallelism 3: Independent Memory Components

- Even after memory address is delivered, it still takes time for the DRAM to return the data.
- Rather than let the memory bus sit idle while waiting, pipeline a bunch of memory requests to different memory components.
 - K&H(3) calls these "banks".
 - Mark called these "tiles".
 - Consecutive memory chunks are assigned first to channels / banks (see previous slides).
 - These subsystems allow concurrent access because they have independent communication lines.
 - Then assign next set of consecutive chunks to banks / tiles.
 - These subsystems allow sequential but pipelined access because they share communication lines
- Pipelining increases throughput (although latency remains).
- Only relevant for global memory.
 - Shared memory achieves dramatically lower latency with SRAM.

Is This on the Final?

- No (sort of). This is not a course on memory system design and implementation.
- Mark and I are far from experts.
 - Details depend on the particular GPU chip and card, and change regularly.
 - Program correctness does not depend on getting it right.

- Yes (sort of): By following some simple rules, speed can be improved dramatically.
- Design global memory access pattern to allow accesses from threads in the same warp to be **coalesced** into a single memory transaction.
 - Design memory access pattern to avoid channel / bank / tile **collisions**.

Implications for Shared Memory

- See [CUDA Toolkit Documentation C Programming Guide Figure 17](#) and [Figure 18](#).
- Consider shared memory address bits:
 - 48KB / thread block requires 16 bits to address.
 - Bottom two bits specify the byte within a 32-bit word of data.
 - Next five bits specify which of 32 banks.
 - Top nine bits specify which word within the bank.
 - Key takeaway: (if two threads in a warp access a memory location in the same bank (same middle five bits of address):
 - If threads access the same location (same top nine bits), then broadcast (on read) or one value wins (on write).
 - If threads access different location, access is serialized (slower but still correct).

Implications for Global Memory

- Try to get memory access addresses from threads in a warp to be very close together.
- Accesses to consecutive (or nearly so) addresses are coalesced into a single transaction on the off-chip memory bus.
 - You should already be doing this for your CPU designs so that your caches can take advantage of spatial locality.
 - Best coalescing occurs when the set of addresses is naturally aligned.
 - For two and higher dimensional arrays, that may mean padding thread block and array width allocation in memory to be a multiple of the warp size.
 - Possibility of channel / bank collisions would argue for avoiding addresses with the same "middle" bits.
 - I could not find nVidia documentation of these details.
 - How do caches interact with channels / banks?

Comments from Mark?

SMS and Thread Occupancy

- Occupancy: how many warps are available for the SM
 - Why we care: the SP pipelines have long latencies.
 - The CUDA approach is to run lots of threads simultaneously to keep the pipelines busy.
- Limits to occupancy
 - How many blocks per SM.
 - How much shared-memory per block.
 - How many threads per block.
 - How many registers per thread.
- Figuring it out
 - `nvcc -O3 -c --ptxas-options -v examples.cu`
 - The nVidia occupancy calculator: [CUDA.Occupancy.calculator.xls](#)
 - But we can do it manually?

Occupancy with CUDA 2.1

- Different GPUs at level CUDA 2.1 have differing numbers of SMS.
 - But the SMS all look the same, even for different GPUs.
- CUDA 2.1 SMS
 - An SM has warps of 32 threads
 - An SM can simultaneously execute up to 1536 threads (48 warps).
 - An SM has 32K (2¹⁵) 32-bit registers (128K/bytes, 1K registers/SP).
 - An SM has 48K bytes of shared memory.
 - An SM can simultaneously execute up to 8 blocks.
 - Each block can have up to 1024 threads.


Why all these numbers?

- When designing a new generation of GPUs, the GPU architects run lots of simulations to estimate the performance for various choices of the architectural parameters.
- For example, if more warps are allowed in the scheduling pool
 - The SM will have useful instructions to dispatch more often → better performance.
 - BUT the on-chip circuitry to hold and manage the scheduling pool will be larger.
 - This means instruction scheduling will be slower → a longer clock period.
 - Instruction scheduling will use more power → a longer clock period, or fewer SMS, or more expensive chip cooling.
 - The real-estate on the chip could have been used for something else. Is this the best use of that area.
 - Note that CUDA 5 made the increase to 64 warps/SM.
- Architects explore these trade-offs to optimize performance for graphics applications, the main source of revenue.
- Architects are also risk-averse: make the chip as much like the last one that worked as you can.
- These hard-wired constraints have a large impact on program performance.

SMS, blocks, and threads

- A SM can have simultaneously execute most 8 blocks.
- All blocks have the same number of threads.
- Thus, a SM can execute at most
$$\min \left(8, \left\lfloor \frac{1536}{\text{threadsPerBlock}} \right\rfloor \right)$$
blocks.
- The ratio of the number of threads executing to the maximum possible is called the "thread occupancy":
$$\text{threadOccupancy} \leq \frac{1536}{\min \left(8, \left\lfloor \frac{1536}{\text{threadsPerBlock}} \right\rfloor \right)} \frac{\text{threadsPerBlock}}{1536}$$

SMS, blocks, and threads – the plot



- I get 100% occupancy when $\text{threadsPerBlock} \in \{192, 384, 768\}$, but the CUDA calculator doesn't.
- I'll have to try some experiments – stay tuned.
- This assumes the grid had enough blocks to keep the SMS busy.
- A grid with a single block will have poor performance.

SMS, threads, and registers

- Each SM has 32K registers – that's 1K registers per SP.
- This is another constraint:
$$\text{nbkbs} \leq \frac{1024}{\text{registersPerThread}}$$
- An SM can run 48 warps simultaneously
 - But only if each warp uses at most 21 registers.

Hitting the register constraint
What if each thread uses 22 registers?

- 22 * 48 = 1056 > 1024 → can't run 48 warps.
- $\lfloor \frac{1024}{22} \rfloor = \lfloor 46.54 \rfloor = 46$.
- Can we run 46 warps?
 - One block with 46 warps would have 46 * 32 = 1472 > 1024 threads. Not allowed.
 - Two block with 23 warps each would each have 736 threads. That should work.
 - But, the plot with the occupancy calculator only shows warp counts that are multiples of 8.
 - Have I overlooked another architectural constraint?
 - probably
- Let's assume that with 23 registers per thread, the SM can run at most 40 warps simultaneously.
 - Then either each thread must have enough instruction-level parallelism to keep the SPs busy.
 - Or, we'll see a drop in performance.

- use the `--ptxas-options -v` option for `nvcc`

```
nvcc -ptxas-options -v -O3 -c examples.cu
ptxas info : 0 bytes smem
ptxas info : Compiling entry function 'shmem2j1ip' for 'sm20'
ptxas info : Function properties for 'shmem2j1ip'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info : Used 17 registers, 4096 bytes smem, 56 bytes cmm[0]
ptxas info : Compiling entry function 'shmem2j1ip' for 'sm20'
ptxas info : Function properties for 'shmem2j1ip'
0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info : Used 14 registers, 4096 bytes smem, 56 bytes cmm[0]
```
- Translation:
 - kernel sh_mem_2 uses 17 registers per thread.
 - kernel sh_mem_1 uses 14 registers per thread.
 - both kernels use 4024 bytes of shared memory per block.
 - neither kernel spills registers to global memory (good).

Greenstreet & Michael
CUDA: Performance Considerations
CS 418 – March 17, 2017
19 / 22

Greenstreet & Michael
CUDA: Performance Considerations
CS 418 – March 17, 2017
20 / 22

Greenstreet & Michael
CUDA: Performance Considerations
CS 418 – March 17, 2017
21 / 22

Greenstreet & Michael
CUDA: Performance Considerations
CS 418 – March 17, 2017
16 / 18

Loop Limitations

- It takes two or three instructions per loop iteration to manage the loop:
 - One to update the loop index
 - One or two to check the loop bounds and branch.
 - If `do something` is only three or four instructions, then 40-50% of the execution time is for loop management.
- If each iteration of `do something` depends on the previous one
 - Then the long latency of the SP pipelines can limit performance.
 - Even if we have 48 warps running.

Loop Unrolling

- Have each loop iteration perform multiple copies of the loop body

```
__global__ myKernel(int m, ...) {
    for(int i = 0; i < m; i += 4) {
        do something 1
        do something 2
        do something 3
        do something 4
    }
}
```
- More "real work" for each time the loop management code is executed.
- Need to make sure that `m` is a multiple of four, or handle end-cases separately.
- Often, we need more registers.

Unrolling – the plots

This example is from [last year's HW3](#), Q1.

Bigger Ranges

__global__ myKernel(...) {
 do something
}

Unless `do something` is big, kernel launch takes most of the time.

- We can launch a big-grid
 - If we have a huge number of array elements than each need a small amount of work, this can be a good idea.
 - BUT** we're likely to create a memory-bound problem.
- Or, we can make each thread do many somethings.

```
__global__ myKernel(int m, ...) {
    for(int i = 0; i < m; i++) {
        do something
    }
}
```

Preview

March 20: Matrix multiplication, Part 1
March 22: Matrix multiplication, Part 2
March 24: Complete CUDA
March 27 – April 3: this may change
March 27: Using Parallel Libraries
March 29 – April 3: Verification of/and Parallel Programs
April 5: Party: 50 th Anniversary of Amdahl's Law

Matrix Multiplication – Algorithms

Mark Greenstreet

CpSc 418 – Mar. 22, 2017

Outline:

- Sequential Matrix Multiplication
- Parallel Implementations, Performance, and Trade-Offs.

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license <http://creativecommons.org/licenses/by/4.0/>

Objectives

Apply concepts of algorithm analysis, parallelization, overhead, and performance measurement to a real problem.

- Design sequential and parallel algorithms for matrix multiplication.
- Analyse algorithms and measure performance.
- Identify bottlenecks and refine algorithms.

Matrix representation in Erlang

- I'll represent a matrix as a list of lists.
- For example, the matrix $\begin{bmatrix} 1 & 2 & 3 & 4 \\ 1 & 4 & 9 & 16 \\ 1 & 8 & 27 & 64 \end{bmatrix}$ is represented by the Erlang nested-list:

```
[ [ 1, 2, 3, 4 ],
  [ 1, 4, 9, 16 ],
  [ 1, 8, 27, 64 ] ]
```
- The empty matrix is `[]`.
 - This means my representation can't distinguish between a 2×0 matrix, a 0×4 matrix, and a 0×0 matrix.
 - That's OK. This package is to show some simple examples.
 - I'm not claiming it's for advanced scientific computing.

Sequential Matrix Multiplication

```
mult(A, B) ->
    BT = transpose(B),
    lists:map(
        fun(RowOfA) ->
            lists:map(
                fun(ColOfB) ->
                    dot_prod(RowOfA, ColOfB)
                end, BT)
            end, A).
```

```
dot_prod(V1, V2) ->
    lists:foldl(
        fun({X,Y}, Sum) -> Sum + X*Y end,
        0, lists:zip(V1, V2)).
```

- Next, we'll use `list comprehensions` to get a more succinct version.

Matrix Multiplication, with comprehensions

```
mult(A, B) ->
    BT = transpose(B),
    [ [ dot_prod(RowA, ColB) | ColB <- BT ] | RowA <- A ].
```

```
transpose([]) -> []; % special case for empty matrices
transpose([[]:[]]) -> []; % bottom of recursion, the columns are empty
transpose(M) ->
    [ [ H | [ H | T ] <- M ] <- M ] % create a row from the first column of M
    | transpose([ T | [ T | T ] <- M ]) % now, transpose what's left
    ].
```

- Parallelize the outer-loop.
- Each iteration of the outer-loop multiplies a row of `A` by all of `B` to produce a row of `A × B`.
- Divide `A` (and `B`) into blocks.
- Each processor sends its blocks of `B` to all of the other processors.
- Now, each processor has a block of rows of `A` and all of `B`. The processor computes its part of the product to produce a block of rows of `C`.
- Note: OpenMP does this kind of parallelization automatically.

Performance – Modeled

- Really simple, operation counts:
 - Multiplications: $n_rows_a * n_cols_b * n_cols_a$.
 - Additions: $n_rows_a * n_cols_b * (n_cols_a - 1)$.
 - Memory-reads: $2 * n_Multiplications$.
 - Memory-writes: $n_rows_a * n_cols_b$.
 - Time is $O(n_rows_a * n_cols_b * n_cols_a)$.
 - If both matrices are $N \times N$, then it's $O(N^3)$.
- But, memory access can be terrible.
 - For example, let matrices `a` and `b` be 1000×1000 .
 - Assume a processor with a 4M L2-cache (final cache), 32 byte-cache lines, and a 200 cycle stall for main memory accesses.
 - Observe that a row of matrix `a` and a column of `b` fit in the cache. (a total of ~40K bytes).
 - But, all of `b` does not fit in the cache (that's 8 Mbytes).
 - So, on every fourth pass through the inner loop, **every** read from `b` is a cache miss!
 - Cache miss dominates everything else.
- This is why there are carefully tuned numerical libraries.

Performance – Measured

- Cubic of best fit: $T = (107N^3 + 134N^2 + 173N - 32)\text{ns}$.
- Fit to first six data points.
- Cache misses effects are visible, for $N=1000$:
 - model predicts $T = 107\text{seconds}$,
 - but the measured value is $T = 142\text{seconds}$.

Tiling Matrices

An Example

- Let `A`, `B`, and `C = AB` be 16×16 matrices.
- Let `A1 = A[1 : 4, 1 : 16]`, i.e. the first four rows of `A`.
 - In our Erlang representation, `[A1, _] = lists:split(4, A)`.
- Let `A2 = A[5 : 8, 1 : 16]`; `A3 = A[9 : 12, 1 : 16]`; `A4 = A[13 : 16, 1 : 16]`; and likewise for `C1`, `C2`, `C3`, and `C4`.
- Big important fact:

$C_1 = A_1 B$

$C_2 = A_2 B$

$C_3 = A_3 B$

$C_4 = A_4 B$
- In **sequential** Erlang:

```
[C1, C2, C3, C4] = [mult(AA, B) | AA <- A]
```

 - To make it **parallel**, we compute each of the `Cj = Aj B` with a separate process.

Parallel Algorithm 1

- Parallelize the outer-loop.
- Each iteration of the outer-loop multiplies a row of `A` by all of `B` to produce a row of `A × B`.
- Divide `A` (and `B`) into blocks.
- Each processor sends its blocks of `B` to all of the other processors.
- Now, each processor has a block of rows of `A` and all of `B`. The processor computes its part of the product to produce a block of rows of `C`.
- Note: OpenMP does this kind of parallelization automatically.

Parallel Algorithm 1 in Erlang

```
% mult(W, Key1, Key2) - create a matrix associated with Key
% that is the product of the matrices associated with Key1 and Key2.
mult1(W, Key1, Key2) ->
    Nproc = workers:workers(W),
    workers:update(W, Key1,
        fun(PS, I) ->
            A = workers:get(PS, Key1), % my rows of A
            B = workers:get(PS, Key2), % my rows of B
            [WW : [B, _] | WW <- W], % send my rows of B to everyone
            B_full = lists:append(), % receive BB from everyone
            [ receive [BB, _] -> BB end
              | J <- lists:seq(1, Nproc) ]),
            matrix:mult(A, B_full) % compute my part of the product
        end
    ).
```

Performance of Parallel Algorithm 1 – Modeled

- CPU operations: same total number of multiplies and adds, but distributed around P processors. Total time: $O(N^3/P)$.
- Communication: Each processors sends (and receives) $P - 1$ messages of size N^2/P . It takes time to send a message is $t_0 + t_1 \times M$ where M is the size of the message, then the communication time is $(P - 1) \left(t_0 + t_1 \frac{N^2}{P} \right) = O(N^2 + \lambda P)$ but, **beware of large constants** $= O(N^2)$, $N^2 > P$
- Note: I'm assuming t_0 corresponds to λ , and that t_1 is roughly the same as the time for "typical" sequential operations..
- Memory: Each process needs $O(N^2/P)$ storage for its block of `A` and the result. It also needs $O(N^2)$ to hold **all** of `B`.
 - The simple algorithm divides the computation across all processors, but it doesn't make good use of their combined memory.

Performance of Parallel Algorithm 1 – Measured

Using Memory More Efficiently

- Main idea: each process works on one "slab" of `B` at a time.

$C[i, j] = \sum_{k=1}^{N/4} A[i, k] B[k, j]$, a dot-product

$= \left(\sum_{k=1}^{N/4} A[i, k] B[k, j] \right) + \left(\sum_{k=N/4+1}^{N/2} A[i, k] B[k, j] \right) + \left(\sum_{k=N/2+1}^{3N/4} A[i, k] B[k, j] \right) + \left(\sum_{k=3N/4+1}^N A[i, k] B[k, j] \right)$
- Each process does each of its four summations when it holds the corresponding slab of `B`.
 - Each holds one slab of `A` for the whole computation.
 - Each process only needs to hold one slab of `B` at a time.
- The algorithm generalizes to having any number of slabs for `A` and `B` in the obvious way.
 - Should be "obvious" if I've explained this clearly.
 - If it isn't obvious, that's my bad – please ask a question.

Parallel Algorithm 2 (illustrated)

Parallel Algorithm 2 (code sketch)

- Each processor first computes what it can with its rows from `A` and `B`.
 - It can only use N/P of its columns of its block from `A`.
 - It uses its entire block from `B`.
 - We've now computed one of P matrices, where the sum of all of these matrices is the matrix `AB`.
- We view the processors as being arranged in a ring.
 - Each processor forwards its block of `B` to the next processor in the ring.
 - Each processor computes an new partial product of `AB` and adds it to what it had from the previous step.
 - This process continues until every block of `B` has been used by every processor.

Performance of Parallel Algorithm 2

- CPU operations: Same as for parallel algorithm 1: total time: $O(N^3/P)$.
- Communication: Same as for parallel algorithm 1: $O(N^2 + P)$.
 - With algorithm 1, each processor sent the same message to $P - 1$ different processors.
 - With algorithm 2, for each processor, there is one destination to which it sends $P - 1$ different messages.
 - Thus, algorithm 2 can work efficiently with simpler interconnect networks.
- Memory: Each process needs $O(N^2/P)$ storage for its block of `A`, its current block of `B`, and its block of the result.
 - Note: each processor might hold onto its original block of `B` so we still have the blocks of `B` available at the expected processors for future operations.
- Do the memory savings matter?

Bad performance, pass it on

- Consider what happens with algorithm 2 if one processor, P_{slow} takes a bit longer than the others one of the times its doing a block multiply.
 - P_{slow} will send it's block from `B` to its neighbour a bit later than it would have otherwise.
 - Even if the neighbour had finished its previous computation on time, it won't be able to start the next one until it gets the block of `B` from P_{slow} .
 - Thus, for the next block computation, both P_{slow} and its neighbour will be late, even if both of them do their next block computation in the usual time.
 - In other words, tardiness propagates.
- Solution: forward your block to you neighbour **before** you use it to perform a block computation.
 - This overlaps computation with communication, generally a good idea.
 - We could send two or more blocks ahead if needed to compensate for communication delays and variation in compute times.
 - This is a way to save time by using more memory.

Tiling in Real-Life

- Why? If there's time, I'll explain in class.

Summary

- Matrix multiplication is well-suited for a parallel implementation.
- Need to consider communication costs.
- In the previous algorithms, compute time grows as N^3/P , while communication time goes as $(N^2 + P)$.
- Thus, if N is big enough, computation time will dominate communication time.
- Connection of theory with actual run time is pretty good:
 - But the matrices have to be big enough to amortize the communication costs.


Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license <https://creativecommons.org/licenses/by/4.0/>

- ▶ wall-clock $\sim \alpha N^2$ with α determined mainly by global memory bandwidth.
- ▶ measured: $T(1024) \approx 0.0986s$; $T(2048) \approx 0.797s$; $T(3072) \approx 2.7s$; $T(4096) \approx 6.3s$.

$N^2/T(N) \approx 11/ns$ – i.e. about 20×10^9 multiply-adds per second. Well below GPU peak floating point capacity. Demonstrates global memory bandwidth bottleneck (with a little help from the on-chip caches).

Mark Greenstreet	CUDA: Matrix Multiplication	CS 418 – Mar. 24, 2017	3 / 12
------------------	-----------------------------	------------------------	--------

Tiles vs. Slabs



- Matrix multiplication: each processor (color) has tiles at (i,j) and (j,i) .
 - Can compute all products for the main diagonal, and stripes at spacings of P .
 - Use a reduce to combine results to get the main diagonal and the stripes.
 - Rotate B one block to the left, and compute the next set of stripes.
 - After P rounds, the computation is done.
 - Same amount of work (and communication) as the improved slab method from Wednesday.
- Other algorithms such as LU-Decomposition
 - Rows and columns are eliminated from the left and the top.
 - Tiles provide better load balancing.

Mark Greenstreet	CUDA: Matrix Multiplication	CS 418 – Mar. 24, 2017	4 / 12
------------------	-----------------------------	------------------------	--------

Performance of mmult2

- $T(1024) = 0.027s$; $T(2048) = 0.214s$; $T(3072) = 0.742s$;
 $T(4096) = 1.73s$.
- Still cubic in N , of course.
 - ▶ $N^3/T(N) \approx 40/ns$ – about 40 billion multiply-adds per second.
 - ▶ About four times faster than `mmult1`.

Mark Greenstreet	CUDA: Matrix Multiplication	CS 41B – Mar. 24, 2017	8 / 12
------------------	-----------------------------	------------------------	--------

Preview

Mark Greenstreet	CUDA: Matrix Multiplication	CS 41B – Mar. 24, 2017	12 / 12
------------------	-----------------------------	------------------------	---------

Basic Linear Algebra Subprograms (BLAS)

- Encourages "structured programming": Modularization of common code sequences.
- Code will be more self-documenting: Other programmers will recognize the subprogram names.
- Subprograms can be coded in assembly to improve efficiency, and if the majority of computational effort is within the subprograms that will significantly benefit the whole application.
- Subprograms can be coded by experts to deal with "algorithmic and implementation subtleties."
- Code becomes portable while still maintaining efficiency.

While the details may differ, similar benefits still accrue today.

Greenstreet & Mitchell Linear Algebra Libraries and CUDA CPSC 418 – March 27, 2017 4 / 15

Decypering BLAS Function Names

- Often a prefix, such as `BLAS_` or `cblas_`.
- One character to denote data type; for example:
 - ▶ `s`: single precision.
 - ▶ `d`: double precision.
- Operations involving a matrix add two characters to denote matrix type; for example:
 - ▶ `ge`: general dense matrix.
 - ▶ `tb`: triangular banded.
- Short mnemonic string to denote operation; for example
 - ▶ `sxpy`: $ax + y$.
 - ▶ `mm`: matrix multiply.
- Put them all together:
 - ▶ `BLAS_SAXPY()`: Fortran single precision vector summation.
 - ▶ `cblas_dgemv()`: C double precision dense matrix product.

Greenstreet & Mitchell Linear Algebra Libraries and CUDA CPSC 418 – March 27, 2017 8 / 15

CUDA and BLAS

- The [cuBLAS](#) library provides an API for running BLAS routines on CUDA GPUs.

- The **cuBLAS** library provides an API for running BLAS routines on CUDA GPUs.
- Basic pattern of use:
 - Initialize the cuBLAS library and allocate hardware resources using `cudaSetDevice()`.
 - Allocate memory using `cudaMalloc()`.
 - Copy data from host to GPU using `cudablasSetVector()` or `cudablasSetMatrix()`.
 - Perform BLAS operations; for example `cudablasSaxpy()` or `cudablasGemm()`.
 - Copy data from GPU to host using `cudablasGetVector()` or `cudablasGetMatrix()`.
 - Release memory using `cudaFree()`.
 - Release hardware resources using `cudablasDestroy()`.
- Example(s)

Greenstreet & Mitchell	Linear Algebra Libraries and CUDA	CPSC 418 – March 27, 2017	12 / 15
------------------------	-----------------------------------	---------------------------	---------

on CUBLAS

- Always uses column-major ordering
 - So be careful with data layout.
- Always uses 1-based indexing.
 - Usually irrelevant since you do not index into arrays.
- All cuBLAS code is called from the host.
 - You do not write any kernel code.
 - You do not have to worry about grids, blocks, shared memory, ...
- Need to link against cuBLAS library.
 - Check that environment variable `LD_LIBRARY_PATH` includes CUDA library directory.
 - `/usr/local/lib64/pkgconfig/cudatoolkit/lib64` on `linux` machines.
 - Add `-lcublas` to compile command.

Efficiency of cuBLAS

Matrix product example from 2017-03-24 lecture (in seconds):

	1024	2048	3072	4096
Brute force <code>mmtul1</code>	0.079	0.648	2.190	5.152
Tiled <code>mmtul2</code>	0.027	0.208	0.724	1.690
<code>cublas_sgemv</code>	0.007	0.052	0.176	0.421

- Brute force `mmtul1` achieves ~ 13 GFLOPS.
- `cublas_sgemv` achieves ~ 160 GFLOPS.
- (Ian's tiled implementation `mmtul2` is buggy).

Linear Algebra PACKage (LAPACK).

- Implements the more complex linear algebra operations.
- Designed to call BLAS for basic computational steps.

For your CPU:

- Intel's Math Kernel Library (MKL) implements core functions from BLAS, LAPACK, FFTs, etc.
- Automatically Tuned Linear Algebra Software (ATLAS) generates a BLAS library tuned to a machine's memory hierarchy.

Many other accelerated libraries available for CUDA devices.

- For example: cuFFT, cuSPARSE, cuRAND, cuDNN, MAGMA (supports LAPACK), ...

You are almost always better off learning to use the library.

Model Checking

Mark Greenstreet

CpSc 418 – Mar. 29, 2017

- Motivation
- Today's paper
- Applications of Model Checking

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

Model-Checking: Motivation

- What is "model checking"?
 - Construct a "model" for a piece of hardware or software – typically a finite-state machine.
 - Give a precise, mathematical definition of properties that the design is supposed to have.
 - Show that that model satisfies the specification.
 - For example, find all reachable states of the model.
 - Show that every reachable state satisfies a desired property – for example, mutual exclusion.
- Why use model checking?
 - Find bugs.
 - Hardware bugs are very expensive.
 - Software bugs are very common, but
 - Finding bugs in concurrent software is hard.
 - The challenges of finding bugs motivates using more systematic approaches.
- A simple example: Dekker's Mutual Exclusion algorithm

Dekker's Algorithm

Problem statement: ensure that at most one thread is in its critical section at any given time.

thread 0:

```
PC0= 0; while(true) {
PC0= 1; non-critical code
PC0= 2; flag[0] = true;
PC0= 3; while(flag[1]) {
PC0= 4; if(turn != 0) {
PC0= 5; flag[0] = false;
PC0= 6; while(turn != 0);
PC0= 7; flag[0] = true;
PC0= 8; }
PC0= 9; }
PC0=10; critical section
PC0=11; turn = 1;
PC0=12; flag[0] = false;
PC0=13; }
```

thread 1:

```
PC1= 0; while(true) {
PC1= 1; non-critical code
PC1= 2; flag[1] = true;
PC1= 3; while(flag[0]) {
PC1= 4; if(turn != 1) {
PC1= 5; flag[1] = false;
PC1= 6; while(turn != 1);
PC1= 7; flag[1] = true;
PC1= 8; }
PC1= 9; }
PC1=10; critical section
PC1=11; turn = 0;
PC1=12; flag[1] = false;
PC1=13; }
```

See http://en.wikipedia.org/wiki/Dekker's_algorithm.

Is Dekker's algorithm correct?

- Dijkstra (Turing Award 1972), presented the algorithm, with a proof in 1965.
- We'll use it as an example for model-checking:
 - Construct a finite-state machine model of the algorithm.
 - Determine the set of reachable states.
 - Verify that all reachable states satisfy mutual exclusion.
 - We could check other properties as well.
 - For example, freedom from starvation: show that if a process is attempting to enter it's critical region, it will eventually succeed.

Model Checking Dekker's algorithm

- Represent each state with 9-bits:
 - three for the location of each process (6 locations)
 - three for `flag` and `turn`
 - I'll show a simple python version that uses python tuples
- Pseudo-code:

```
initialState = (1,1,0,0,0,0) // (loc0, loc1, flag0, flag1, turn)
worklist = queue() // initially empty
knownStates = set() // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    for s' in next_states(s):
        if s' not in knownStates:
            check s' for mutual exclusion;
            add s' to worklist and knownStates
```
- Model-checking finds 48 reachable states for Dekker's algorithm and verifies mutual exclusion.

A Brief History of Model Checking

- Proposed by Clarke and Emerson (1981) and independently by Sipakis (1982).
 - They shared the 2007 Turing Award.
 - Their approach was essentially the one described above.
- Symbolic methods introduced by McMillan (1987) using binary-decision diagrams, a DAG representation of boolean formulas.
- Widespread adaptation of model-checking for hardware design took place in the 1990s and continues today.
 - The `murϕ` model checker is a landmark in this work.
- Model-checking of software is now gaining industrial acceptance
 - Based on "predicate abstraction" methods of Clarke and Grumberg, and independently Ball.
 - Enabled by advances in boolean SAT solvers and interpolation-based model checking (McMillan).

Today's Paper

Protocol Verification as a Hardware Design Aid

Mark's standard five questions:

- What problem does the paper address?
 - Hardware designs consist of large blocks that communicate using protocols. Mistakes in the protocol design can cause subtle errors that only occur in rare corner cases. Such errors are hard to find by traditional simulation.
- What is the key idea in the paper?
 - Use model checking to exhaustively verify small versions of the design.
- How do the authors validate their idea?
 - They implemented a model checker.
 - This included defining a modeling language so that protocols can be described easily and clearly.
 - They applied their approach to two protocols from real designs in industry.
- Is the paper convincing?
 - Yes: they showed that they could check important properties of two "down scaled" protocols.
 - No: the protocols seem down-scaled to the edge of being trivial.
 - 20/20 hindsight: Definitely! Model-checking methods have evolved and matured and are now widely used in industry for both hardware and software.
- Any other comments?
 - Glad you asked. See the rest of the lecture.

`murϕ`: a guarded command language

- In `murϕ` a guarded command is called a rule and is written:

```
rule guard => action
```

 - When `guard` is satisfied, `action` may be performed.
 - Example: `rule (loc0 == 3) and flag[1] => loc0[] := 4`
- Rules may be quantified using the `Ruleset` construction:

```
Process: scalarset(2);
ruleset i: Process do
    (loc[i] == 3) and flag[1-i] => loc[i] := 4
end
```

`murϕ`: execution model

- A program defines a fixed set of rules.
- A Parallel Algorithm for Model Checking
- Implementing a Parallel Model Checker
- Summary, Preview, & Review

Review

Industrial Strength, Parallel Model Checking

Mark Greenstreet

CpSc 418 – Mar. 31, 2017

- The Big Five
- A Parallel Algorithm for Model Checking
- Implementing a Parallel Model Checker
- Summary, Preview, & Review

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

Model Checking: the algorithm

```
initialState = ...;
worklist = queue(); // initially empty
knownStates = set(); // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    check s' for mutual exclusion;
    for s' in next_states(s):
        if s' not in knownStates:
            add s' to worklist and knownStates;
```

How can we make this algorithm parallel?

- Where does the code spend most of the time?

Model Checking: the algorithm

```
initialState = ...;
worklist = queue(); // initially empty
knownStates = set(); // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    check s' for mutual exclusion;
    for s' in next_states(s):
        if s' not in knownStates:
            add s' to worklist and knownStates;
```

How can we make this algorithm parallel?

- Where does the code spend most of the time? `next_states(s)`

Model Checking: the algorithm

```
initialState = ...;
worklist = queue(); // initially empty
knownStates = set(); // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    check s' for mutual exclusion;
    for s' in next_states(s):
        if s' not in knownStates:
            add s' to worklist and knownStates;
```

How can we make this algorithm parallel?

- Where does the code spend most of the time? `next_states(s)`
- What are the **dependencies**?
 - As written, the code is very sequential.
 - BUT, the correctness of the algorithm does not depend on the order in which states are removed from `worklist`.
- What uses most of the memory?

Model Checking: the algorithm

```
initialState = ...;
worklist = queue(); // initially empty
knownStates = set(); // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    check s' for mutual exclusion;
    for s' in next_states(s):
        if s' not in knownStates:
            add s' to worklist and knownStates;
```

How can we make this algorithm parallel?

- Where does the code spend most of the time? `next_states(s)`
- What are the **dependencies**?
 - As written, the code is very sequential.
 - BUT, the correctness of the algorithm does not depend on the order in which states are removed from `worklist`.
- What uses most of the memory? `KnownStates` and `worklist`

Modeling Dekker's algorithm

thread 0: code

```
PC0= 0; while(true) {
PC0= 1; non-critical code
PC0= 2; flag[0] = true;
PC0= 3; while(flag[1]) {
PC0= 4; if(turn != 0) {
PC0= 5; flag[0] = false;
PC0= 6; while(turn != 0);
PC0= 7; flag[0] = true;
PC0= 8; }
PC0= 9; }
PC0=10; critical section
PC0=11; turn = 1;
PC0=12; flag[0] = false;
PC0=13; }
```

thread 0: state machine

Overview of the paper

- How does model the hardware?
 - `murϕ`: a guarded command language.
- How do we state the properties to be verified?
 - Add assertions to the `murϕ` program.
 - Use model checking to show that these assertions hold for all states of all executions.
- How do we perform the model checking?
 - Compile the `murϕ` program to C++.
 - Link with an efficient implementation of a model checking algorithms like the one in [dekker.nc.py](#).
 - Run the model checker to either verify the properties or report counter-examples.

Model checking today

- Lots of progress on handling larger models:
 - Symbolic methods
 - Exploit common model properties: symmetry, commuting-actions, verifiable abstraction
 - Moore's law: faster machines, larger memories.
 - Parallelism (Friday's lecture)
- Applications
 - An essential part of cache-protocol design. Used in many other aspects of hardware design as well.
 - Software: Microsoft uses model checking to verify that driver code conforms to kernel usage rules.
 - Software: Amazon uses model-checking to verify protocols used in their cloud services.
 - Many others.

Mark's Five Questions: 1–3

- What problem does the paper address?
 - Model checking of real problems in industry requires the combined memory of many machines. Existing implementations of parallel model checkers failed to successfully run on real examples. A robust, distributed, parallel model checker is needed.
- What is the key idea in the paper?
 - Combine `murϕ` and Erlang.
 - Use the existing `murϕ` code to implement the computationally intensive part of the code.
 - Use Erlang to handle the communication and coordination between worker processes.
- How do the authors validate their idea?
 - They implemented a model checker.
 - They performed experiments to identify robustness issues and performance bottlenecks. They implemented solutions to these problems.
 - They used the PReach model checker on several benchmark examples and some real designs from Intel.

Model Checking: the algorithm

```
initialState = ...;
worklist = queue(); // initially empty
knownStates = set(); // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    check s' for mutual exclusion;
    for s' in next_states(s):
        if s' not in knownStates:
            add s' to worklist and knownStates;
```

How can we make this algorithm parallel?

- Where does the code spend most of the time? `next_states(s)`
- What are the **dependencies**?

Making it Parallel

```
initialState = ...;
worklist = queue(); // initially empty
knownStates = set(); // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    check s' for mutual exclusion;
    for s' in next_states(s):
        if s' not in knownStates:
            add s' to worklist and knownStates;
```

How can we make this algorithm parallel?

- Divide `knownStates` and `worklist` across the worker processes? How?
 - Sending each new state, `s'` to every worker is a bad idea. Why?
 - We'll use hashing instead: send state `s'` to worker process `hash(s') mod P`.
- Each worker maintains `knownStates` for the states that hash to the worker process.
- Each worker adds each new state it receives to `knownStates` and `worklist`.
- Each worker processes the states on its own `worklist`.

Model Checking

Mark Greenstreet

CpSc 418 – Mar. 29, 2017

- Motivation
- Today's paper
- Applications of Model Checking

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

Model Checking

Mark Greenstreet

CpSc 418 – Mar. 29, 2017

- Motivation
- Today's paper
- Applications of Model Checking

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

Model Checking

Mark Greenstreet

CpSc 418 – Mar. 29, 2017

- Motivation
- Today's paper
- Applications of Model Checking

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

Model Checking

Mark Greenstreet

CpSc 418 – Mar. 29, 2017

- Motivation
- Today's paper
- Applications of Model Checking

Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<http://creativecommons.org/licenses/by/4.0/>)

Model Checking: the algorithm

```
initialState = ...;
worklist = queue(); // initially empty
knownStates = set(); // initially empty
worklist.insert(initialState);
while len(worklist) > 0:
    s = worklist.removeNext();
    check s' for mutual exclusion;
    for s' in next_states(s):
        if s' not in knownStates:
            add s' to worklist and knownStates;
```

How can we make this algorithm parallel?

- Where does the code spend most of the time? `next_states(s)`
- What are the **dependencies**?
 - As written, the code is very sequential.
 - BUT, the correctness of the algorithm does not depend on the order in which states are removed from `worklist`.

Parallel Model Checking: the Pseudo-Code

```
modelCheck(KnownStates, Worklist) ->
receive
    S when member(S, KnownStates) -> % already seen
    modelCheck(KnownStates, Worklist);
    S -> modelCheck(add(S, KnownStates),
        add(S, Worklist))
after 0 ->
    case Worklist of
    [] -> if everyone_else_done();
    [S | T1] ->
        [owner(S') : S' | S' <- next_states(S)],
        modelCheck(KnownStates, T1)
    end
end
```

- See [Parallelizing the `murϕ` model checker](#), U. Stern and D.L. Dill.

Parallel Model Checking: Performance Analysis

- A sequential implementation of the model checking algorithm requires $\mathcal{O}(SR)$ time, where S is the number of reachable states, and R is the number of rules.
- A parallel implementation requires $\mathcal{O}(SR/P)$ compute time, and sends worst-case $\mathcal{O}(SR)$ messages.
 - In practice, the average number of successors of each state (i.e. the degree of the state-graph) is relatively small. If we assume this is a small constant, then we get $\mathcal{O}(S)$ messages.
- Consider the case where each worker process generates σ new successor states, σ^P per second.
 - These are sent to the other processes uniformly at random (if we have a good hash function).
 - Half of these messages cross the bisection of any network.
 - That means we need a bisection bandwidth of $\sigma P/2$.
- For real-life networks, bisection bandwidth grows much more slowly than P .
 - If we scale this algorithm to a large enough number of processors, network bandwidth will be the limiting constraint.
 - This is a common performance pattern in parallel computing.

Mark Greenstreet	Industrial Strength, Parallel Model Checking	CS 418 – Mar 31, 2017	6 / 17
------------------	--	-----------------------	--------

Batching Messages

- If we send each new state, s^i , one at a time to its owner process, communication overhead dominates the run time.
- **Key lesson:** pay attention to λ .
- The Erlang code maintains a separate buffer for each worker process
 - Add states that should be sent to that process to the buffer until we have enough.
 - Then send them as a batch.
 - A process that is running low on work sends requests to the other workers to ask them to flush their buffers.
 - These flush requests are bundled with state batches to avoid extra messages.
- Erlang makes the communication architecture simple and easy to extend.

Flexibility

- The Erlang code for PReach is simple.
 - The version described in the paper is about 1000 lines of code.
- This makes PReach a flexible platform for experiments:
 - Checking response properties: e.g. every request is eventually granted.
 - Exploiting symmetry: there are times we can verify a protocol for two or three nodes and conclude with certainty that it is correct for any number of nodes.
 - And others.
- Applications:
 - Used by Intel architects when exploring protocols for on-chip networks.
 - Used in other companies and universities.
 - It's been run on hundreds of machines with models of hundreds of billions of states.
 - Symbolic methods are faster than PReach for safety properties (showing that the model never reaches a bad state)
 - PReach is faster for handling liveness properties: showing that some condition will eventually be satisfied.

Mark Greenstreet	Industrial Strength, Parallel Model Checking	CS 418 – Mar 31, 2017	14 / 17
------------------	--	-----------------------	---------

Review: for today's lecture

- To make a parallel implementation of a computation, we often need to identify a set (or many set(s)) of operations that can be performed in any order.
 - For model checking, what set of operations did we find that can be performed in any order? Does this exactly replicate the sequential version, or does it perform an equivalent computation?
 - Same questions as above, but for reduce.
 - Scan? Sorting? Matrix multiplication?
- Why did slow processes in PReach tend to become catastrophically slower? What was the solution?
- What is load balancing? Compare the load balancing mechanisms of PReach and Google's map-reduce.
- Is Erlang suitable for large-scale, high-performance, parallel computing? Why or why not?

Parallel Algorithms 2

- Matrix operations
 - matrix multiplication dividing the matrix into blocks
 - Dividing the matrix into blocks
 - Analysis of the compute and communication costs.
 - BLAS and cuBLAS: use a library when you can!
- Map-Reduce
- Model checking
 - We only had two lectures on the topic and no HW.
 - Won't ask any detailed questions, but if there might be some high-level questions with one sentence answers in the review questions, in which case similar questions could be on the exam.
 - Know what model checking is: verifying properties of a hardware or software design, using a finite-state-machine model, finding the reachable states.
 - The idea of distributing work by hashing values and sending each to its owner process.

Mark Greenstreet	Course Summary	CS 418 – Apr 5, 2017	3 / 8
------------------	----------------	----------------------	-------

... and the things we have left undone

- More paradigms and programming frameworks
 - shared memory: Java threads, pthreads.
 - futures: e.g. Scala
 - MPI and OpenMP (for scientific computing)
 - many big-data, machine-learning, and scientific computing frameworks
- Do a bigger project.
- **The good news:**
 - You've got what you need to learn new paradigms, new frameworks, and take on realistic projects.
 - From my experience with research projects that have moved into industry in the past few years, you've got the critical knowledge and skills.
 - Writing industrial-strength, parallel-code with good performance is still more than a homework assignment, but when my students have done it, they've built on the foundation you now have.

Mark Greenstreet	Course Summary	CS 418 – Apr. 5, 2017	7 / 8
------------------	----------------	-----------------------	-------

From Algorithm to Industrial Adaptation

What is needed for real-world verification?

- Lots of memory:
 - Memory and time are both concerns for model checkers, but memory tends to be the more critical concern.
 - A parallel implementation offers the combined memory of a large number of machines.
- Robustness:
 - Simple architecture and re-use stable, well-exercised code.
 - Prevent "overwhelm and crash".
 - Load balancing.
- Flexibility:
 - Solve problems that other tools cannot
 - In particular, liveness properties such as "response".

Overwhelm and Crash

The dangers of using Erlang for high-performance computing

- The Erlang in-box is a list.
 - Newly received messages are prepended to the list.
 - A receive gets the oldest message that matches a pattern of the receive.
 - This means that the time for receive is linear in the number of pending message.
- This leads to a performance catastrophe
 - If a process gets slightly behind, its inbox will fill a little more than the other processes.
 - This means that a process that falls behind will slow down, and its inbox will fill even more.
 - Eventually, the process crashes.

Termination

- How do we know when we're done?
- Well, times up for this lecture.
- More seriously, in PReach we need to know when
 - When every worker process has an empty worklist,
 - **And** there are no messages in flight.
- Both conditions must hold at the same time.
 - This is the topic for Monday's lecture.

Review: for March 29 lecture

- What is model checking?
- What is mutual exclusion?
- How does the model checker presented in the March 29 slides show that Dekker's algorithm guarantees mutual exclusion.
- Describe the role of the `knownStates` and `workList` data structures in the model checking algorithm.
- What is a guarded command?
- Write a `murφ` rule for another statement from Dekker's algorithm.

Parallel Architectures

- pipelining and instruction level parallelism
- shared memory multiprocessors
 - Know what a cache coherence protocol is.
 - Explain the idea of shared-reader or exclusive writer.
 - Be able to point out that real cache coherence protocols aren't as "consistent" as the simple (e.g. MESI) model from class.
- message passing architectures
 - Rings, tori, hypercubes
 - Latency, bisection width.
- data parallel architectures
 - SIMD (and SMT)
 - instruction execution: why so many threads
 - GPU memory hierarchy

Mark Greenstreet	Course Summary	CS 418 – Apr. 5, 2017	4 / 8
------------------	----------------	-----------------------	-------

Erlang to high-performance computing (really)

- Use existing C++ code for `murφ`.
 - It has been carefully optimized – it's fast.
 - It has been widely used over the past 25 years – it's robust.
- Use Erlang to make it parallel
 - Erlang handles the communication between processes.
 - The code is simple: it works and it's flexible.
 - Erlang can't call the C++ functions:
 - The compute intensive part is done in C++.
 - The Erlang code is not a serious bottleneck.

Credits

Preventing "overwhelm and crash"

- Drain the inbox into another buffer whenever possible.
- Maintain a credit system
 - When a process X sends a message to process Y , X decrements its credit-count for Y .
 - If the credit-count is 0, X waits to send its message.
 - When Y moves a message from X out of its inbox, it sends a credit back to X .
 - Of course, these messages are piggy-backed on the new-state messages.

Summary


- PReach shows how the ideas from this class can be used to build real-world, high-performance, large-scale, parallel systems.
- Lessons learned:
 - Erlang is a great environment for building large-scale, parallel/distributed code.
 - Use the C/C++ call interface to use native C/C++ code for the compute intensive parts of the code.
 - Erlang provides three such interfaces!
 - Watch out for overwhelm and crash
 - If you're going to send a lot of messages, you need some kind of flow control mechanism.

Course Summary

Mark Greenstreet

CpSc 418 – Apr. 5, 2017

- The things we have done:
 - parallel algorithms
 - parallel architectures
 - parallel performance
 - parallel paradigms
- and the things we have left undone.
- An exam, but first, a party.

 Unless otherwise noted or cited, these slides are copyright 2017 by Mark Greenstreet and are made available under the terms of the Creative Commons Attribution 4.0 International license (<https://creativecommons.org/licenses/by/4.0/>)

Parallel Performance

- λ
 - Communication costs are critical to understanding parallel performance
 - This is true across all parallel architectures.
- Overheads and losses
 - Communication, synchronization, extra memory, extra computation
 - Idle processors, resource contention, non-parallelizable code
- Speed-up and Amdahl's Law.
- **Understanding real world performance requires experiments and measurements.**

Mark Greenstreet	Course Summary	CS 418 – Apr. 5, 2017	5 / 8
------------------	----------------	-----------------------	-------

Memory

- The worklist is the dominant use of memory (in practice)
 - Why?
 - The worklist needs complete state descriptions.
 - The known-state set can use much smaller hashes.
- Solution: store the worklist on disk.
 - Disks are slow – is this crazy?
 - It works just fine because we can access the worklist in **any** order.
 - Keep a large piece of the worklist in main memory.
 - If the in-memory work list grows too large, then copy a large chunk to disk.
 - If the in-memory work list becomes too small, then read a large chunk from disk.
 - The disk reads and writes can be performed asynchronously.
 - See *Using magnetic disks ... in the `murφ` model checker*, U. Stern and D.L. Dill.
- Storing the known-state set on disk is much less practical because it's a random-access structure.

Mark Greenstreet	Industrial Strength, Parallel Model Checking	CS 418 – Mar 31, 2017	9 / 17
------------------	--	-----------------------	--------

Load Balancing

- Not all processes have the same amount of work, and they don't all run at the same speed.
- This can lead to **idle processors**.
- Solution:
 - Processes include the length of their worklist in their messages to other workers.
 - If a worker has a short worklist, it asks for half of the worklist of the worker with the longest worklists.
- This is a very coarse-grained approach
 - PReach makes no effort to keep worklist lengths equal.
 - The coarse-grained approach requires very few messages: **avoid** λ .
 - The performance is very good.

Preview

April 3: Distributed Termination Detection

April 5: Party: 50th Anniversary of Amdahl's Law

Parallel Algorithms 1

- map, reduce, and scan: simple patterns
 - Easy to parallelize.
 - Learn to recognize when a problem can be solved by these simple methods.
- sorting networks
 - oblivious computation: the control flow doesn't depend on data values.
 - oblivious algorithms are good candidates for parallelism because we can determine the control flow in advance.
 - This lets us identify the data dependencies, and find a parallel solution.
 - The 0-1 principle
 - Bitonic sorting: it's merge sort with an oblivious merge.

Parallel Programming Paradigms

- Message passing: Erlang
- Data Parallel: CUDA
- We've mentioned shared-memory.

Mark Greenstreet	Course Summary	CS 418 – Apr. 5, 2017	6 / 8
------------------	----------------	-----------------------	-------